



Diffusion Models: From Synthetic Data to Detecting the Unknown

A journey through how diffusion models are transforming medical imaging research.

CB

by Cosmin I. Bercea

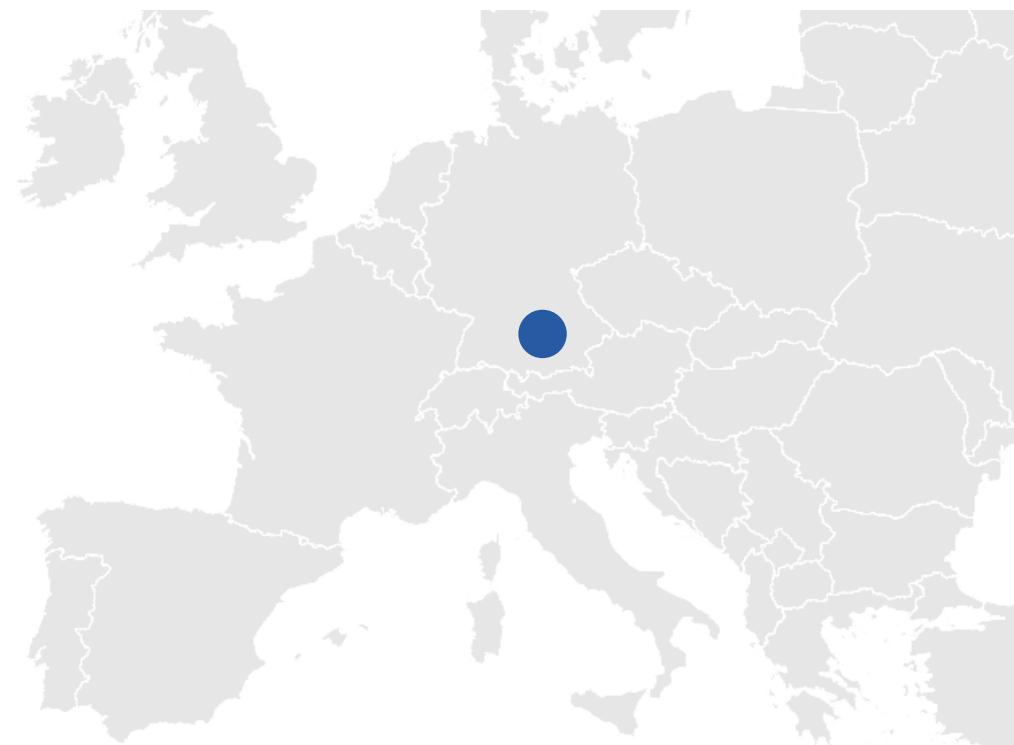
Hello from Munich



<https://compai-lab.github.io>

<https://ai-idt.github.io>

<https://aim-lab.io>



**Julia
Schnabel**



**Bene
Wiestler**



**Daniel
Rueckert**

Hello from Munich



Cosmin I.
Bercea

Sibiu



SIEMENS
Healthineers

B.Sc. / M. Sc.

Informatics



Erlangen

2011-
2018

Barcelona



M.Sc. Intern
AI in Computer
Vision



2016

PhD Intern
AI in Medicine



London

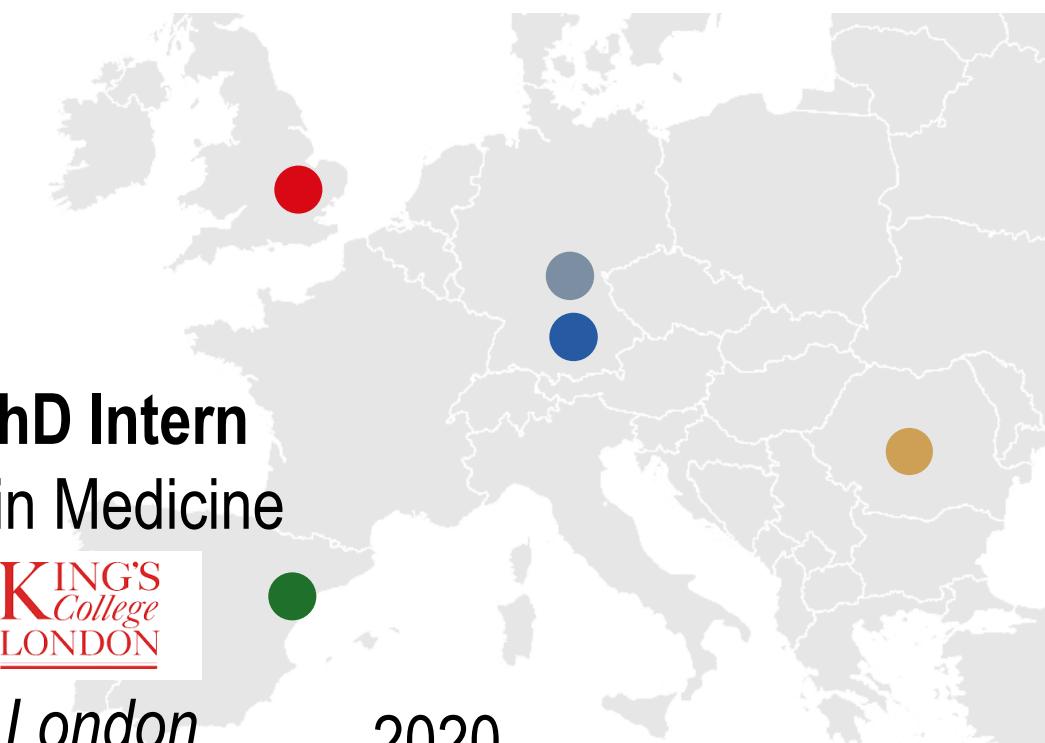
2022

2020-
2024

Munich

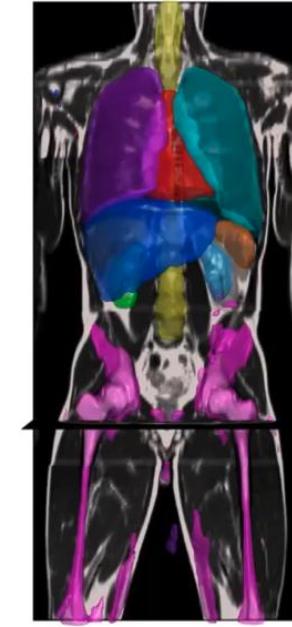
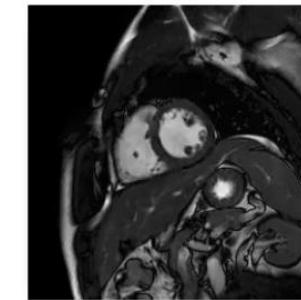
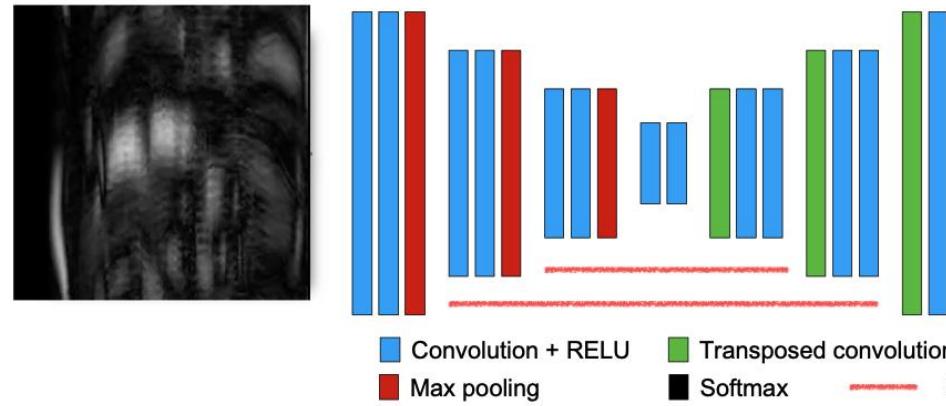


Dr. rer. nat
AI in Medicine

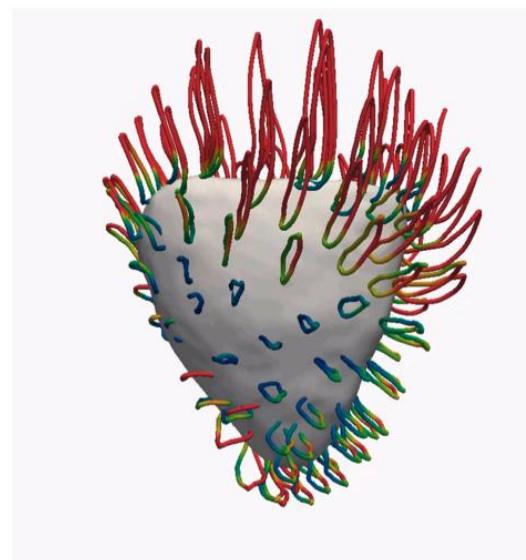


Postdoctoral
researcher

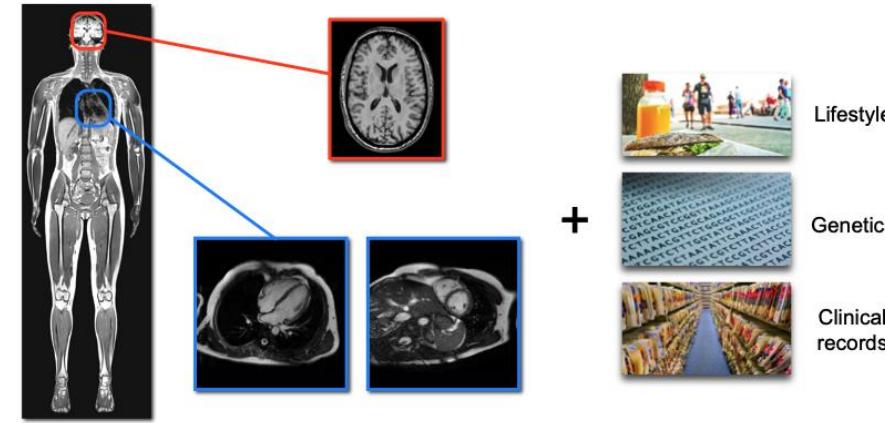
Our Research Topics



Medical Image Computing

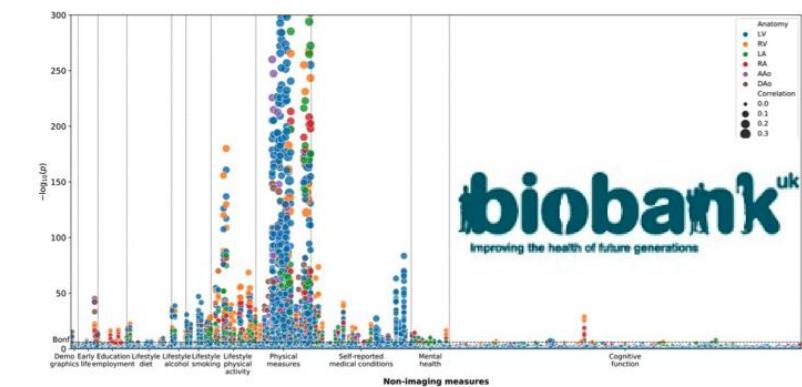


Intelligent Imaging



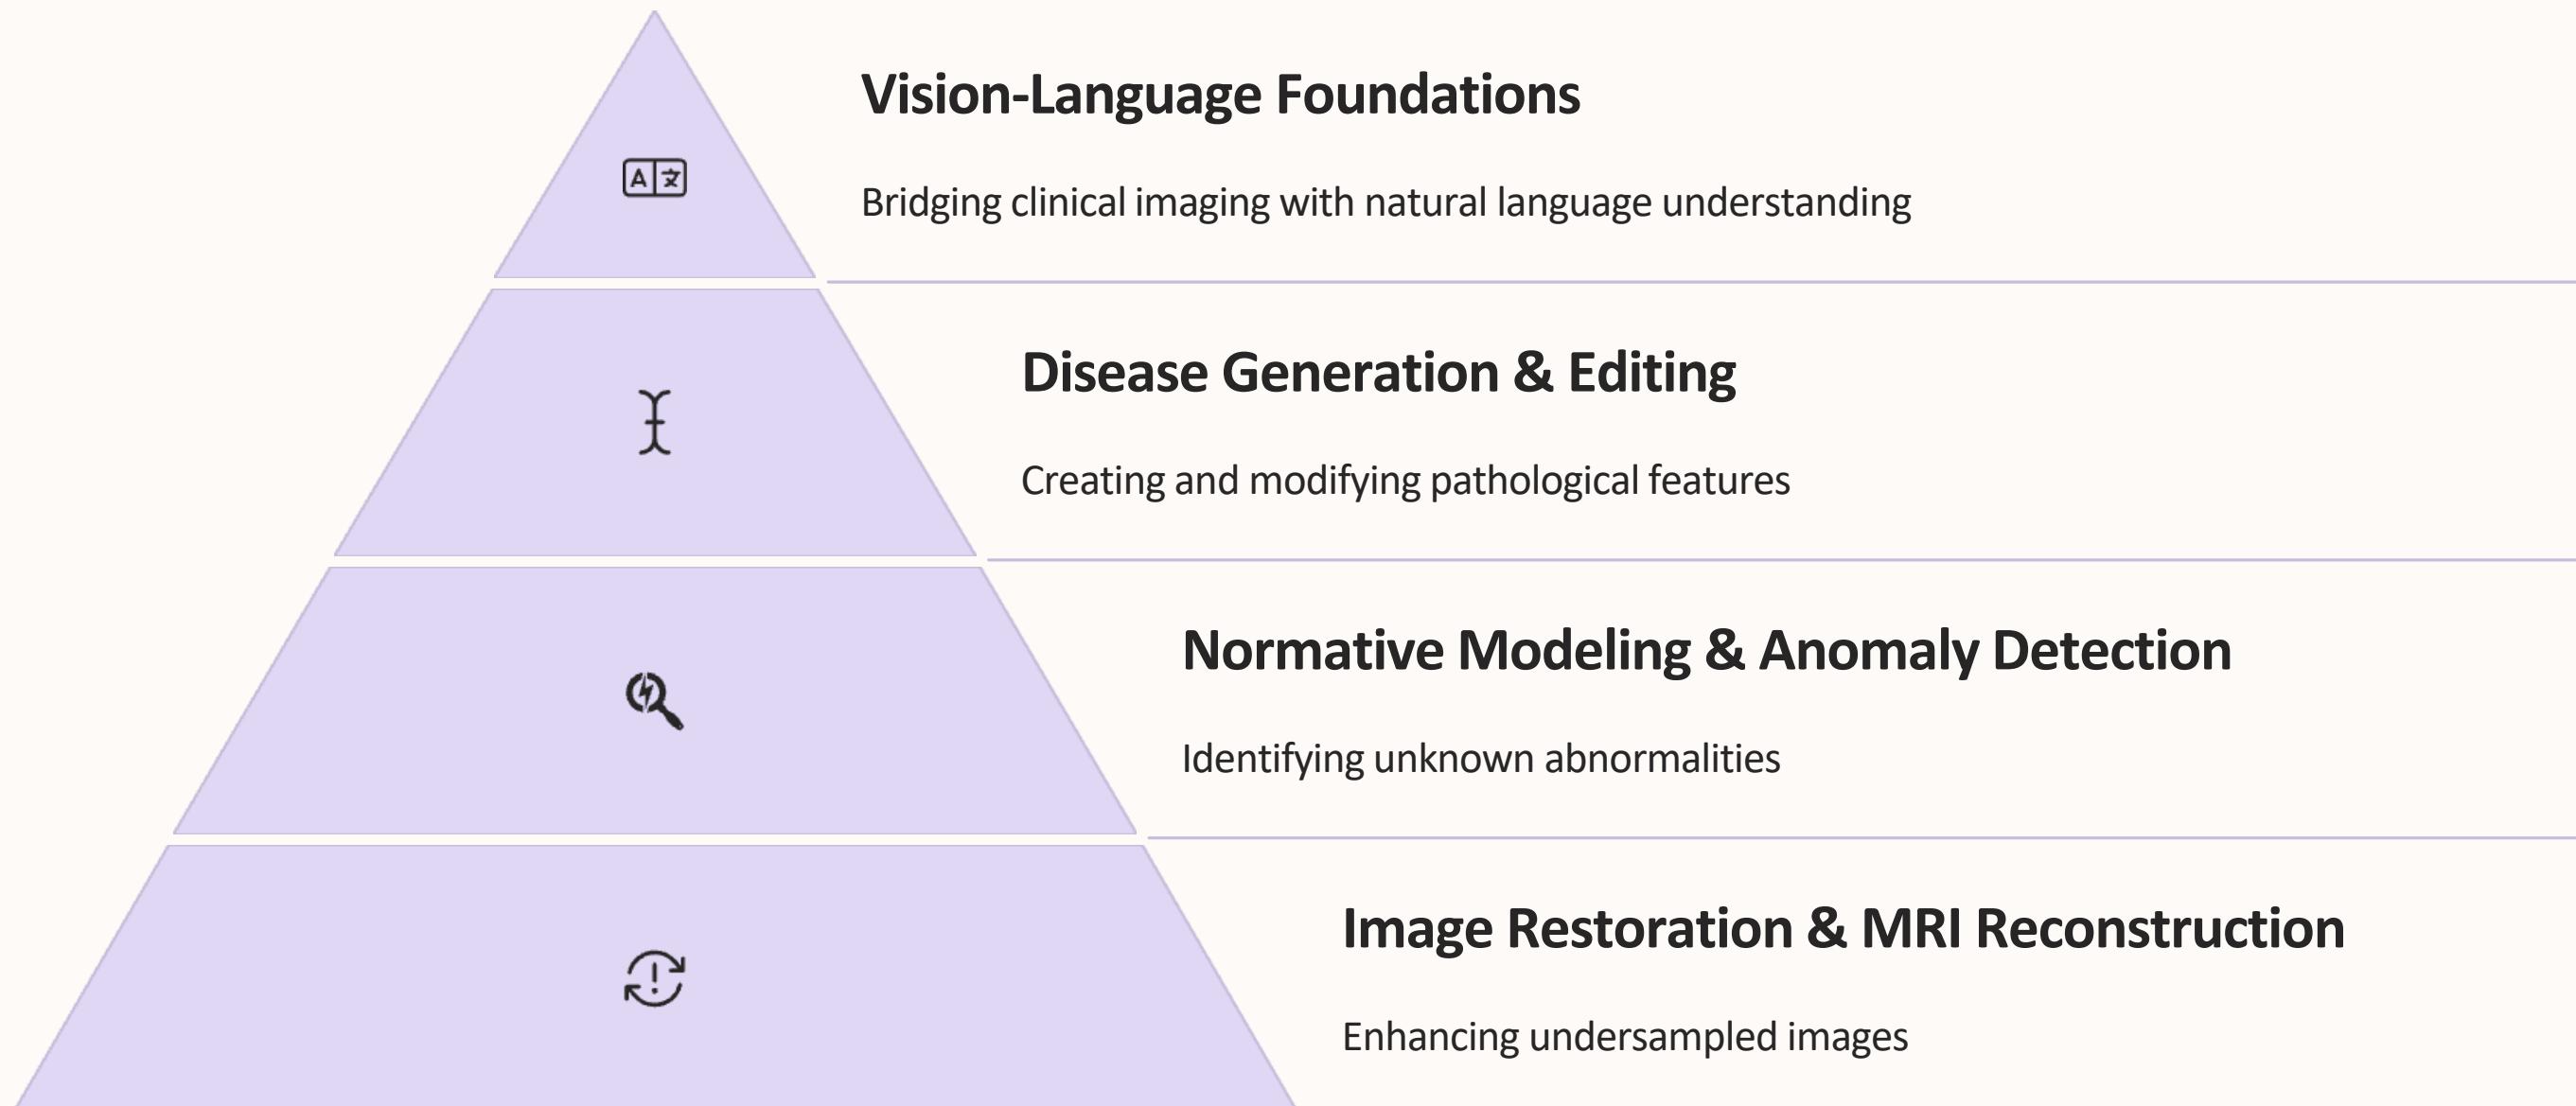
Biomedical Computing and Modelling

Multi-modal learning



Population studies

Our Research Topics



AI in Medicine – Hype vs. Hope



**Self-driving
cars will be
here in two
years.**

2015



**Radiologists
will be
replaced in
5 years**

2016



Radiologists going to work

AI in Medicine – Hype vs. Hope



MIT
Technology
Review

Artificial intelligence / Machine learning

Hundreds of AI tools have been built to catch covid. None of them helped.

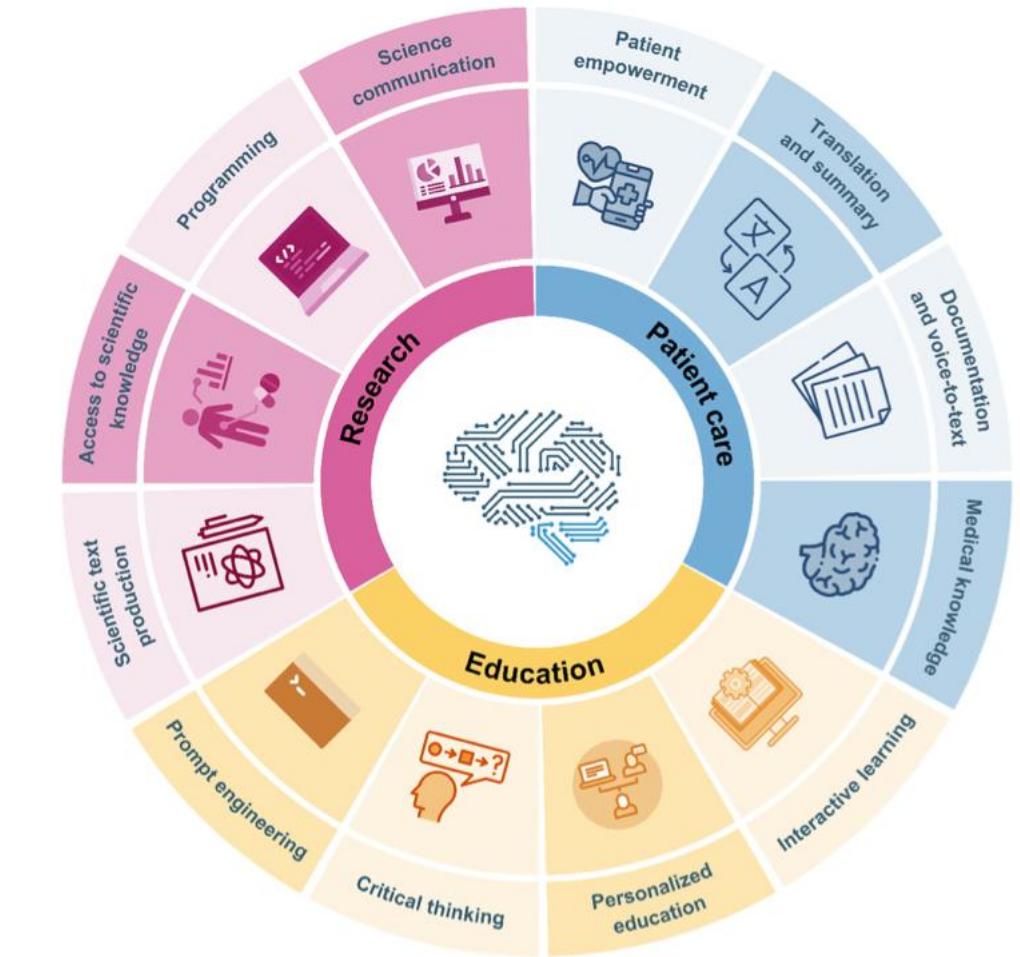
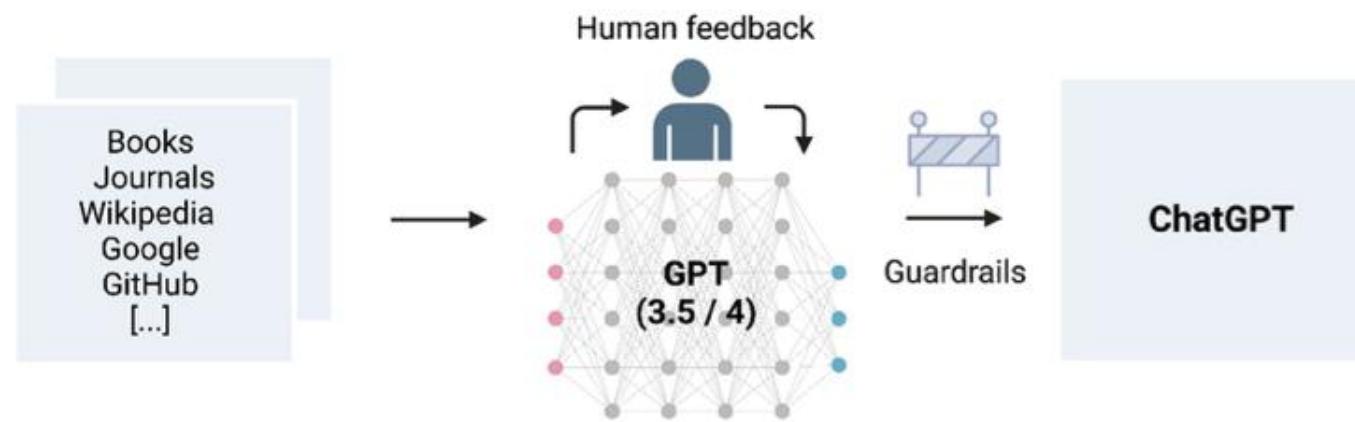
Some have been used in hospitals, despite not being properly tested. But the pandemic could help make medical AI better.

by Will Douglas Heaven

July 30, 2021

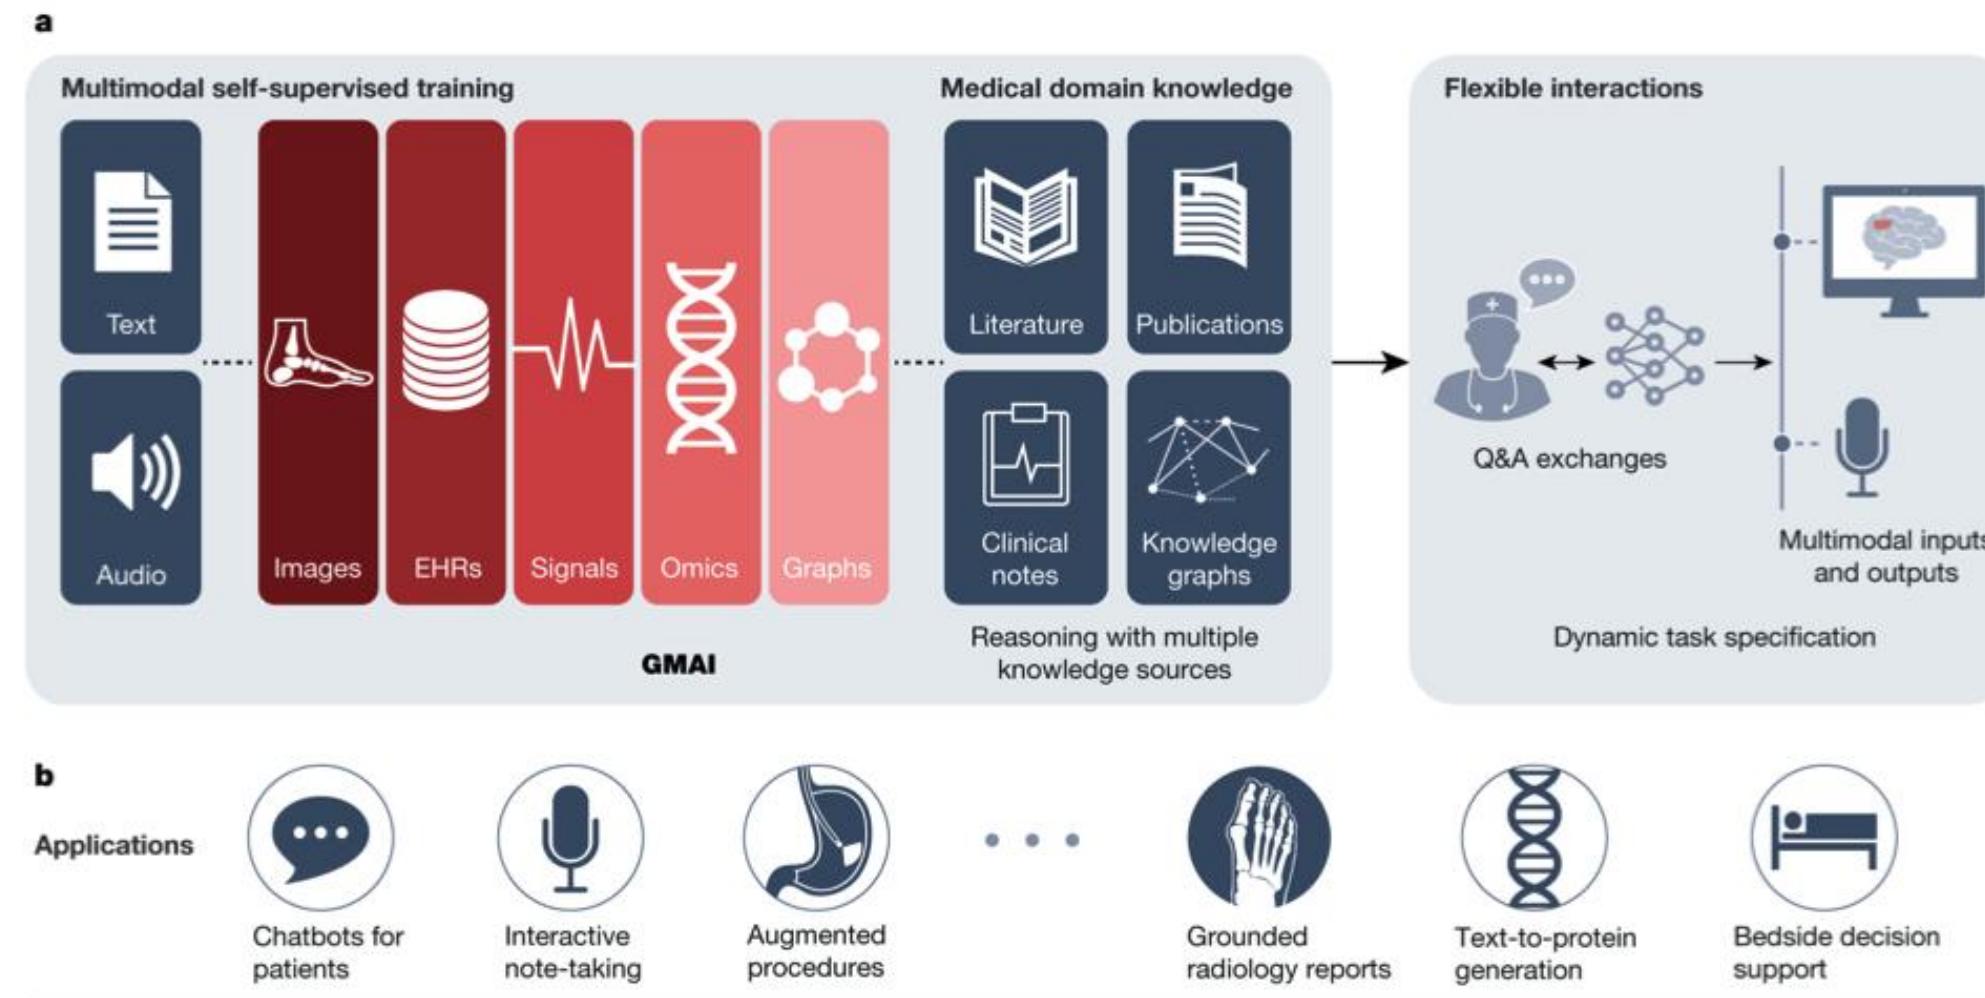
AI Trends Shaping Imaging

LLMs



Overview of potential LLM
applications

AI Trends Shaping Imaging Foundation models



AI Trends Shaping Imaging

Diffusion models



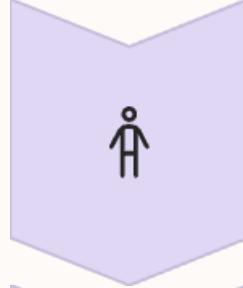
Revolutionizing Diagnostics

Enhance detail in low-quality scans, revealing subtle anomalies.



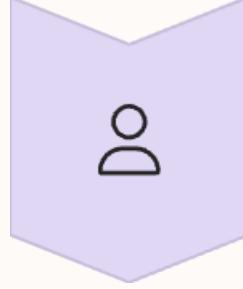
Accelerating Research

Generated datasets accelerate rare disease studies, enabling research without privacy constraints.



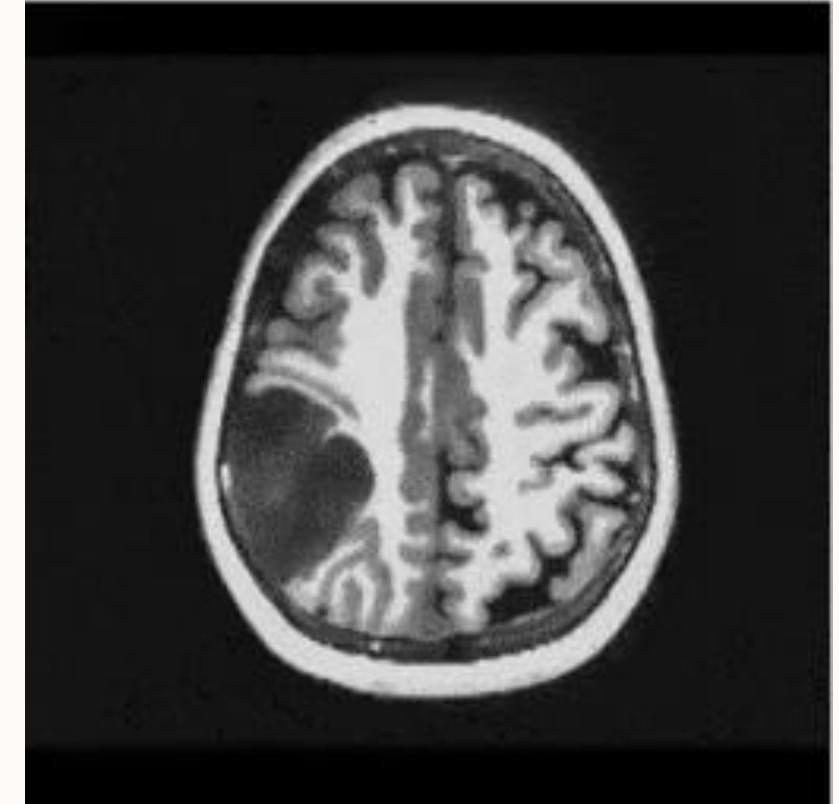
Personalized Medicine

Simulations predict patient-specific treatment responses, visualizing outcomes.

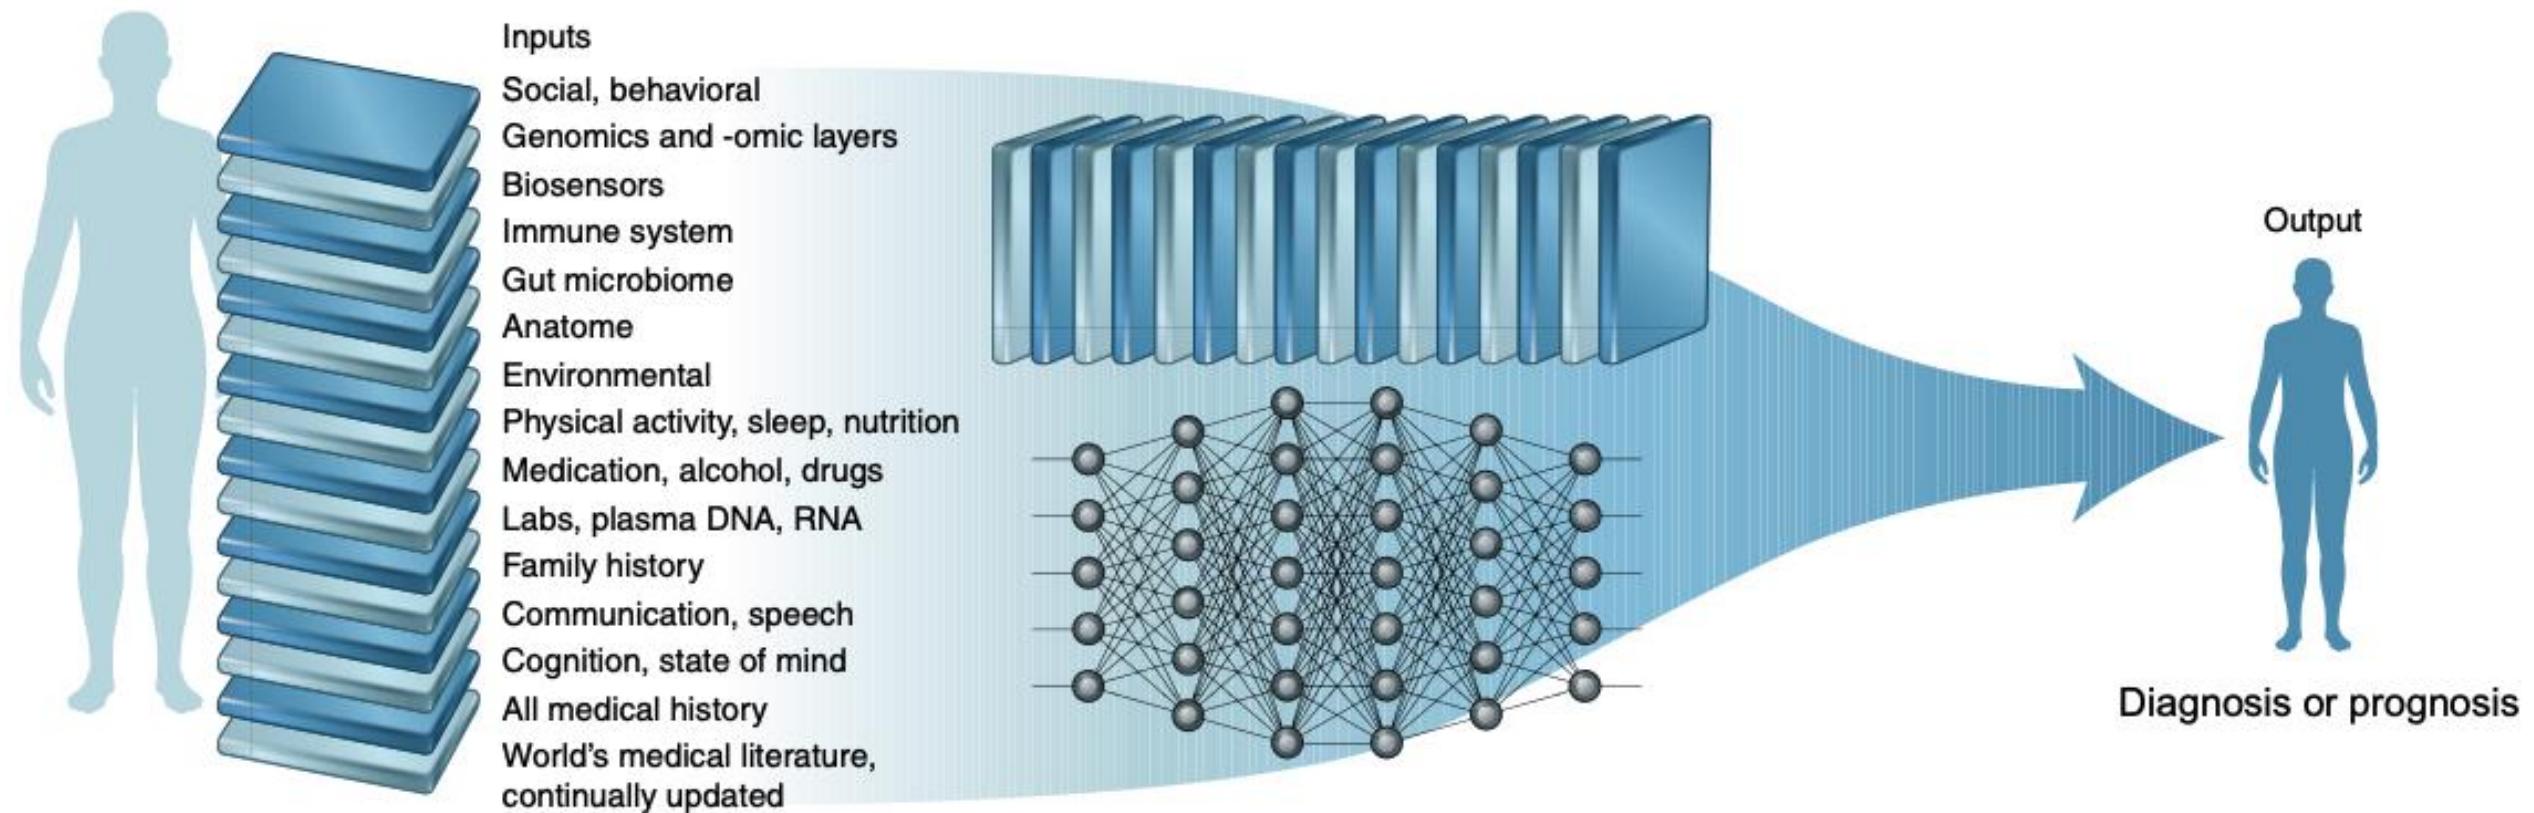


Ethical AI

Synthetic data reduces bias in medical datasets, promoting equitable models.



The Data We Work With



E. Topol, 2017

The Data We Work With



X-Ray & CT

Uses ionizing radiation to create images. Fast but involves radiation exposure.

- CT provides 3D visualization of internal structures
- Widely used for bone and lung imaging

MRI

Uses magnetic fields to generate detailed tissue images. Non-ionizing and safe.

- Superior soft tissue contrast
- Functional MRI captures brain activity

Ultrasound & Nuclear

Ultrasound uses sound waves. Nuclear imaging requires radioactive tracers.

- Ultrasound is real-time and non-invasive
- PET/SPECT show metabolic function

Challenges

- High volume of scans



Challenges

- # Radiology scans
- # Trained radiologists

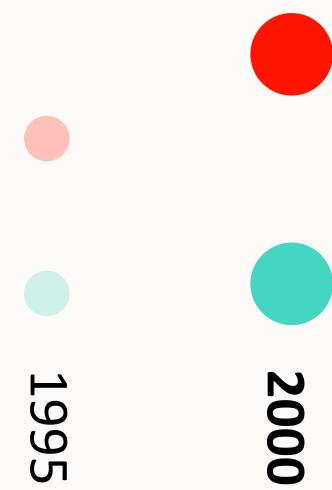
High volume of scans



Challenges

High volume of scans

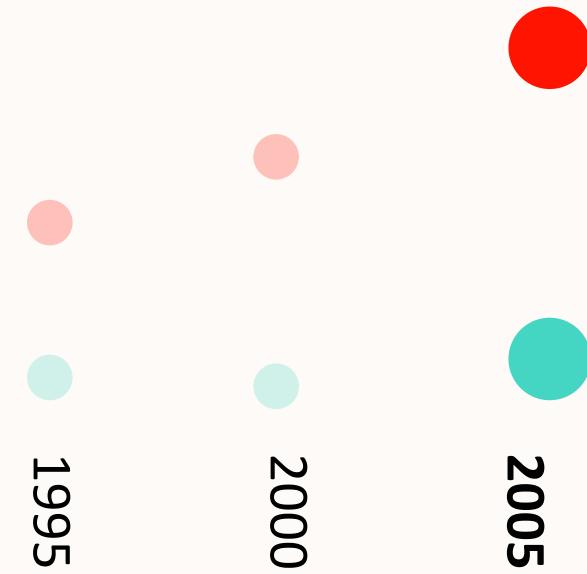
- # Radiology scans
- # Trained radiologists



Challenges

High volume of scans

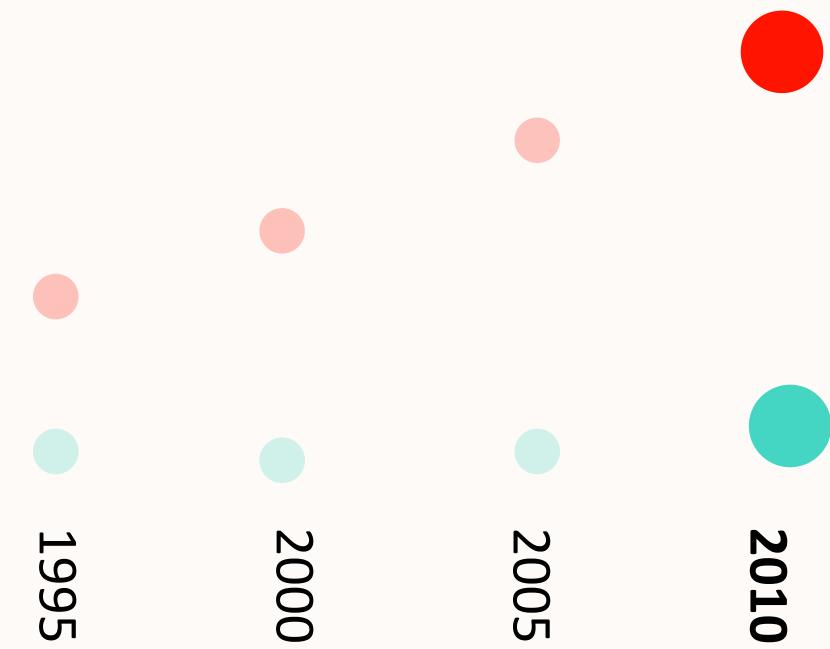
- # Radiology scans
- # Trained radiologists



Challenges

- # Radiology scans
- # Trained radiologists

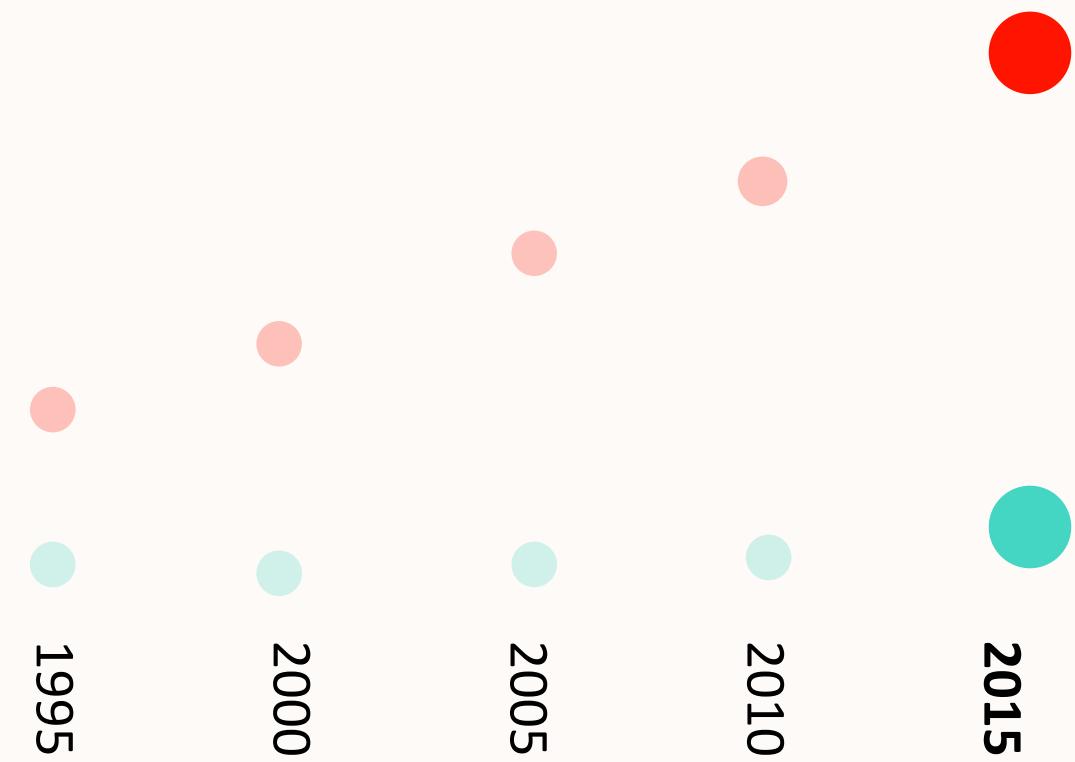
High volume of scans



Challenges

High volume of scans

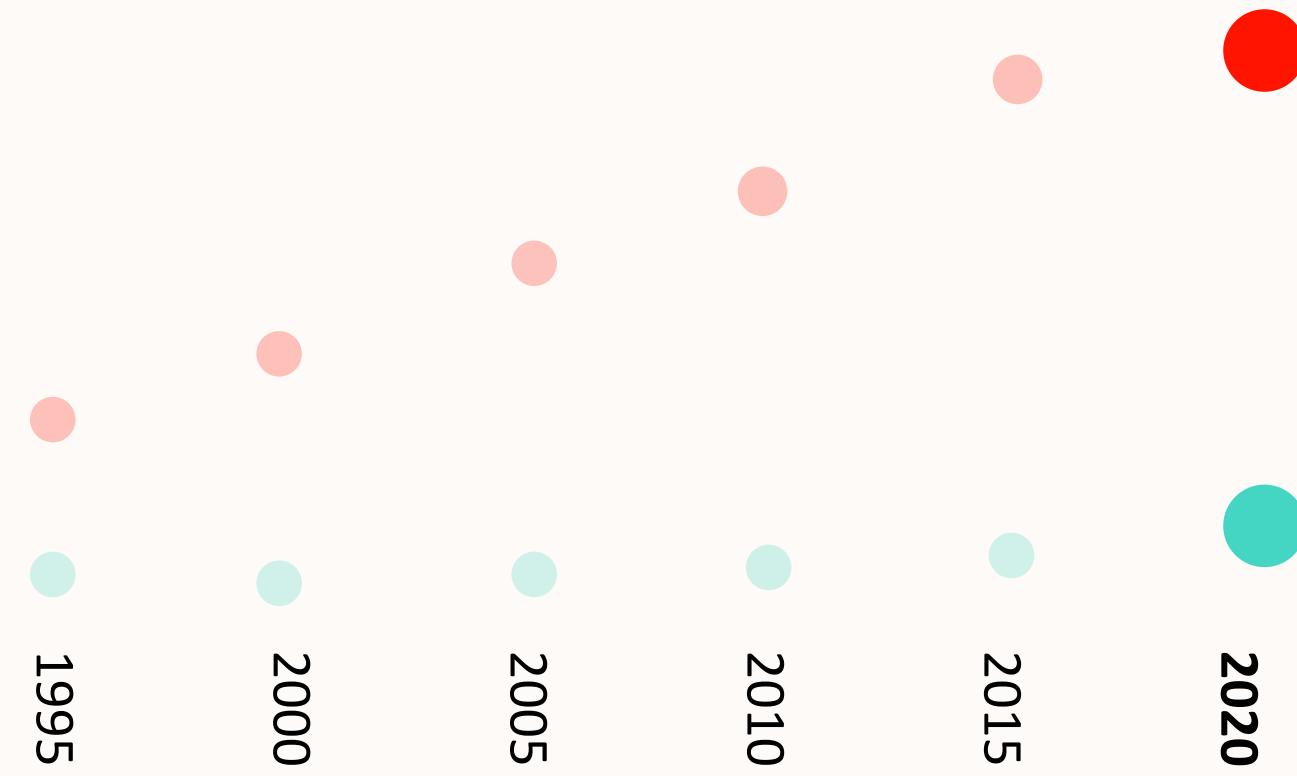
- # Radiology scans
- # Trained radiologists



Challenges

High volume of scans

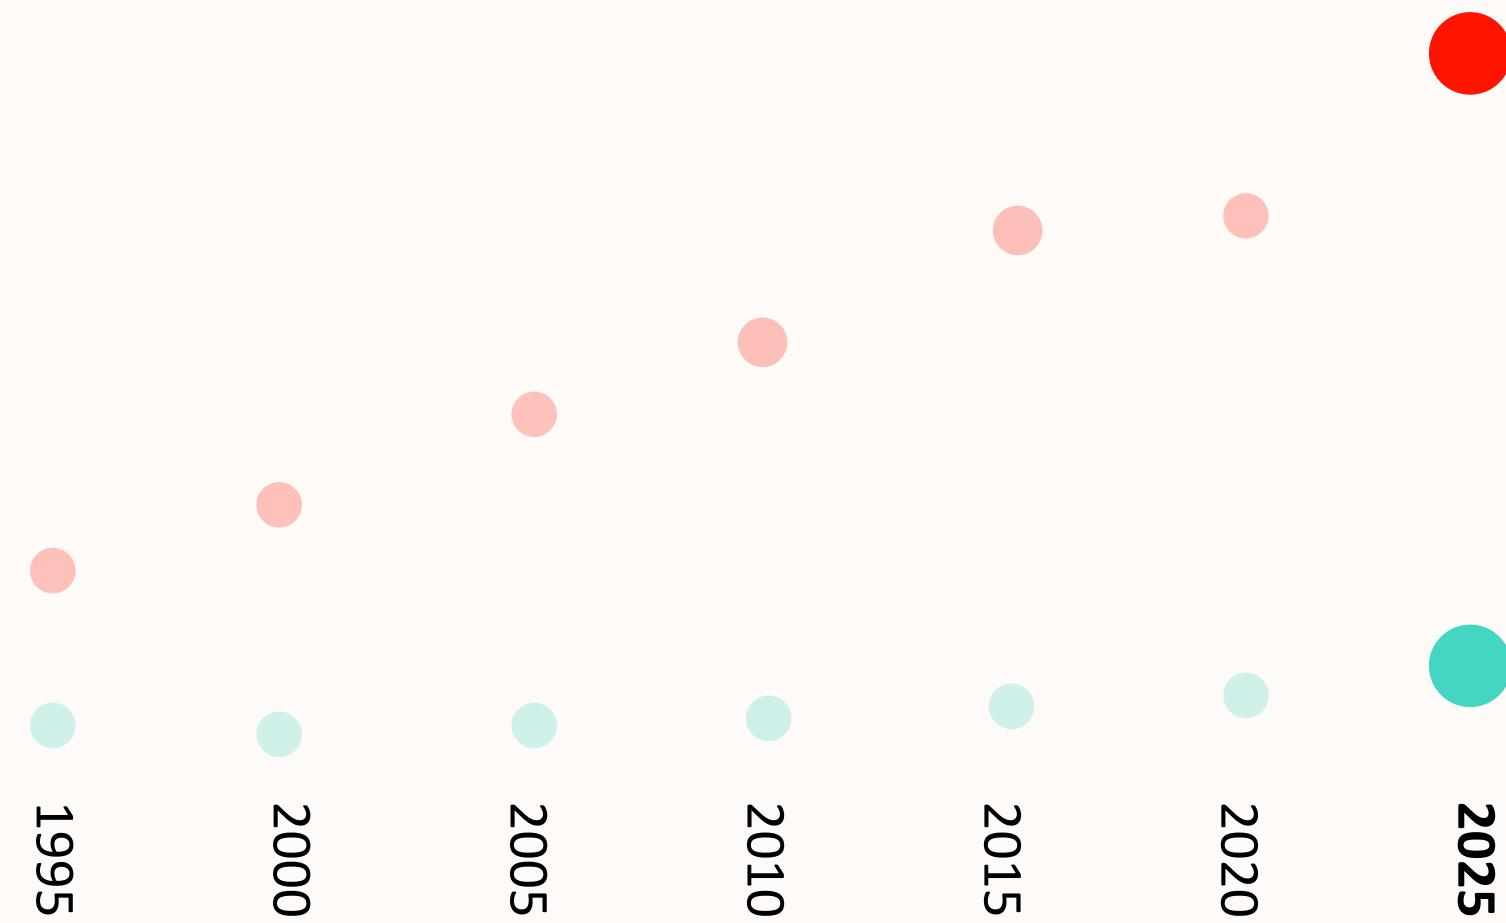
- # Radiology scans
- # Trained radiologists



Challenges

High volume of scans

- # Radiology scans
- # Trained radiologists

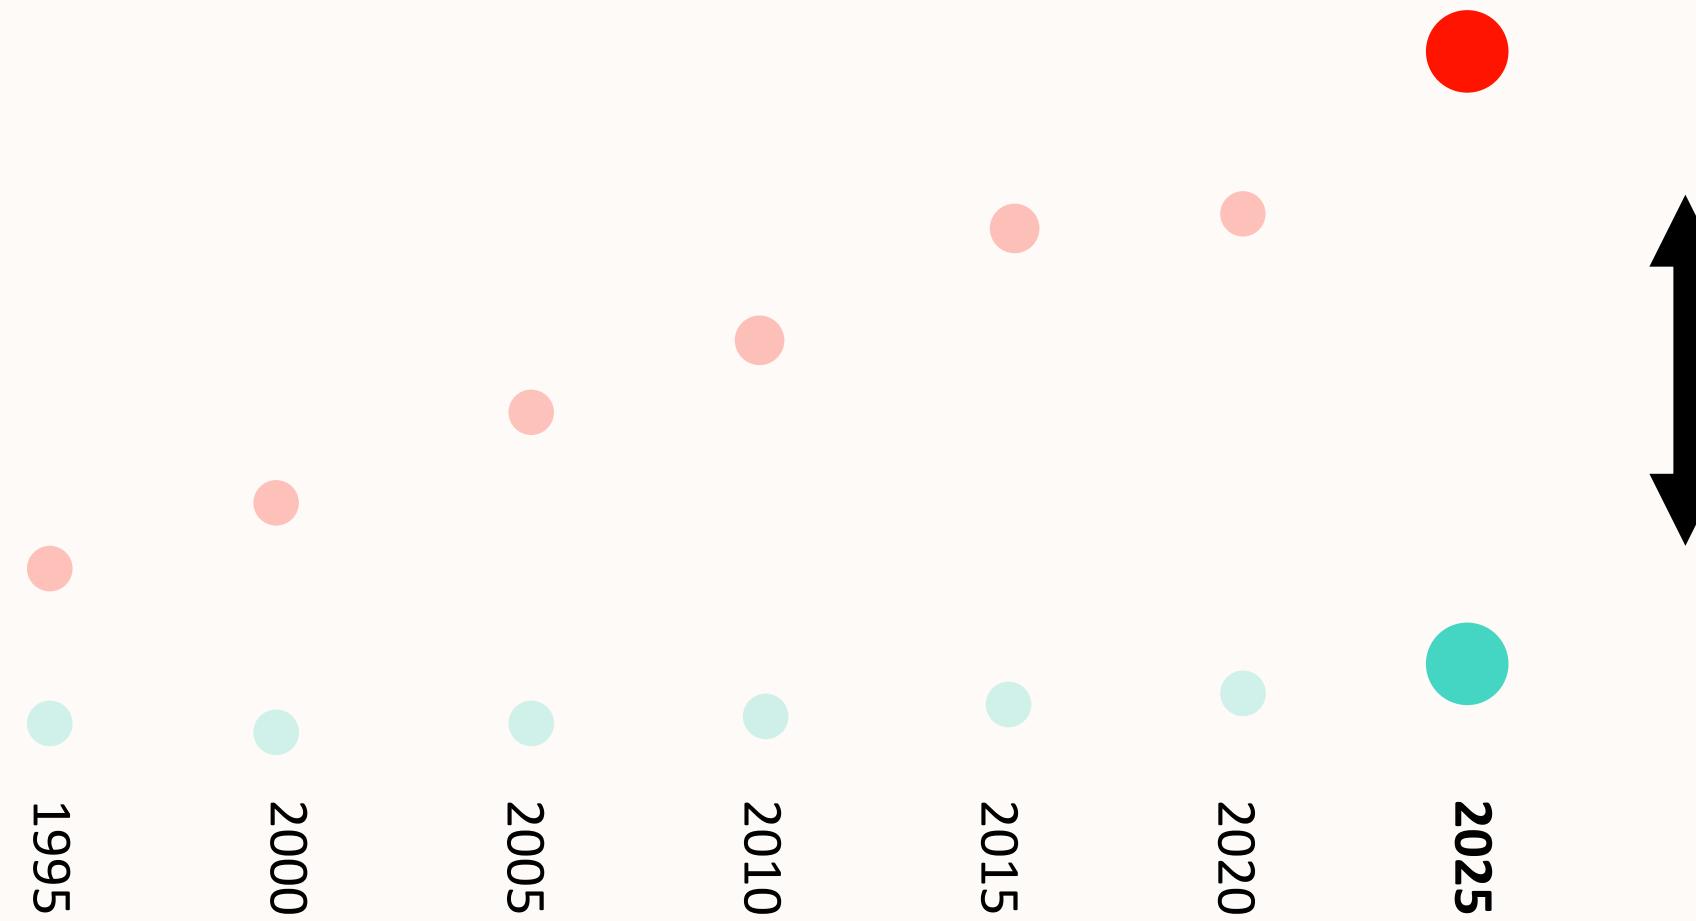


Challenges

High volume of scans

- # Radiology scans
- # Trained radiologists

Trained radiologists



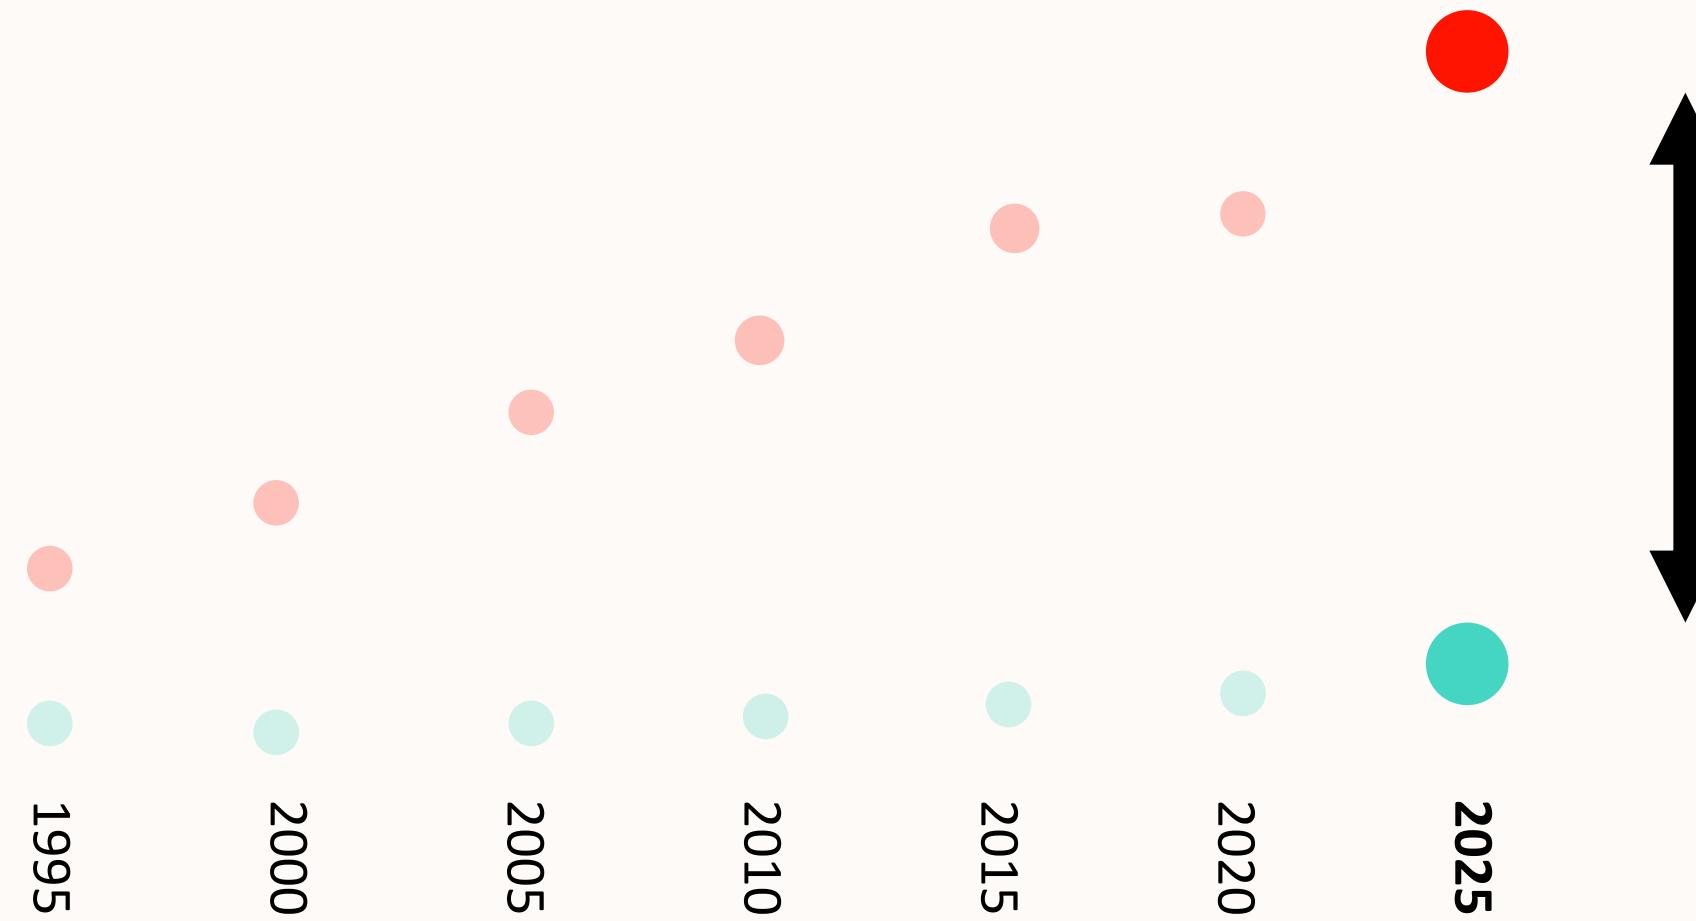
Radiologist
shortage

Challenges

High volume of scans

- # Radiology scans
- # Trained radiologists

Trained radiologists



0 %

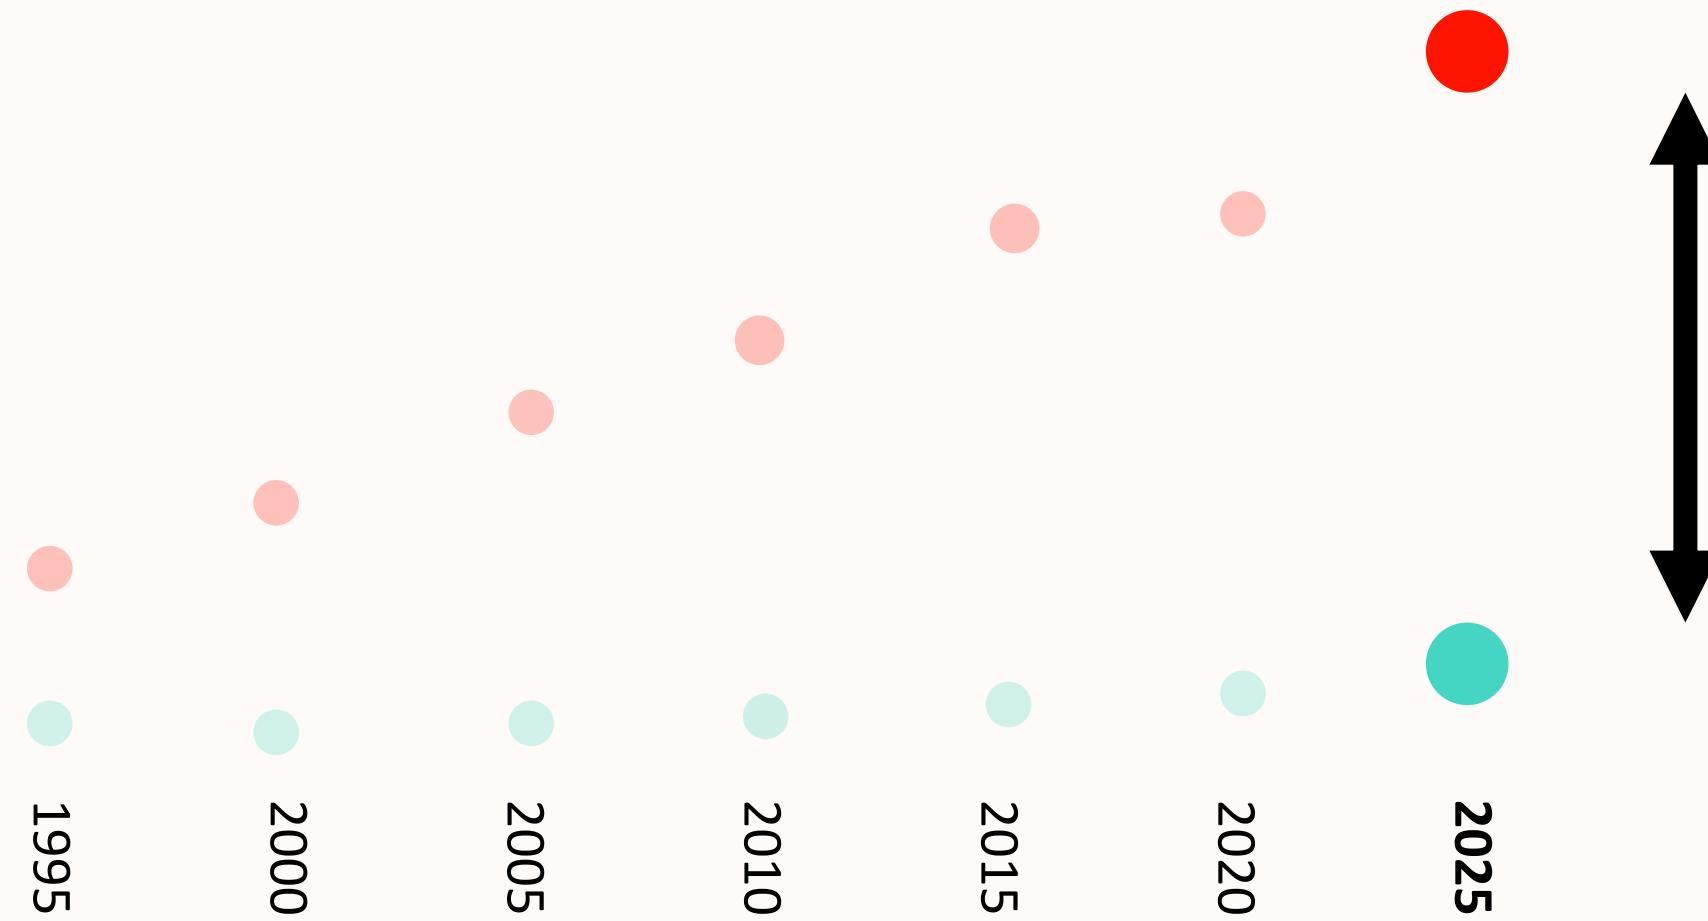
Radiologist
shortage

Challenges

High volume of scans

- # Radiology scans
- # Trained radiologists

Trained radiologists



50 %

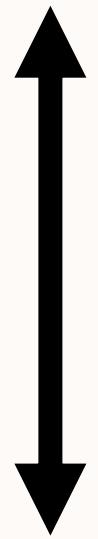
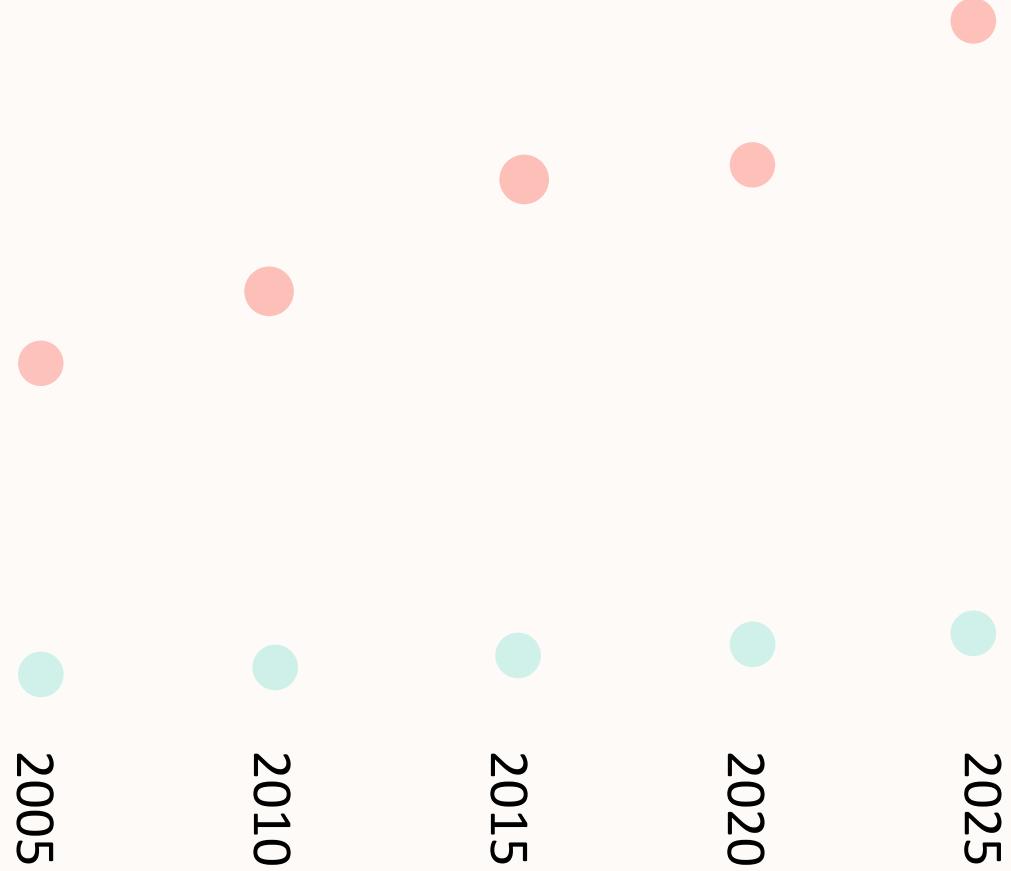
Radiologist
shortage

Challenges

- # Radiology scans
- # Trained radiologists

Trained radiologists

High volume of scans



50 %

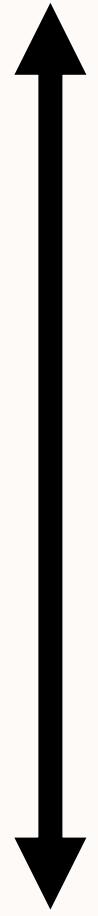
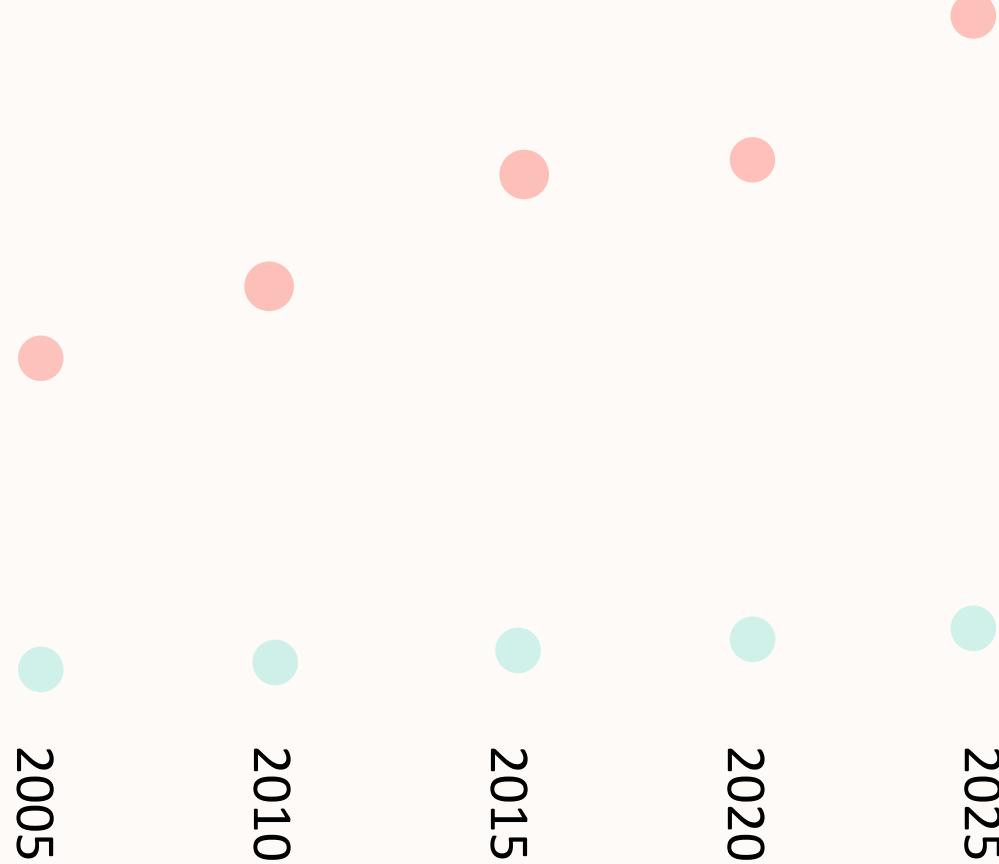
Radiologist
shortage

Challenges

- # Radiology scans
- # Trained radiologists

Trained radiologists

High volume of scans



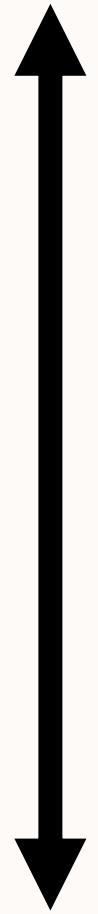
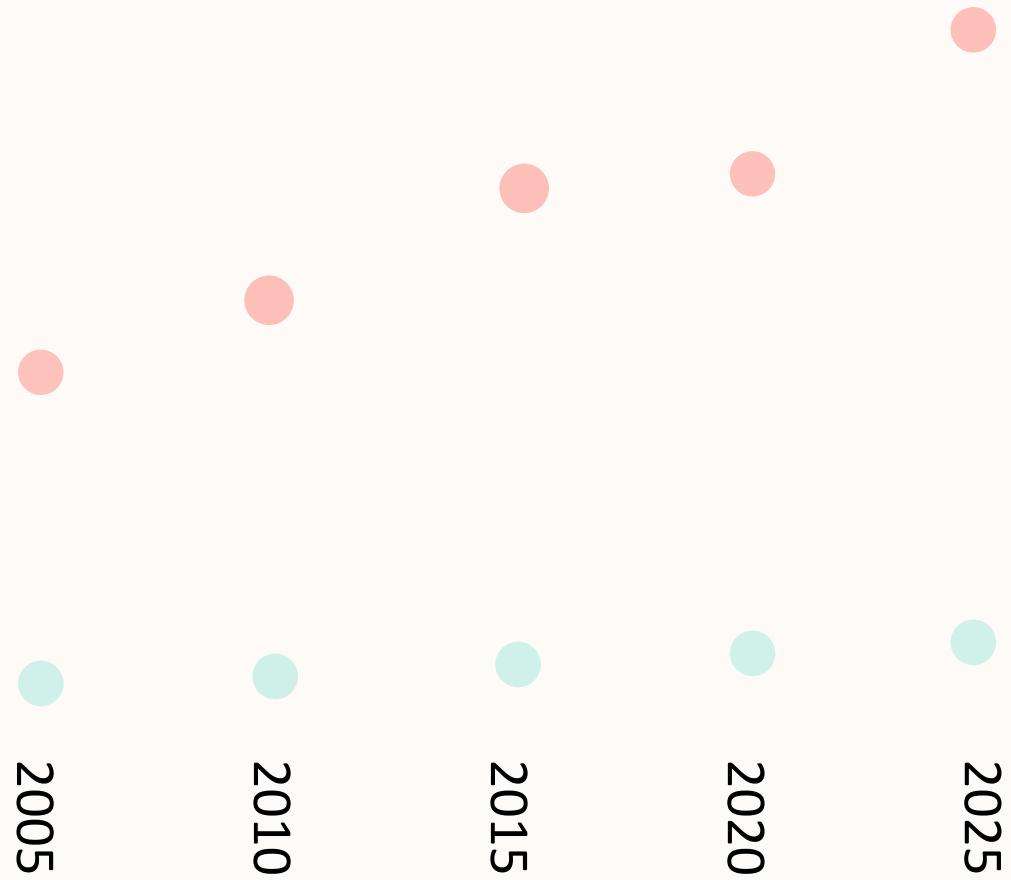
50 %
Radiologist
shortage

Challenges

- # Radiology scans
- # Trained radiologists

Trained radiologists

High volume of scans



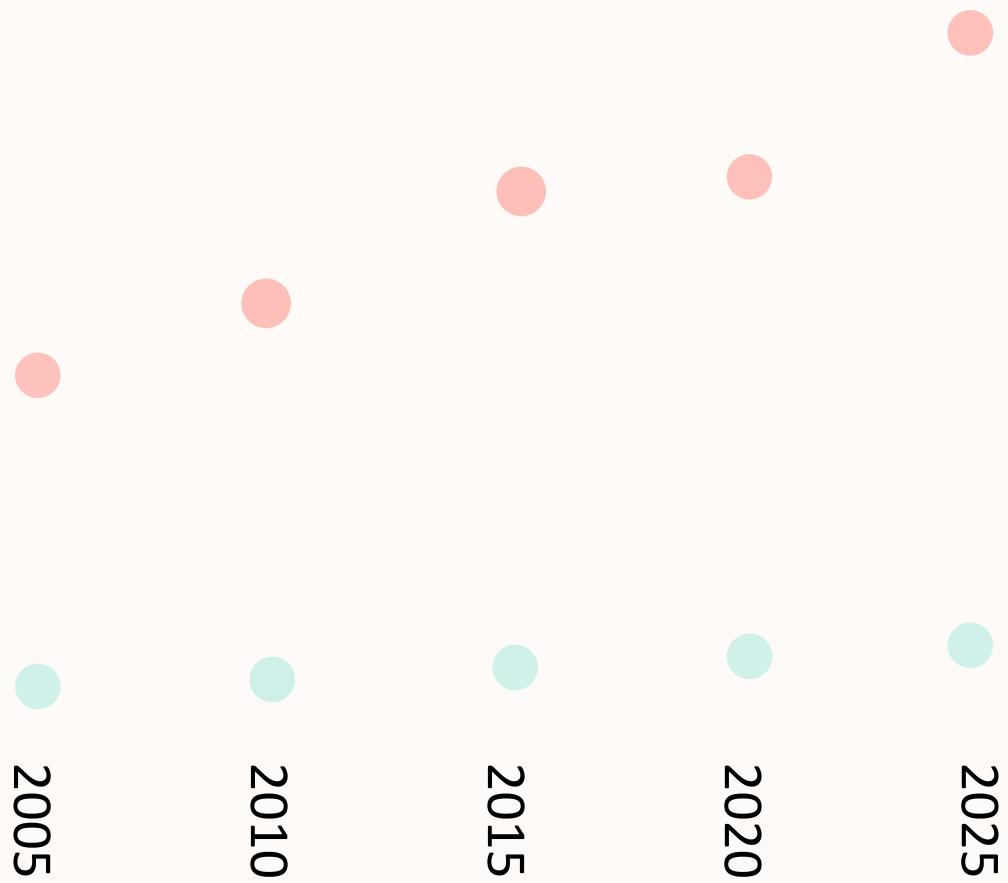
50%
Radiologist
shortage

Challenges

- # Radiology scans
- # Trained radiologists

Trained radiologists

High volume of scans



75 %
Radiologist
shortage

Challenges

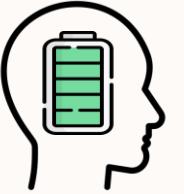
Trained radiologists

75 %

Radiologist
shortage

Challenges

Trained radiologists



75 %

Radiologist
shortage



Challenges

Trained radiologists



75 %

Radiologist
shortage



Challenges

75 %

Radiologist
shortage

Trained radiologists



Increased workloads

Strain on radiologists with
increased workloads



Challenges

75 %
Radiologist
shortage

Trained radiologists



Increased workloads

Strain on radiologists with
increased workloads

Reporting delays

> 220K scans reported after 28
days in the UK in 2024



Challenges

75 %
Radiologist
shortage

Trained radiologists



Increased workloads

Strain on radiologists with
increased workloads

Reporting delays

> 220K scans reported after 28
days in the UK in 2024

Missed pathologies

> 40 Million errors worldwide

Challenges

75 %
Radiologist
shortage

Trained radiologists



0 %

Increased workloads

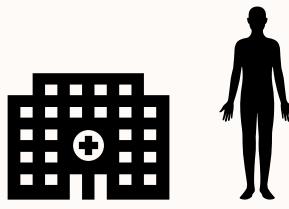
Strain on radiologists with
increased workloads

Reporting delays

> 220K scans reported after 28
days in the UK in 2024

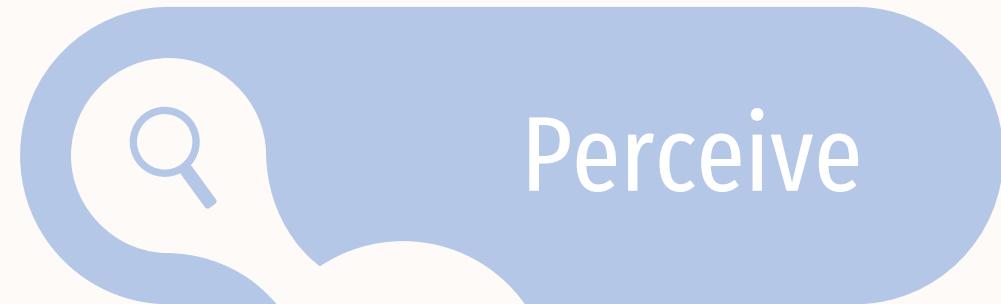
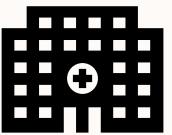
Missed pathologies

> 40 Million errors worldwide



Where do most errors occur?

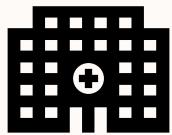
Where do most errors occur?



Detect and localize
radiological findings

Where do most errors occur?

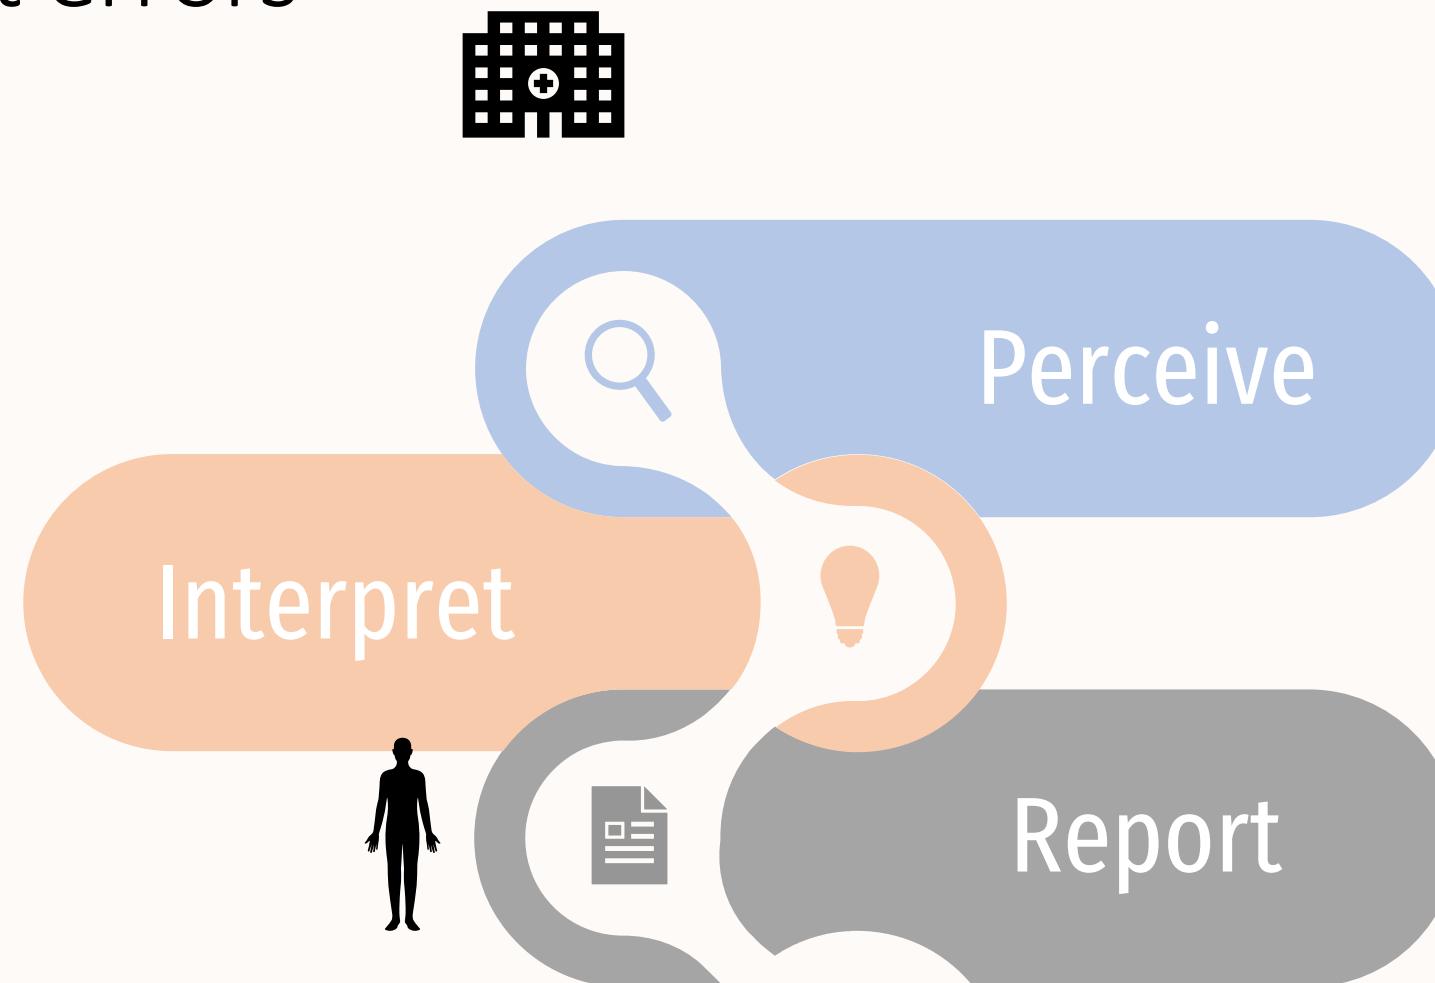
Diagnostic inference
Decision Making



Detect and localize
radiological findings

Where do most errors occur?

Diagnostic inference
Decision Making



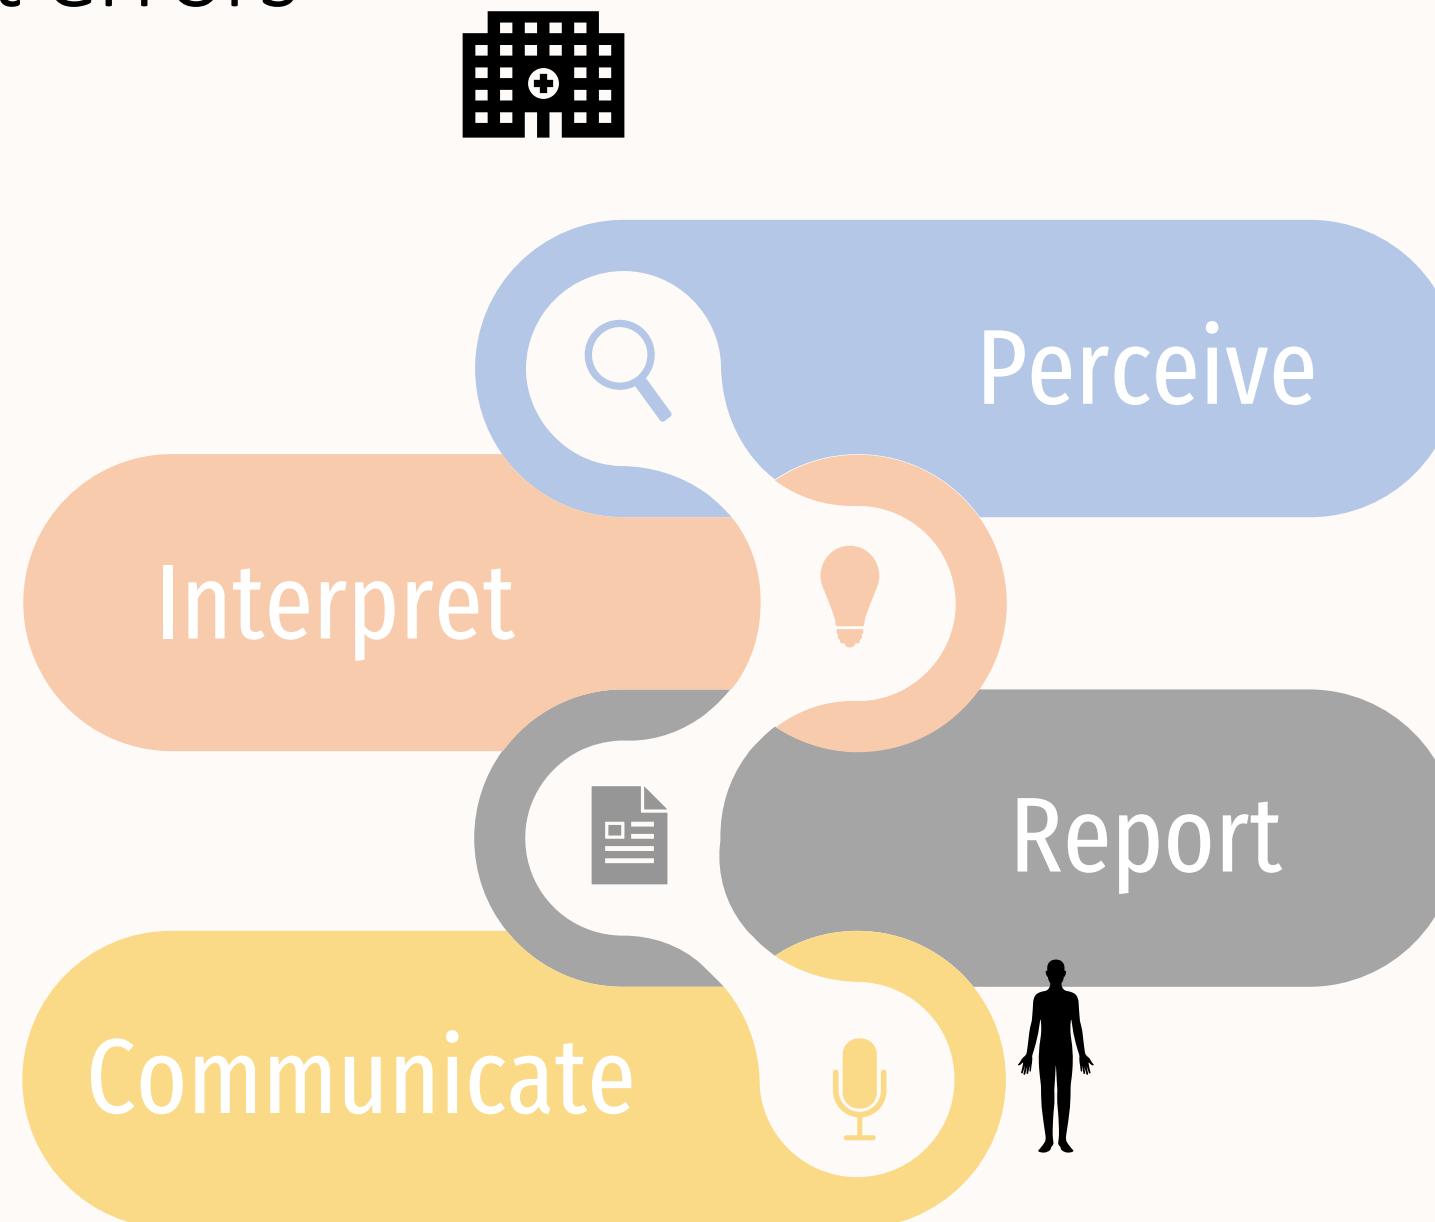
Detect and localize
radiological findings

Documentation of the
radiological findings

Where do most errors occur?

Diagnostic inference
Decision Making

Results communication
Actioning
Treatment planning



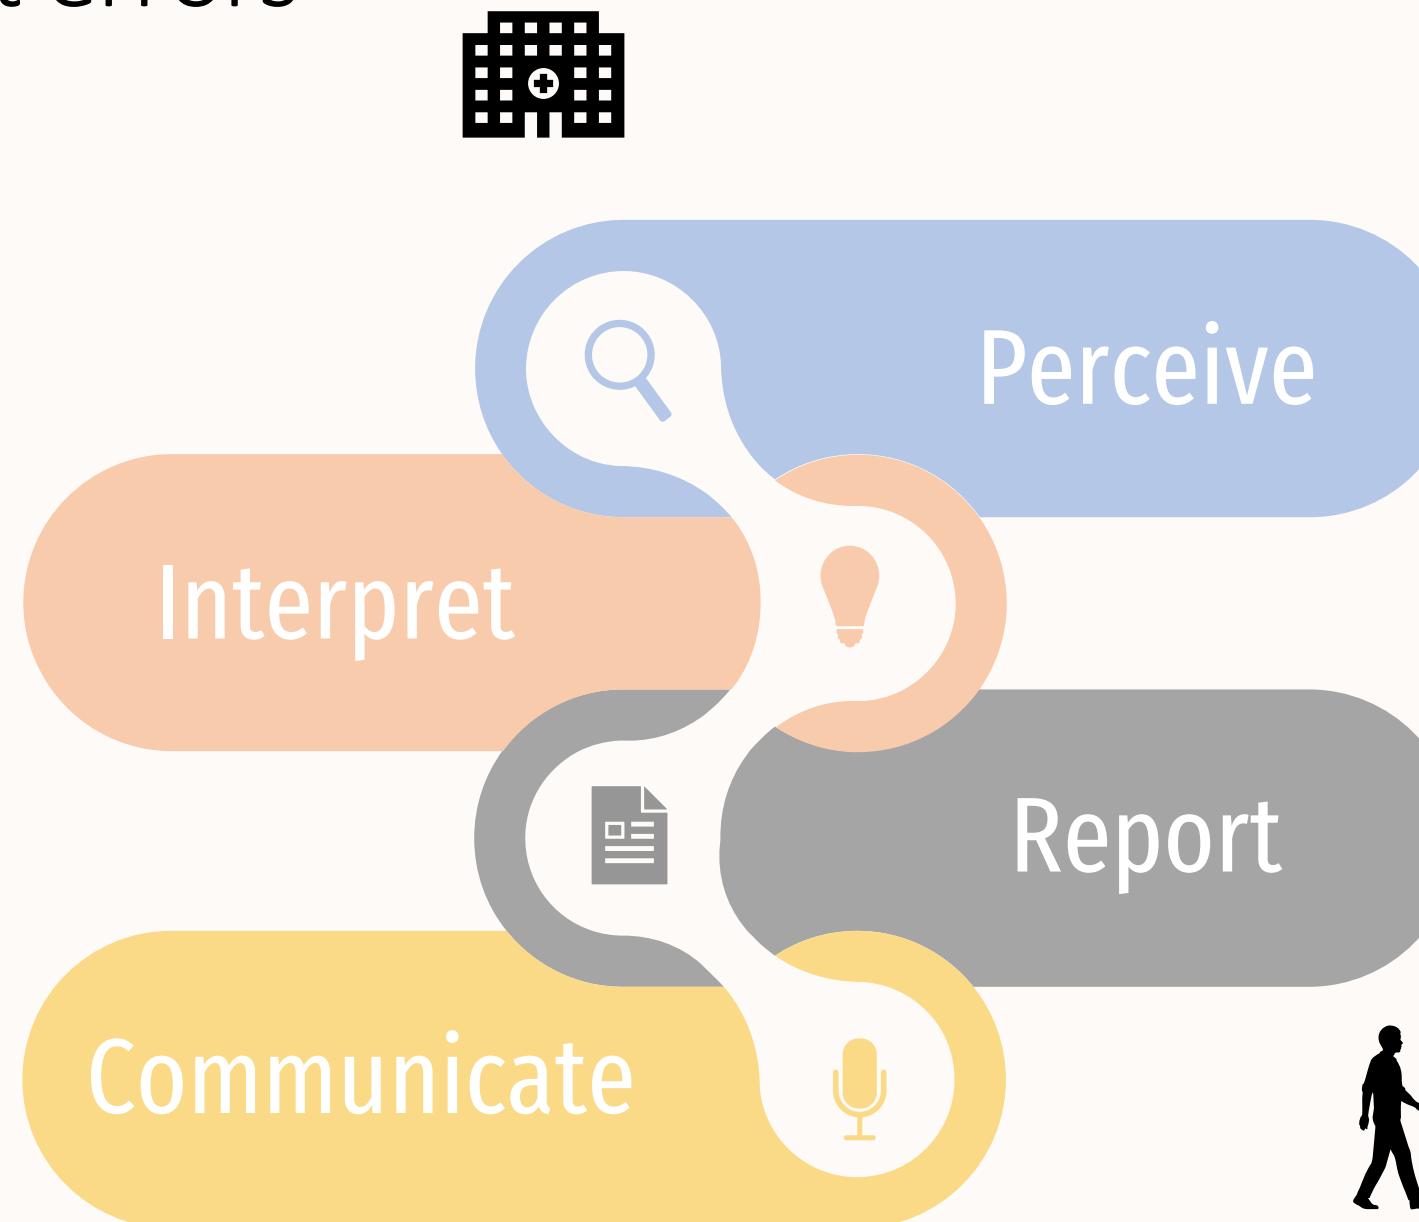
Detect and localize radiological findings

Documentation of the radiological findings

Where do most errors occur?

Diagnostic inference
Decision Making

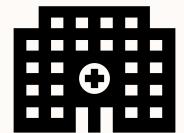
Results communication
Actioning
Treatment planning



Detect and localize
radiological findings

Documentation of the
radiological findings

Where do most errors occur?



Most human errors made

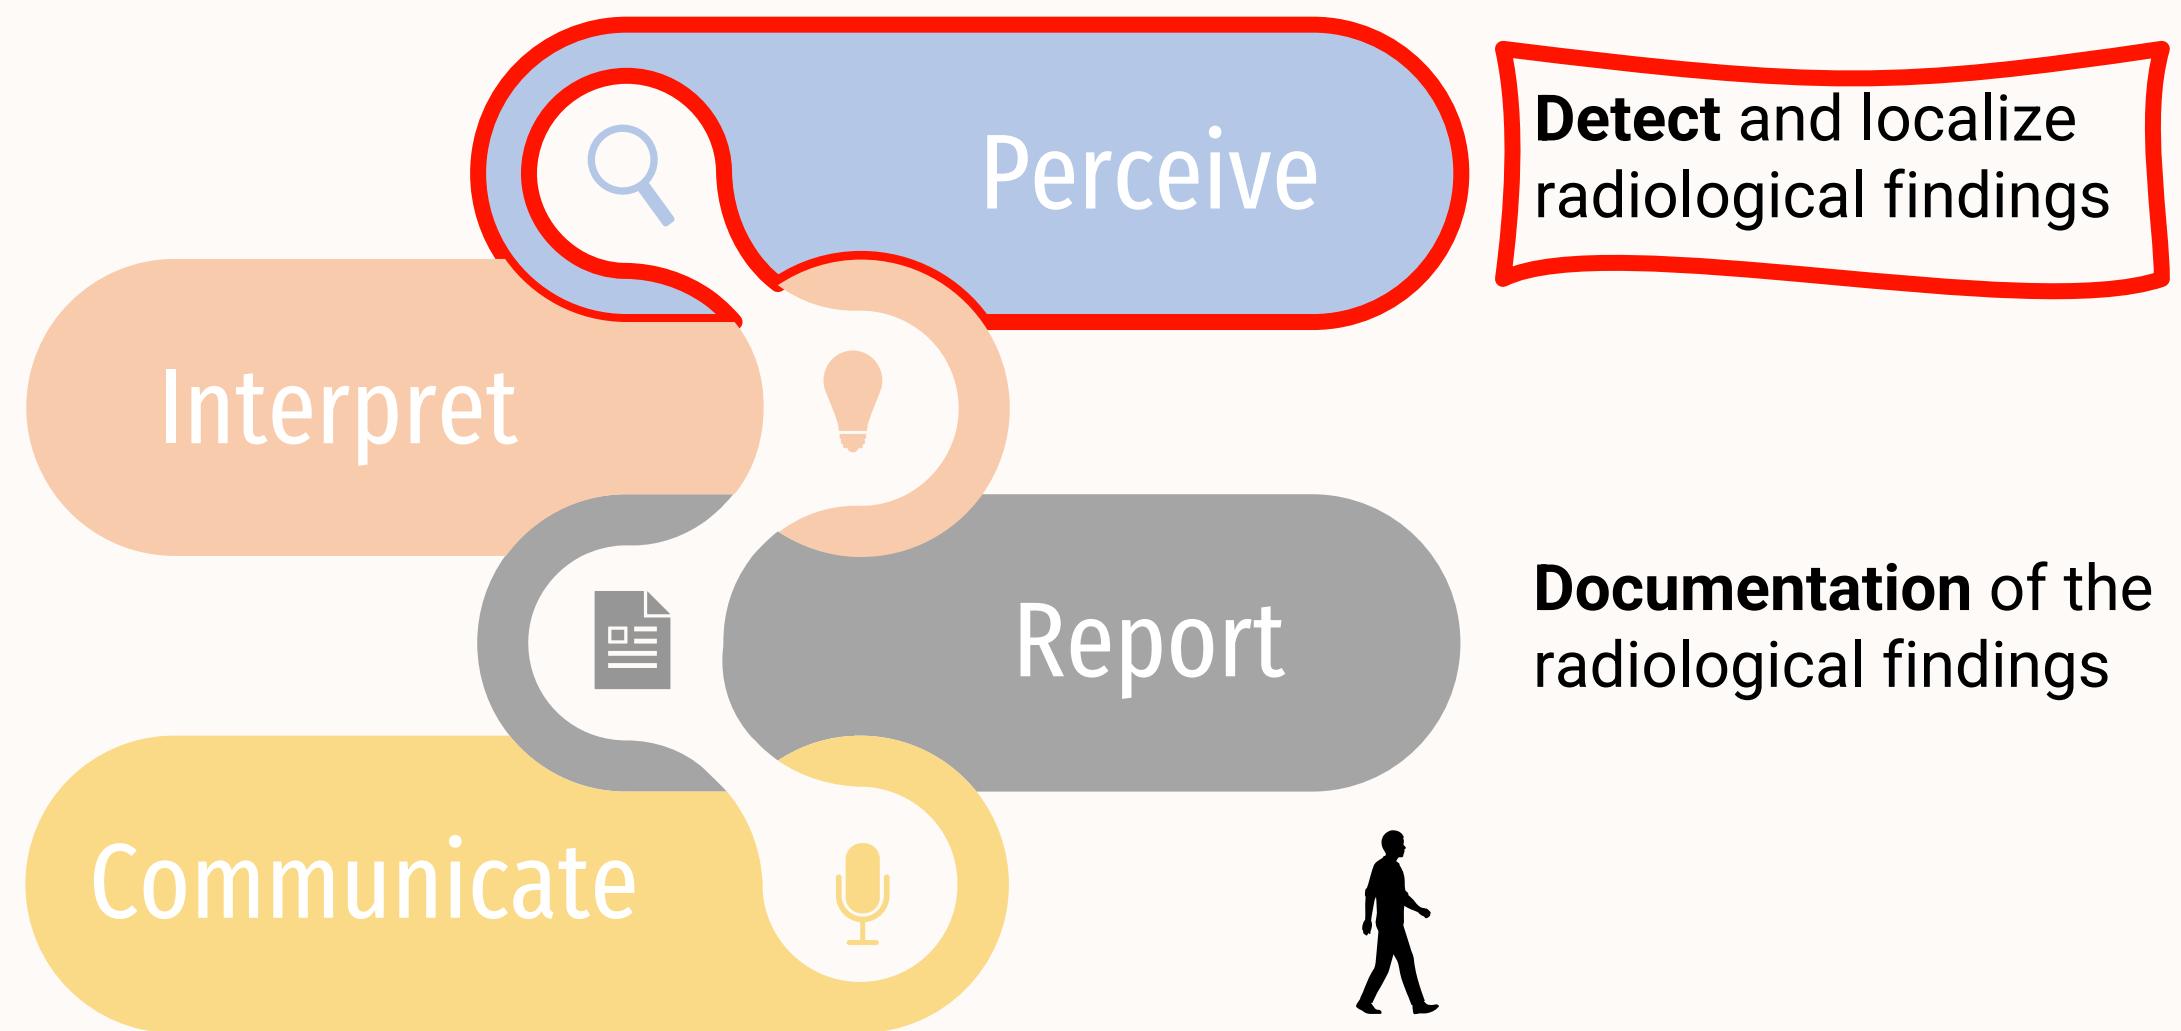
Diagnostic inference
Decision Making

Results communication
Actioning
Treatment planning



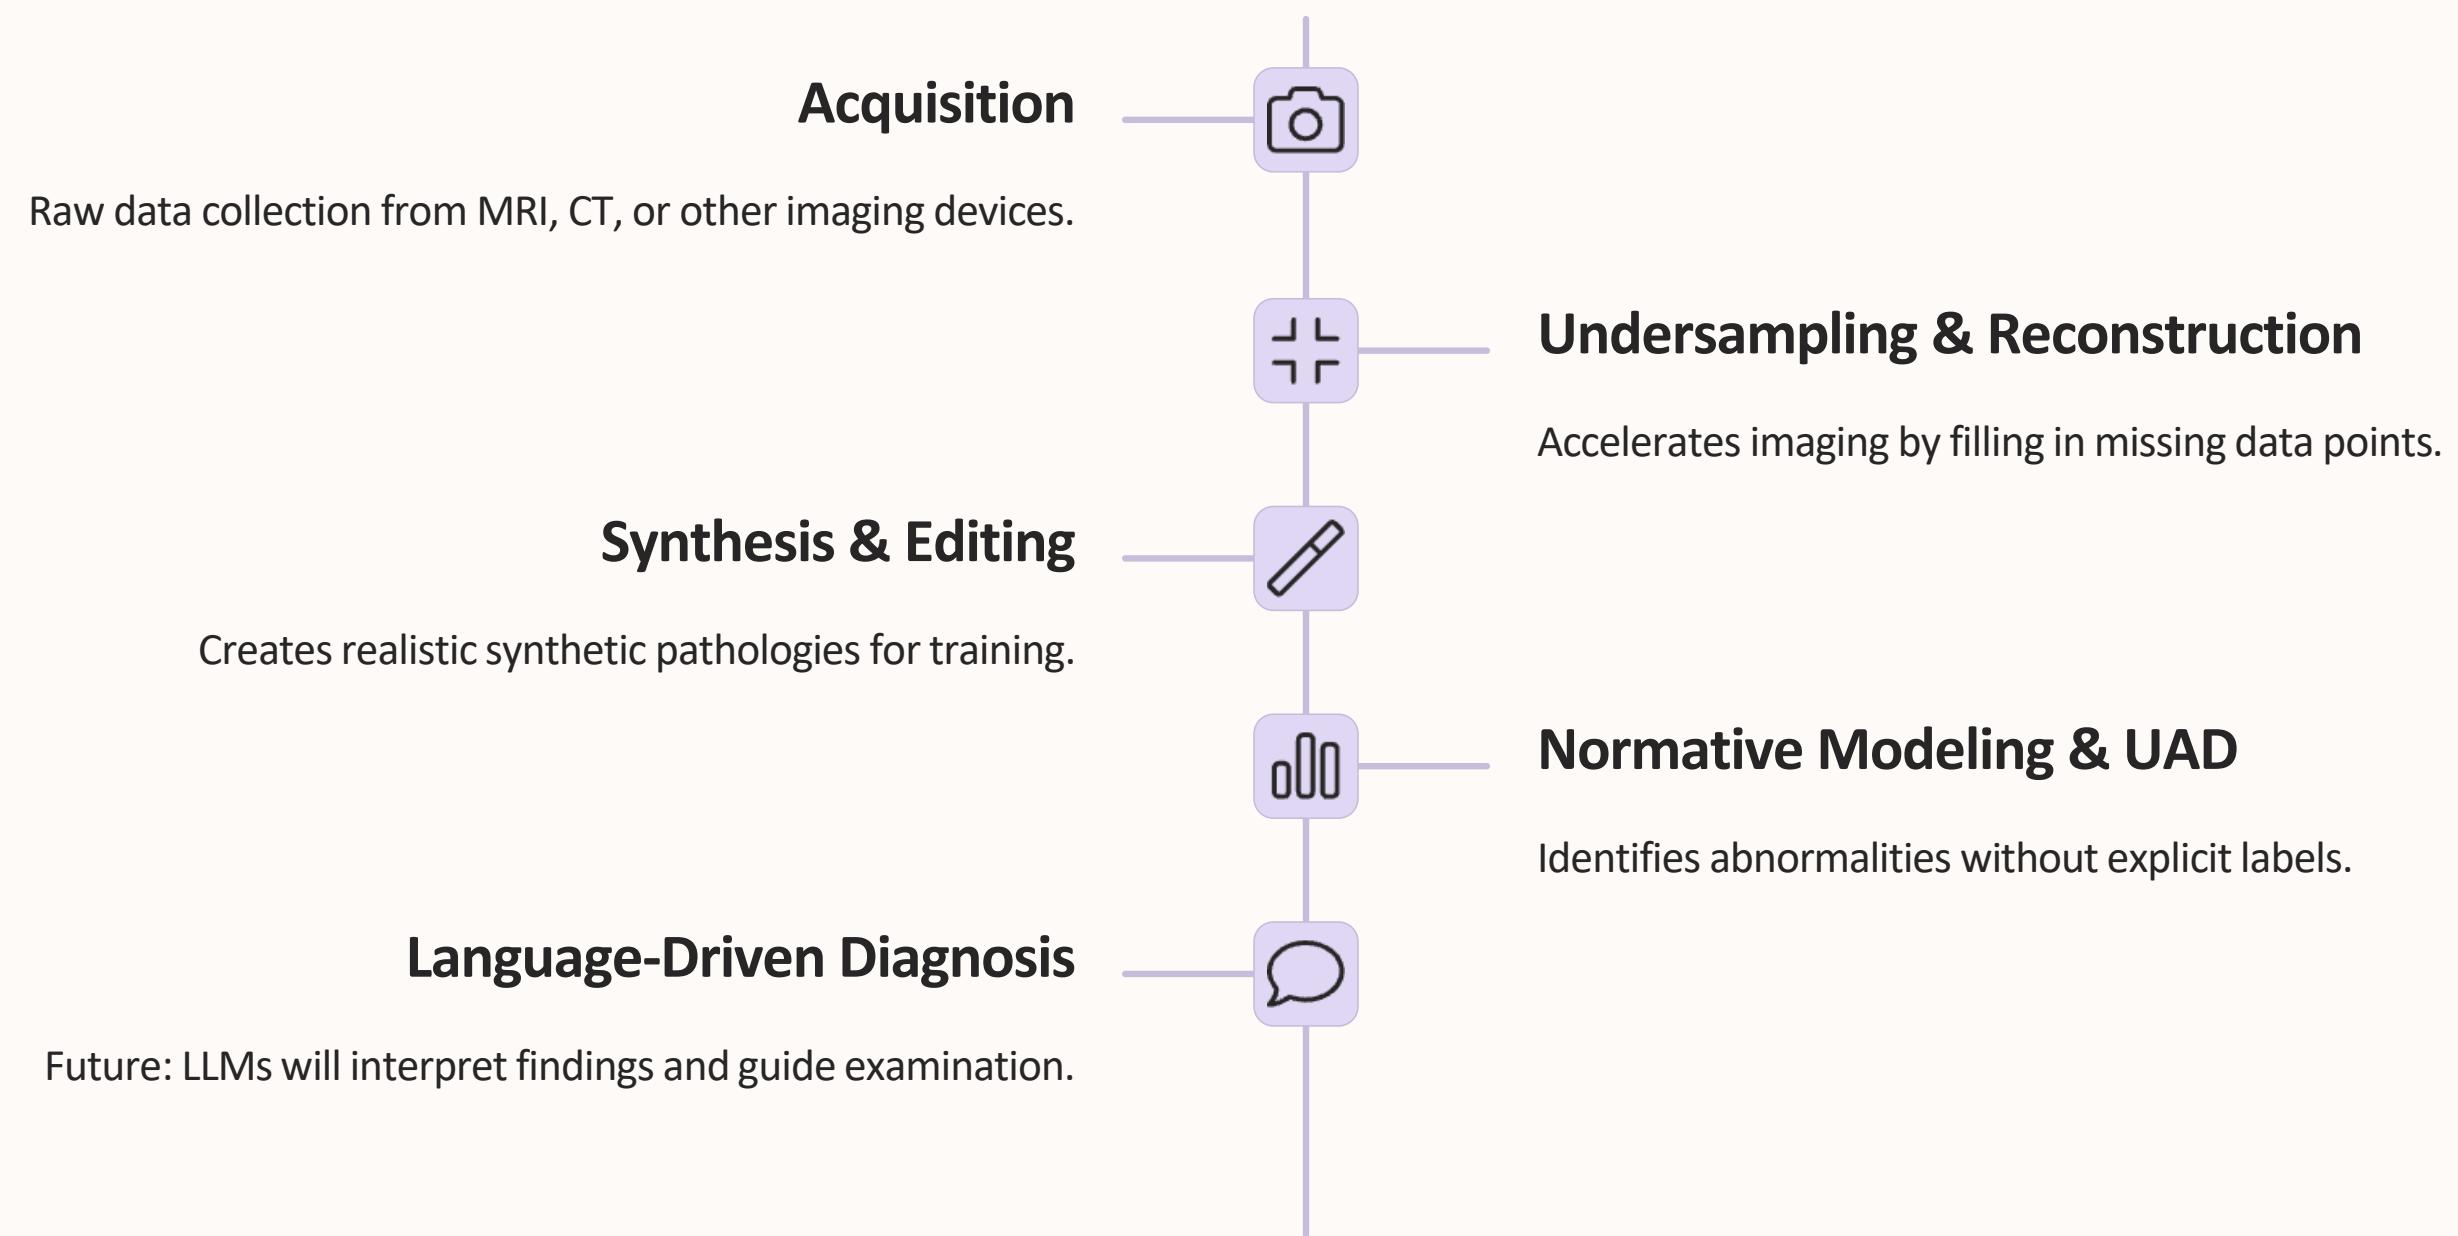
Diagnostic inference
Decision Making

Results communication
Actioning
Treatment planning

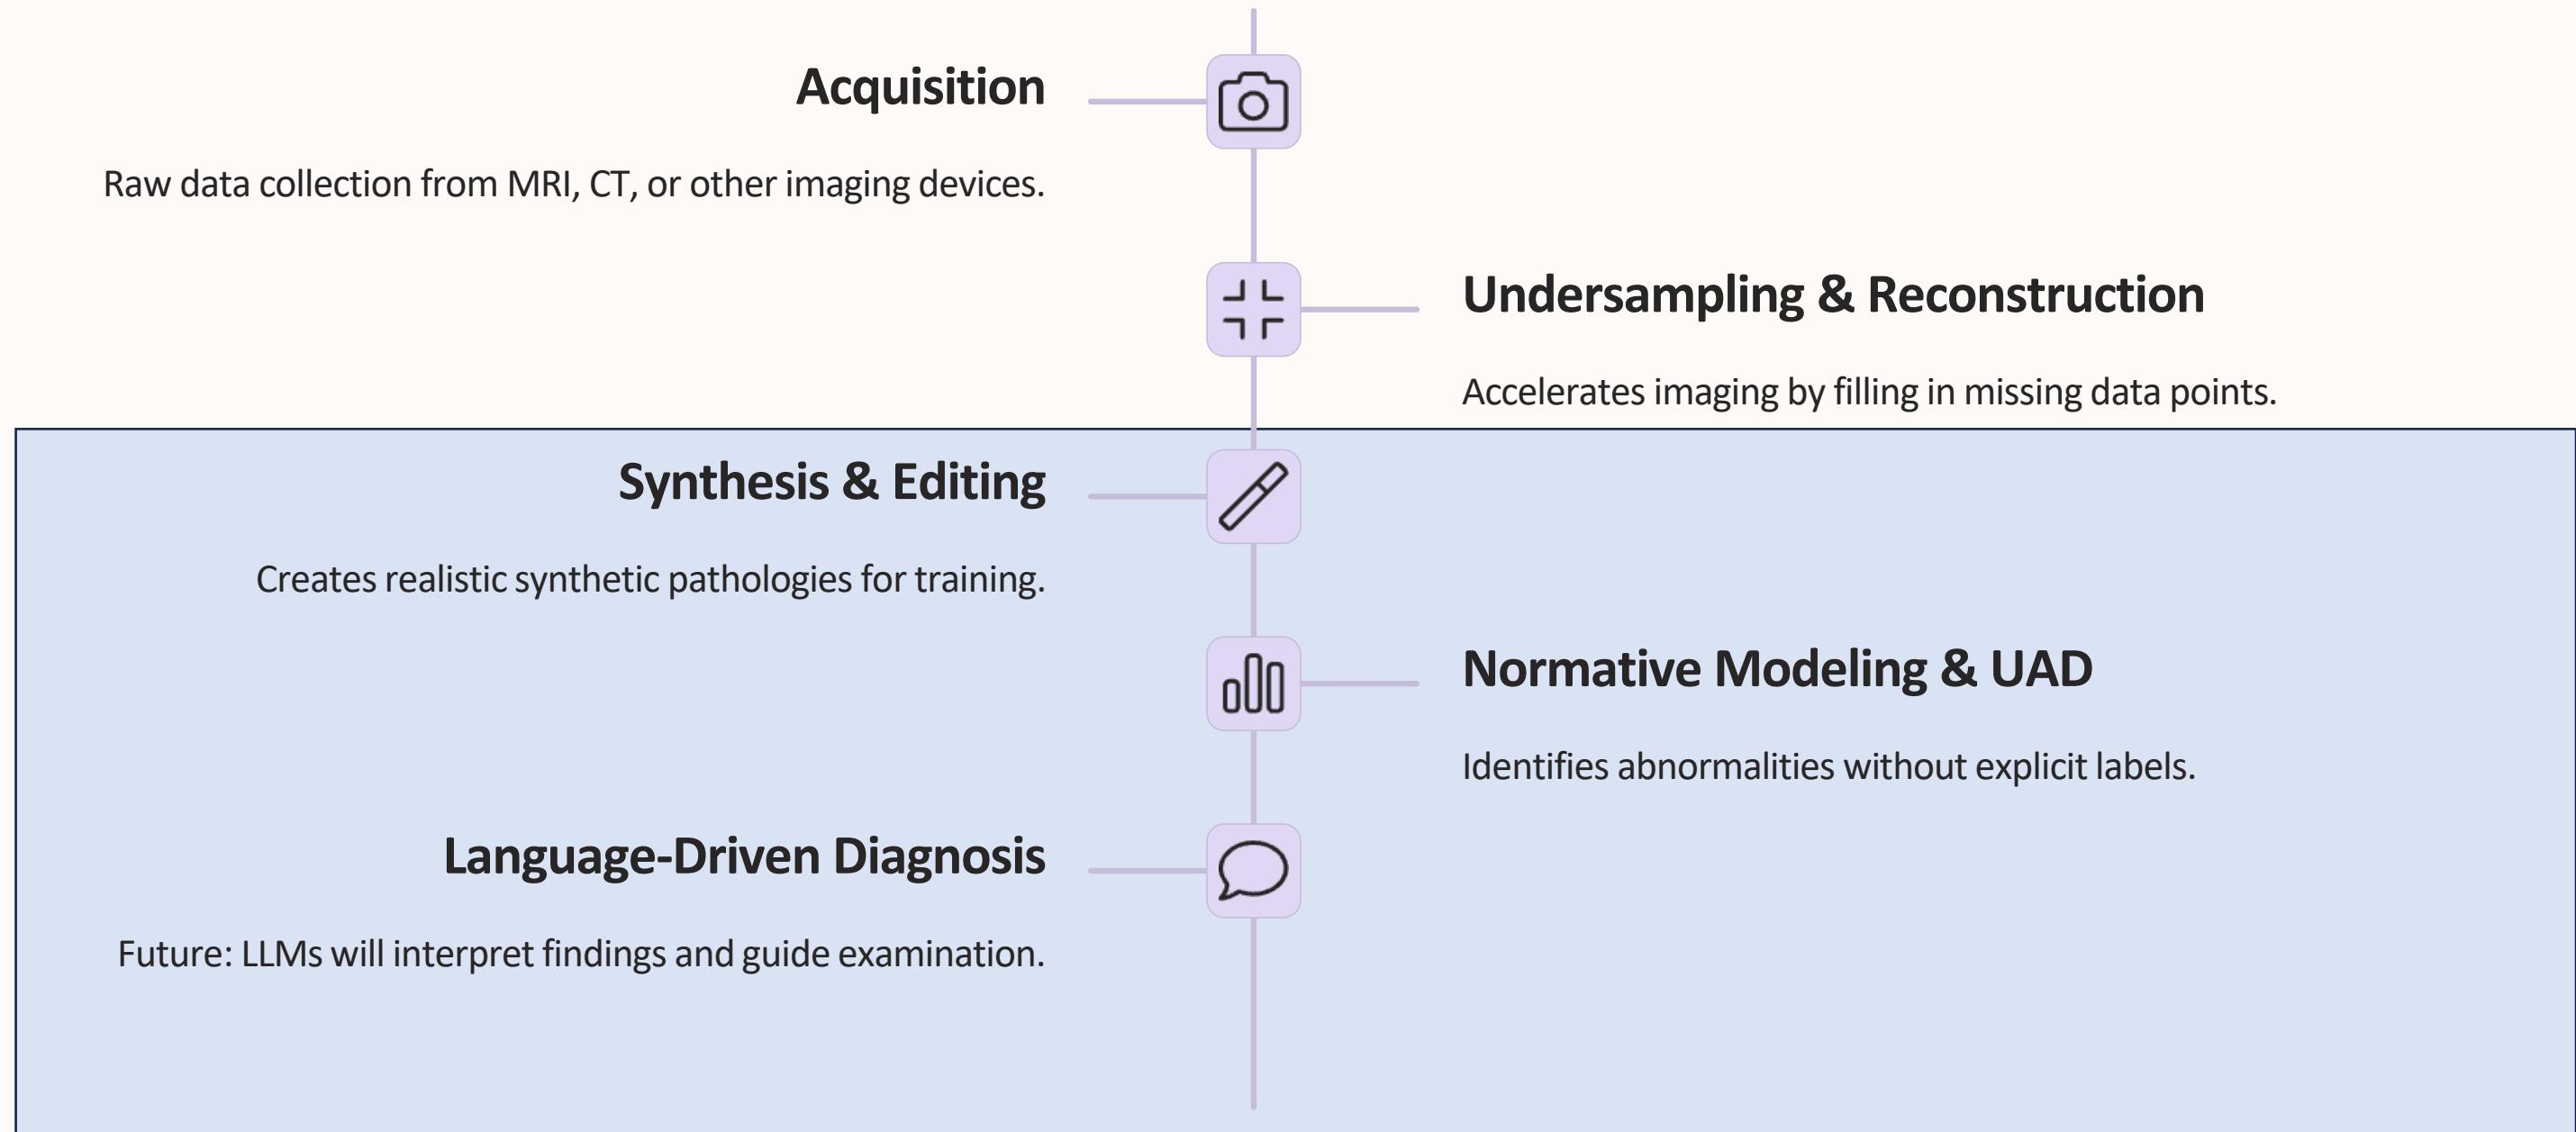


AI can assist radiologists!

Medical Imaging Pipeline

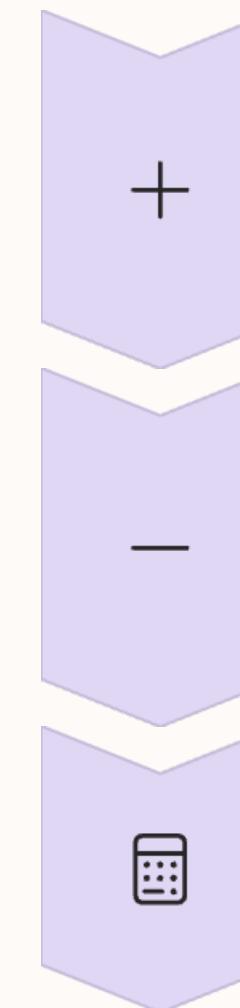


Medical Imaging Pipeline





What Are Diffusion Models?



Forward Process

Gradually add Gaussian noise to data until it becomes pure noise.

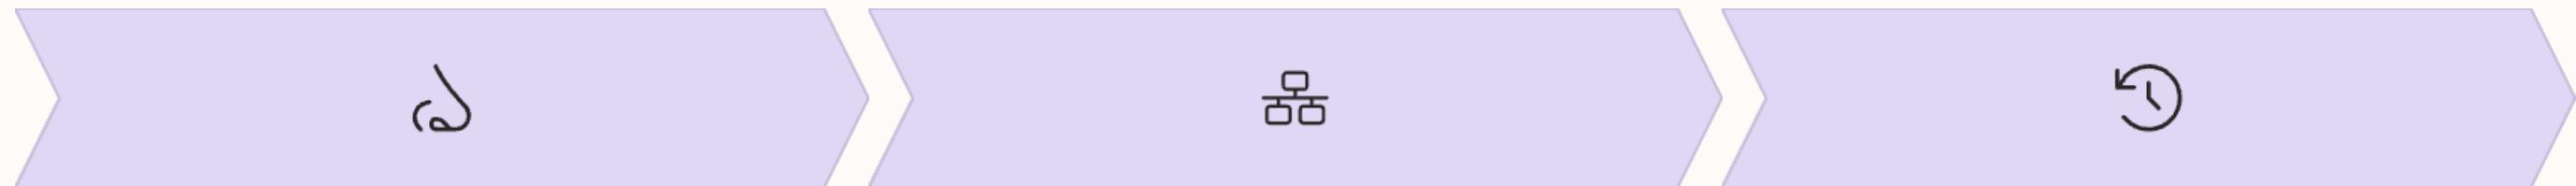
Reverse Process

Learn to denoise step-by-step, reconstructing the original data.

Score Function

Estimate gradient of data distribution log probability.

Diffusion Models



Forward Process

$$q(x_t|x_{\{t-1\}}) = N(x_t; \sqrt{(1-\beta_t)x_{\{t-1\}}, \beta_t I})$$

Gradually add noise to medical images until they become pure Gaussian noise.

Neural Network

$$\varepsilon_{\theta}(x_t, t) \approx \text{noise}$$

Model learns to predict the noise added at each timestep.

Reverse Process

$$p_{\theta}(x_{\{t-1\}}|x_t) = N(x_{\{t-1\}}; \mu_{\theta}(x_t, t), \Sigma_{\theta}(x_t, t))$$

Generate synthetic medical data by reversing noise addition step-by-step.

Synthetic data generation solves the challenge of rare disease underrepresentation and privacy constraints in medical datasets.

Synthetic Data – Why It Matters

Privacy-Preserving

Generated data shares no identifiable information from real patients.

Enables broader sharing and collaboration across institutions.

Rare Disease Representation

Create unlimited examples of conditions with limited real samples.

Balance datasets that would otherwise be heavily skewed.

Data Augmentation

Improve model robustness through diverse synthetic examples.

Test against variations not present in original data.



Conditional Generation

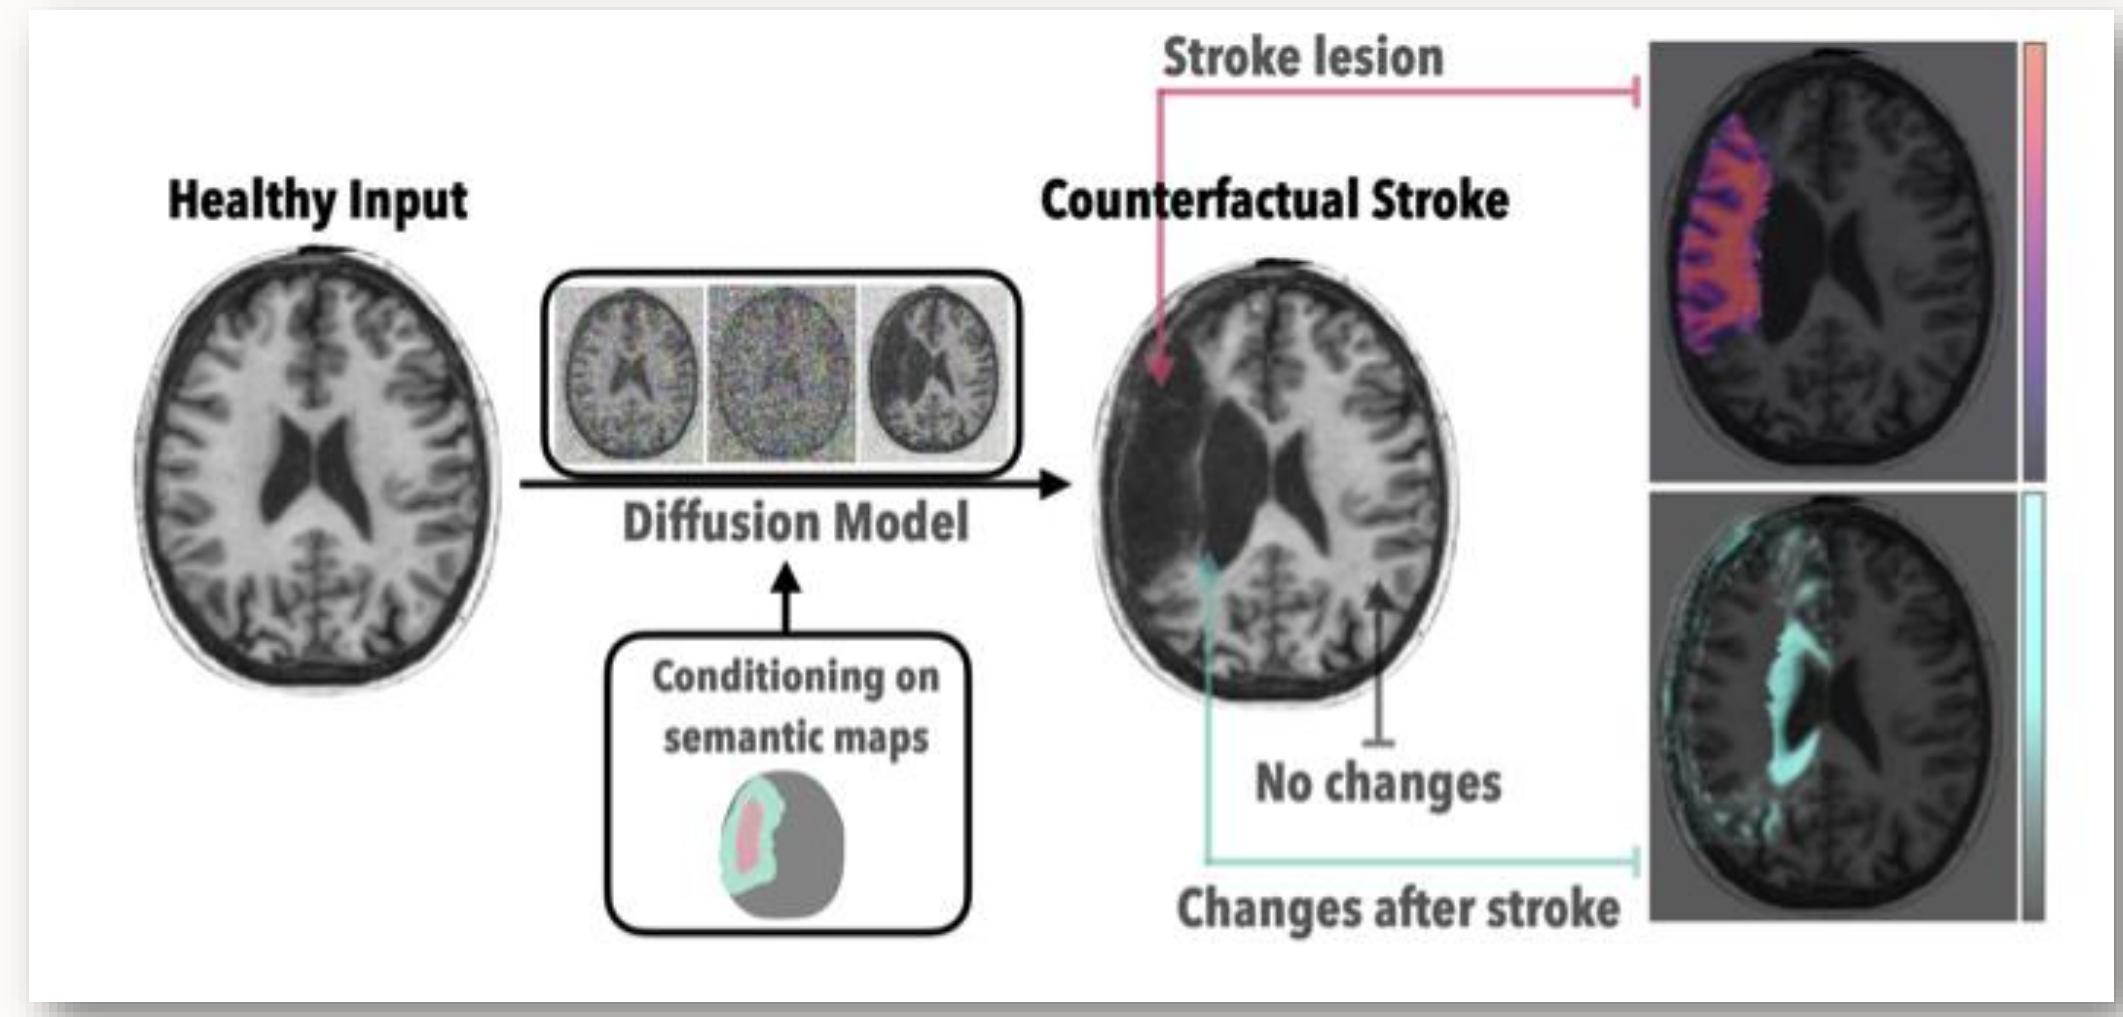


“An image of the pope”

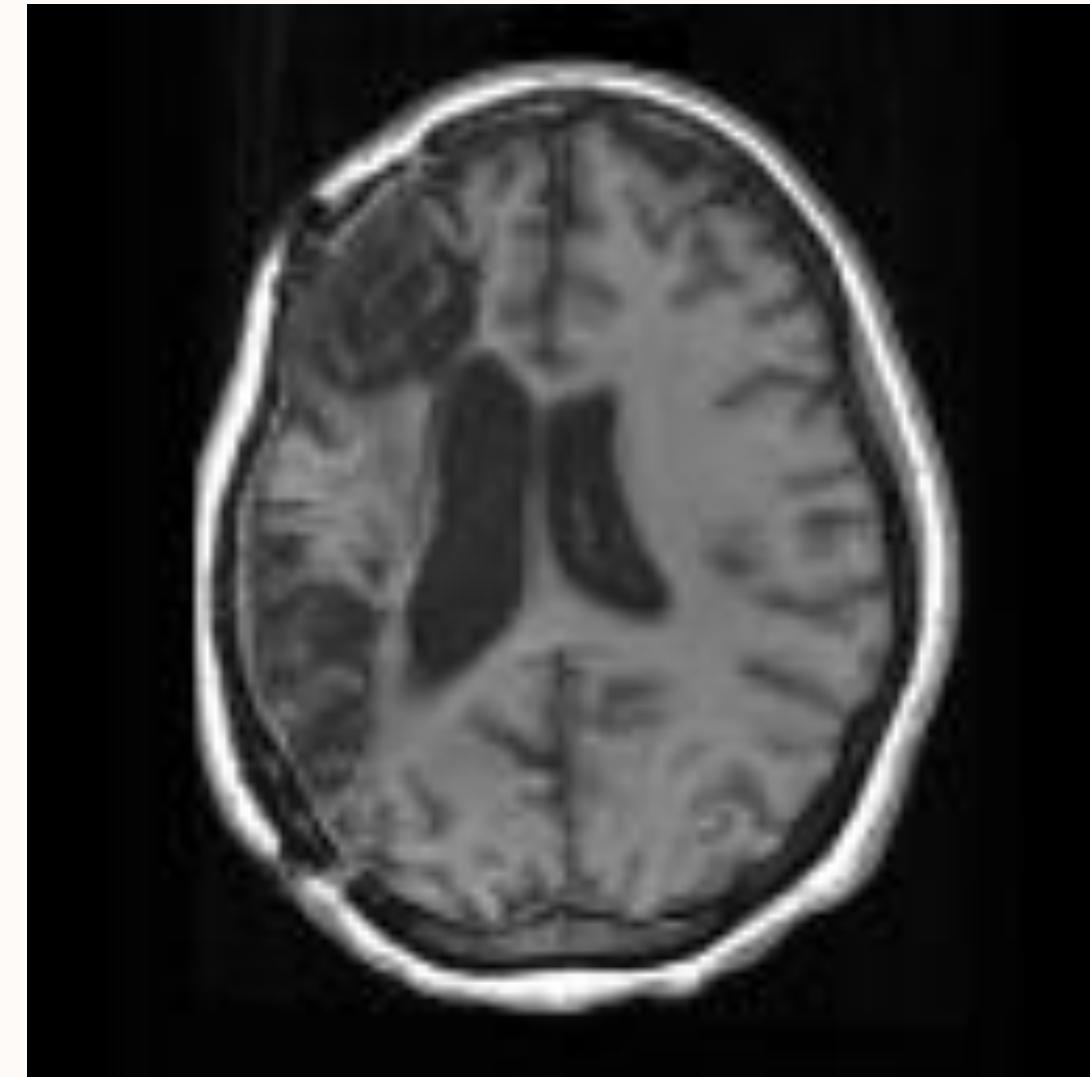
MedEDIT



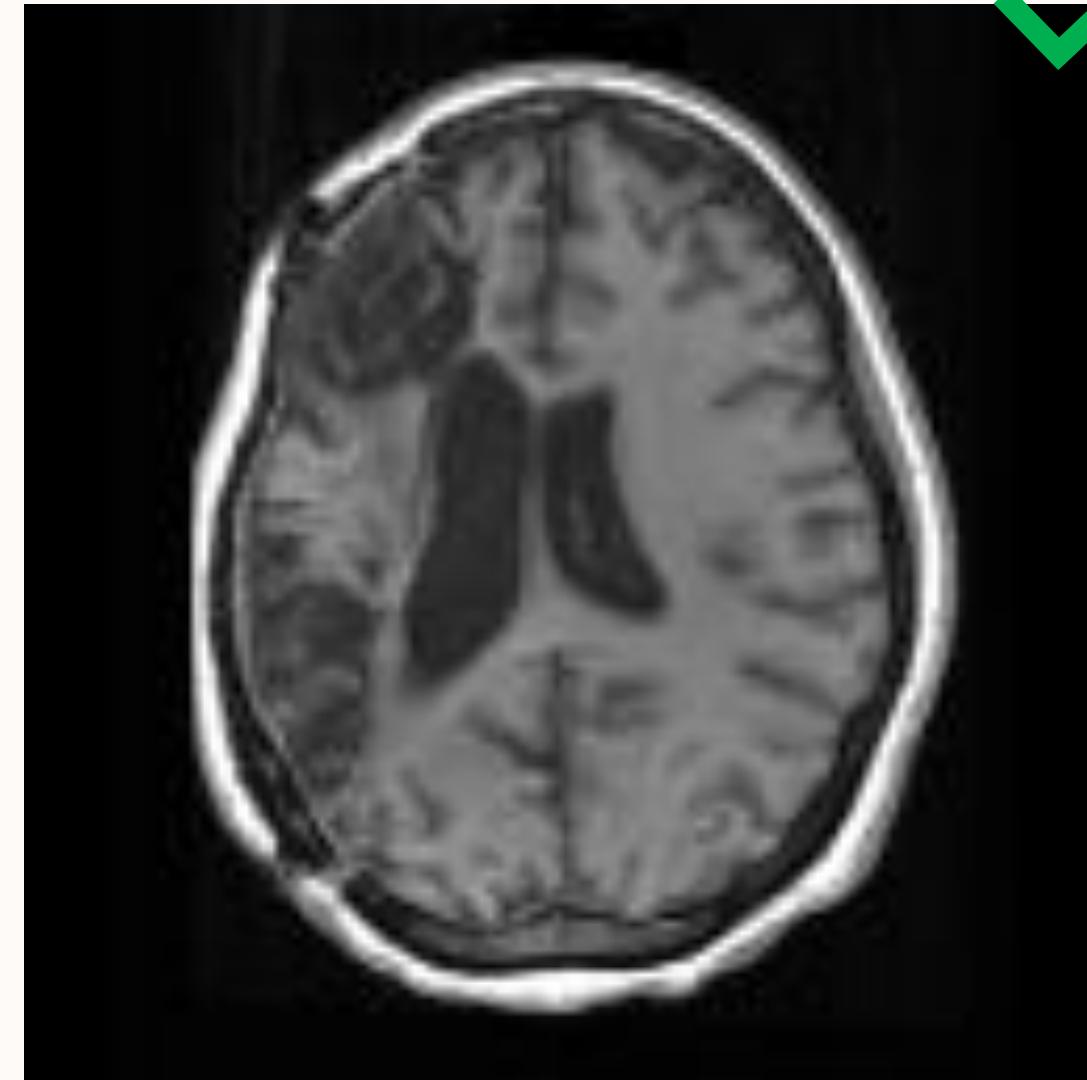
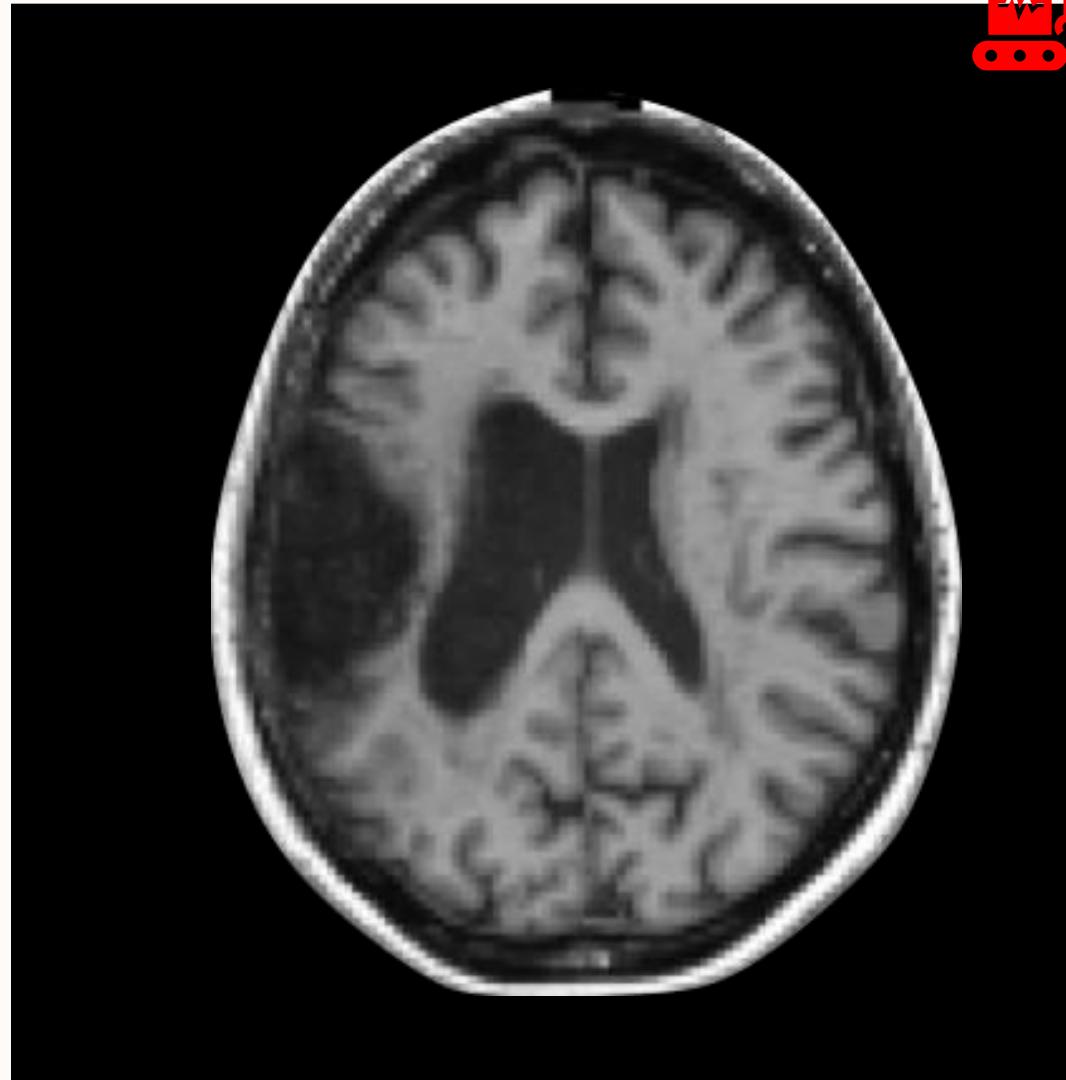
Malek Ben Alaya



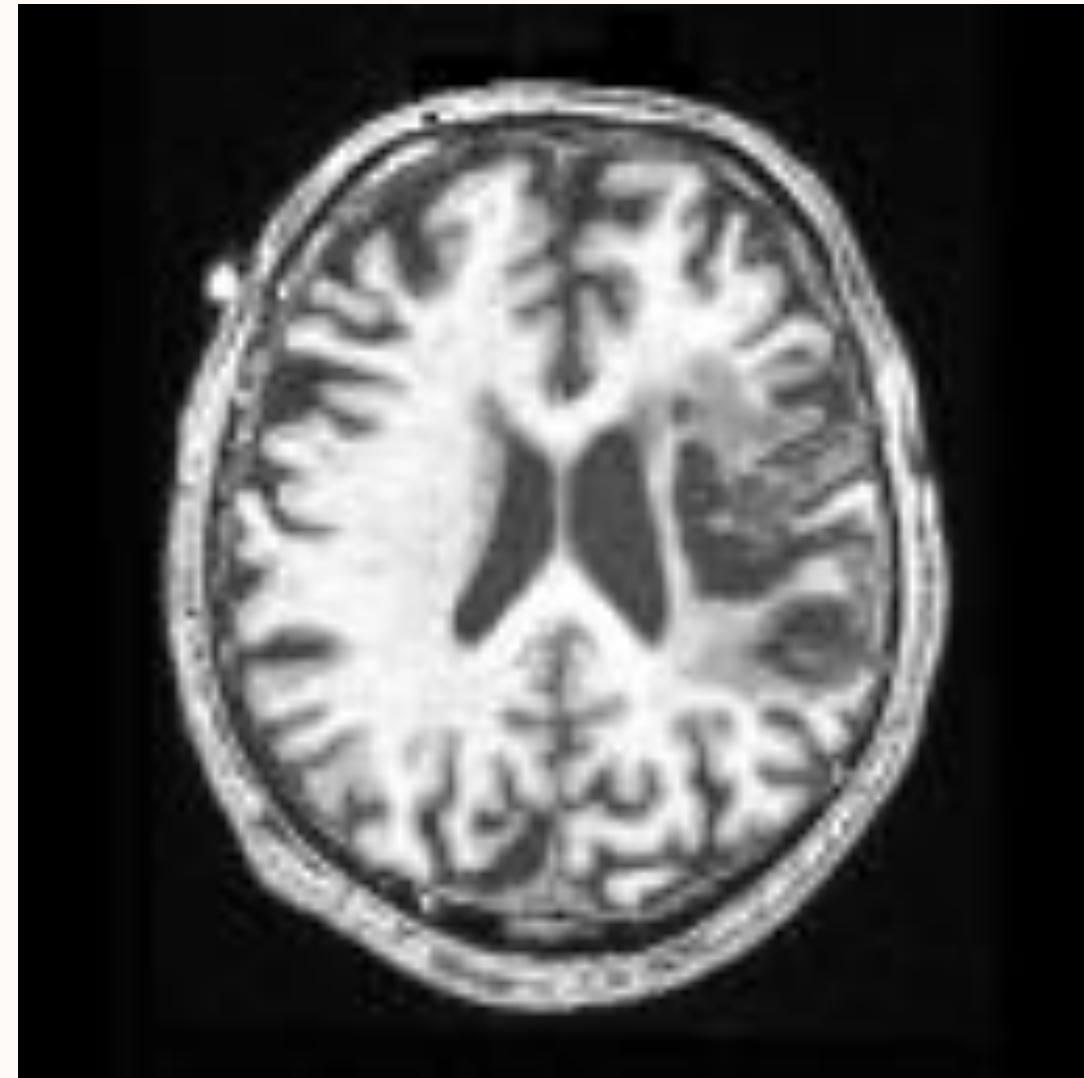
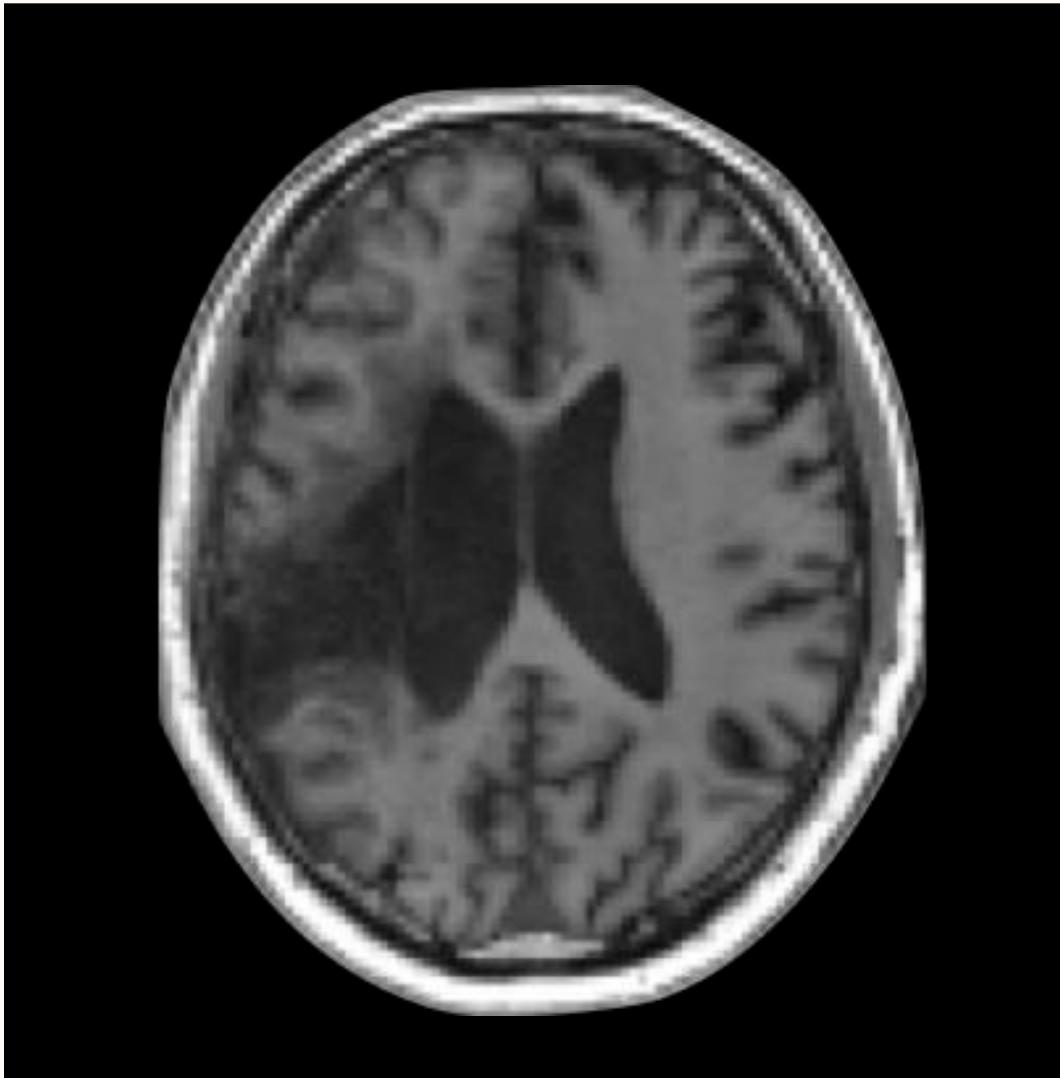
Real or AI?



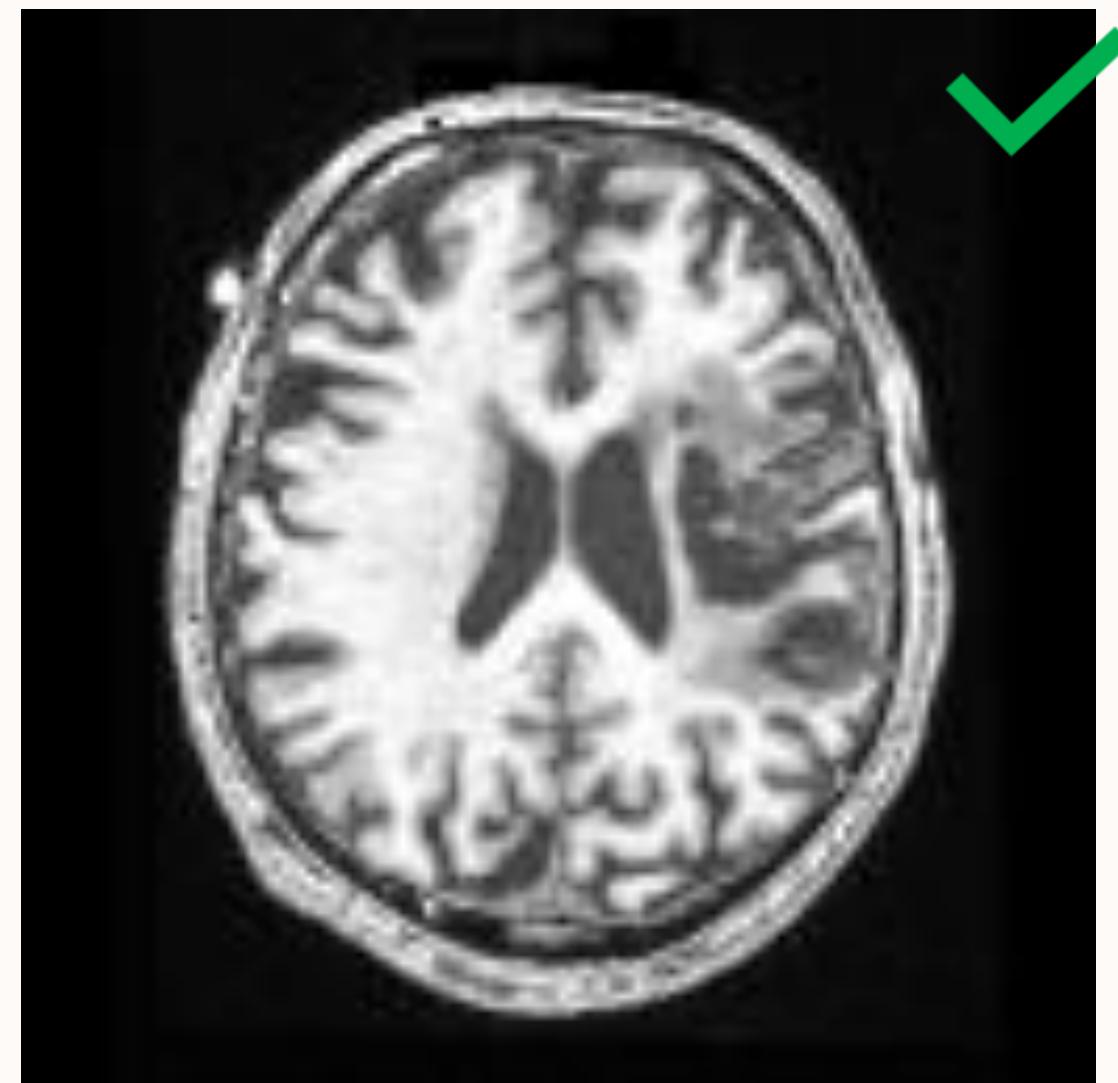
Real or AI?



Real or AI?



Real or AI?



Language-based Image Editing



Language-based Image Editing



“An image of the pope”

Language-based Image Editing



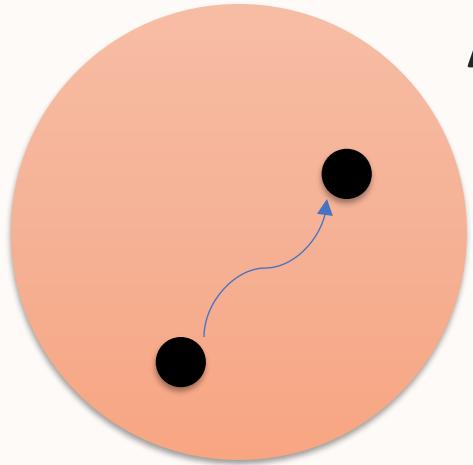
“An image of the
pope... in traditional
portuguese clothes”

Language-based Image Editing



Karim ElGhandour

Language Space

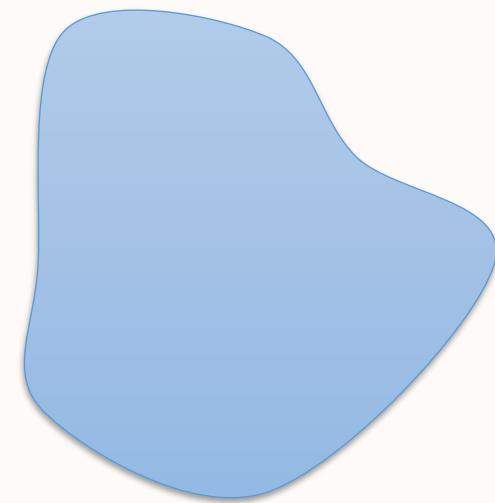


A healthy brain MRI

A brain MRI with a tumor

Diffusion

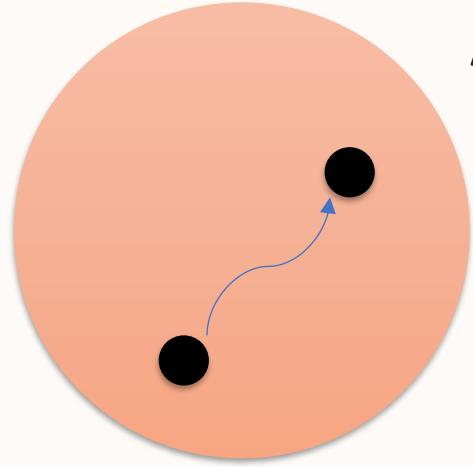
Image Space





Language-based Image Editing

Language Space

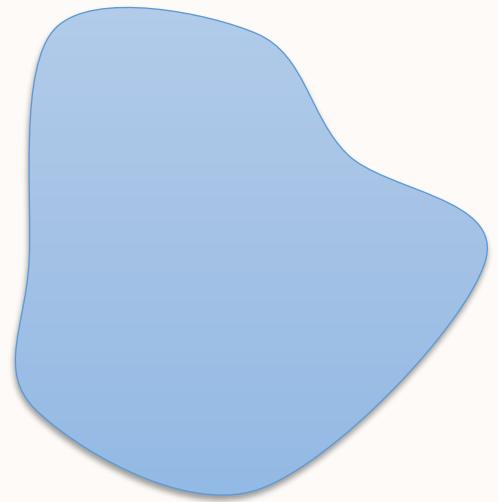


A healthy brain MRI

A brain MRI with a tumor

Diffusion

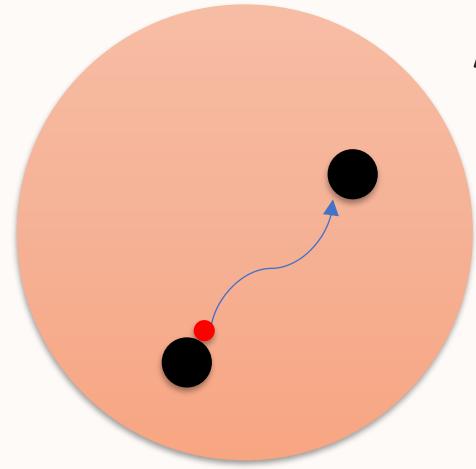
Image Space





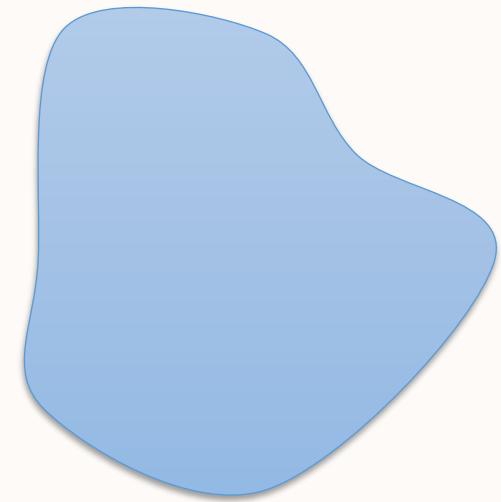
Language-based Image Editing

Language Space



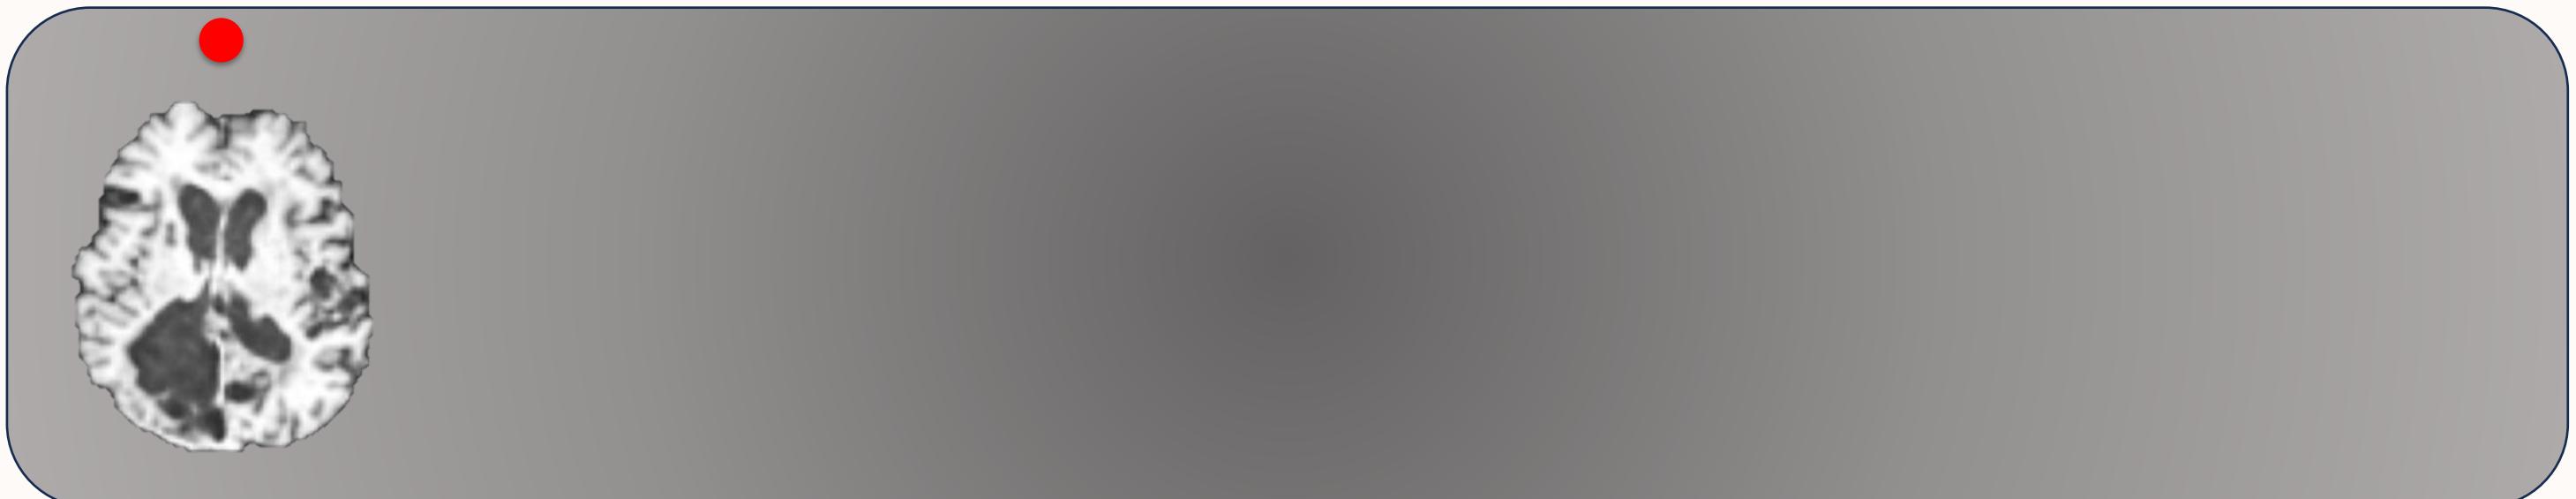
A healthy brain MRI

Image Space



A brain MRI with a tumor

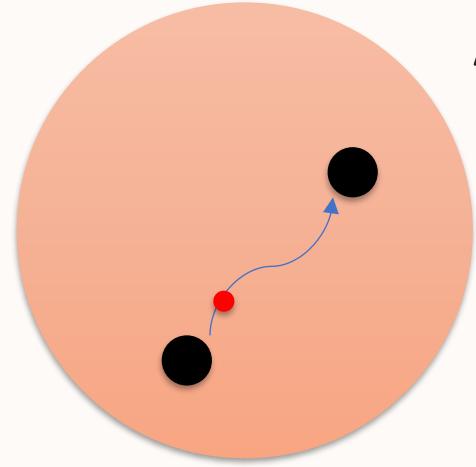
Diffusion





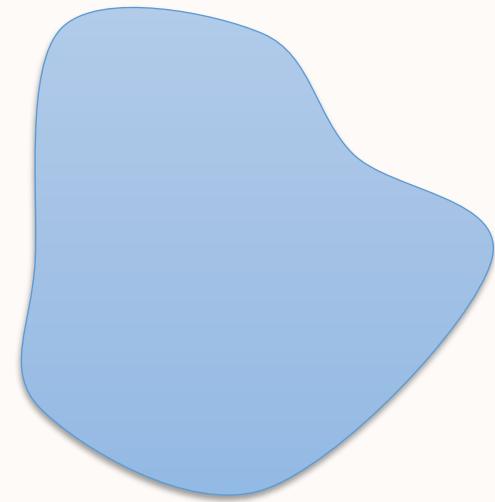
Language-based Image Editing

Language Space



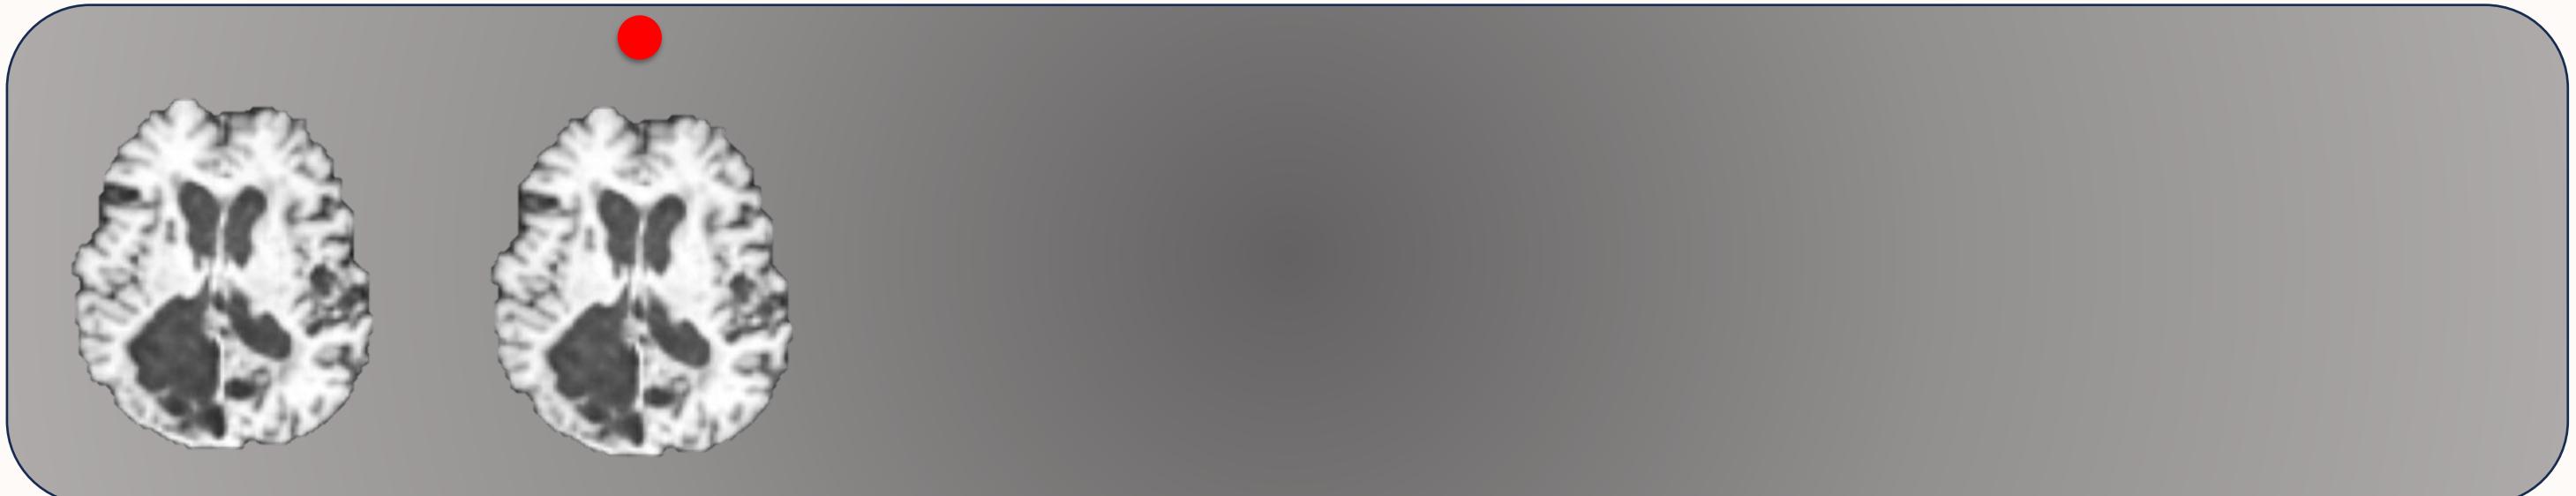
A healthy brain MRI

Image Space



Diffusion

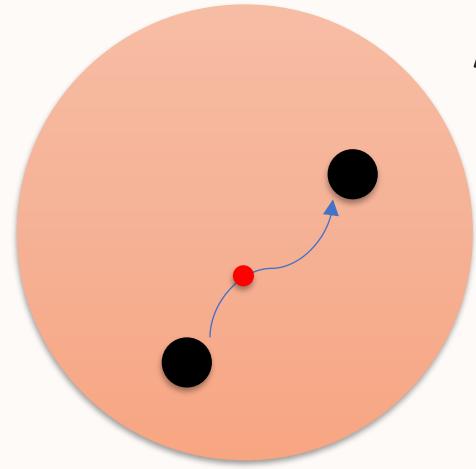
A brain MRI with a tumor





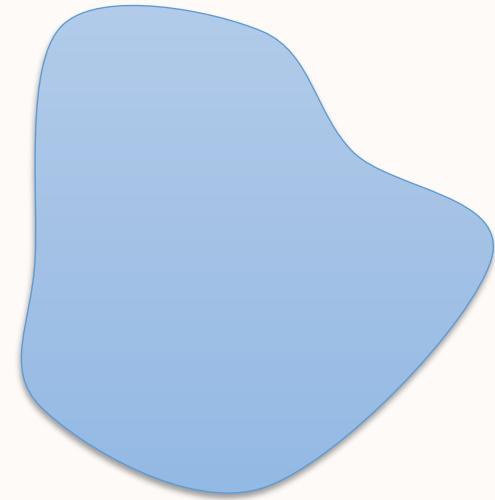
Language-based Image Editing

Language Space



A healthy brain MRI

Image Space



Diffusion

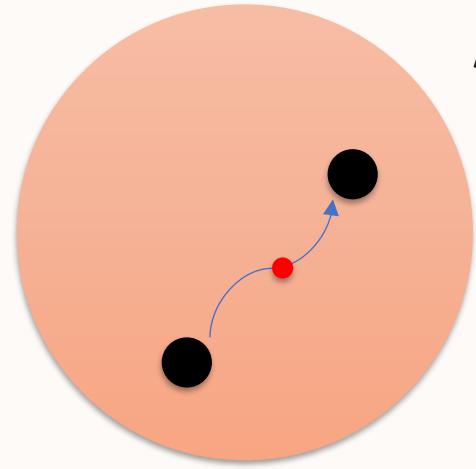
A brain MRI with a tumor





Language-based Image Editing

Language Space



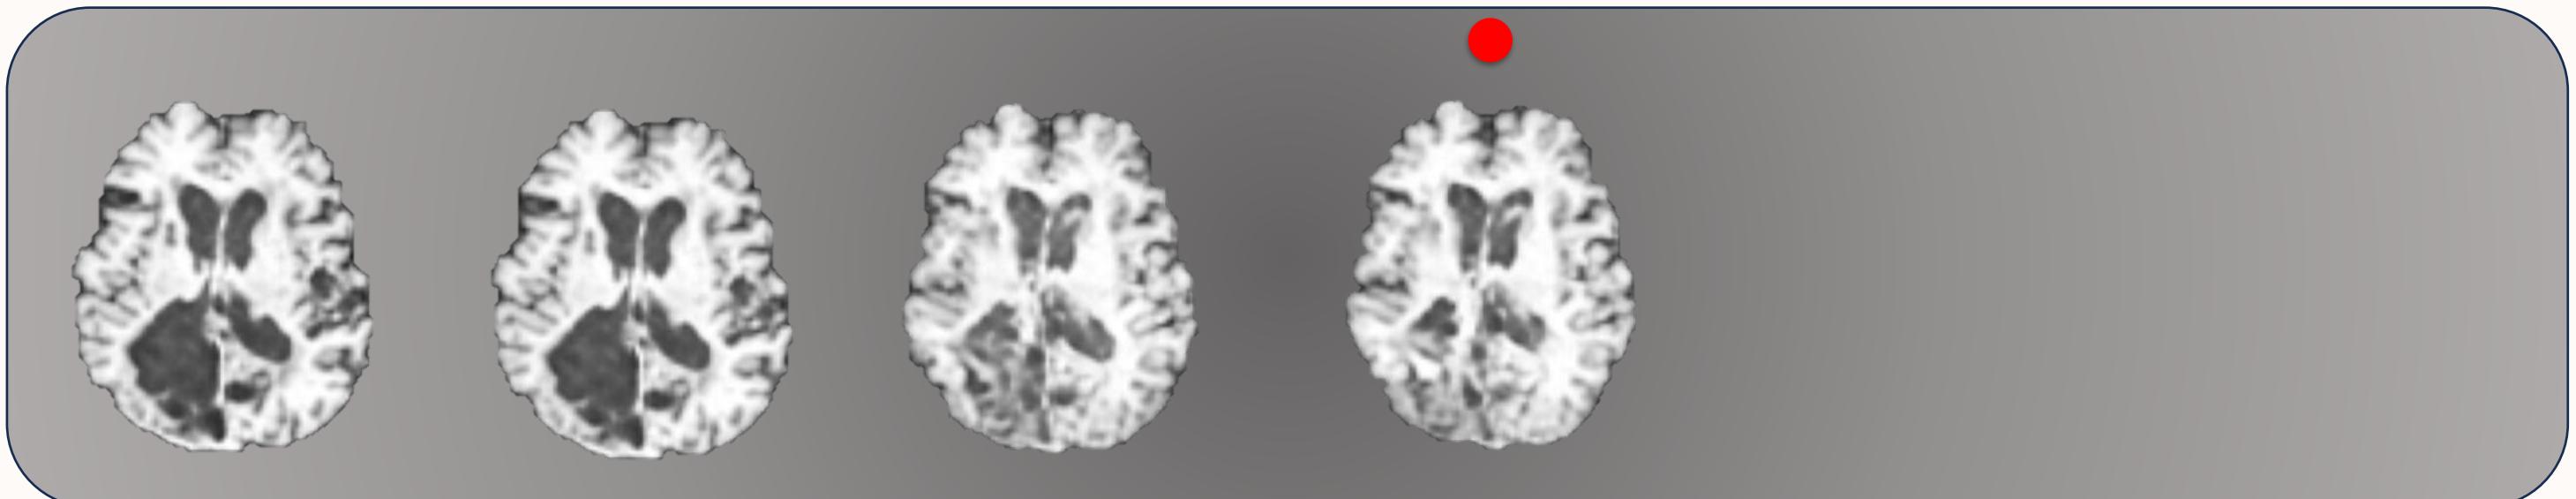
A healthy brain MRI

Image Space



Diffusion

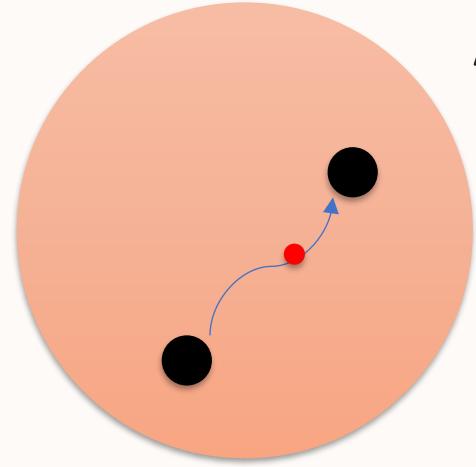
A brain MRI with a tumor





Language-based Image Editing

Language Space



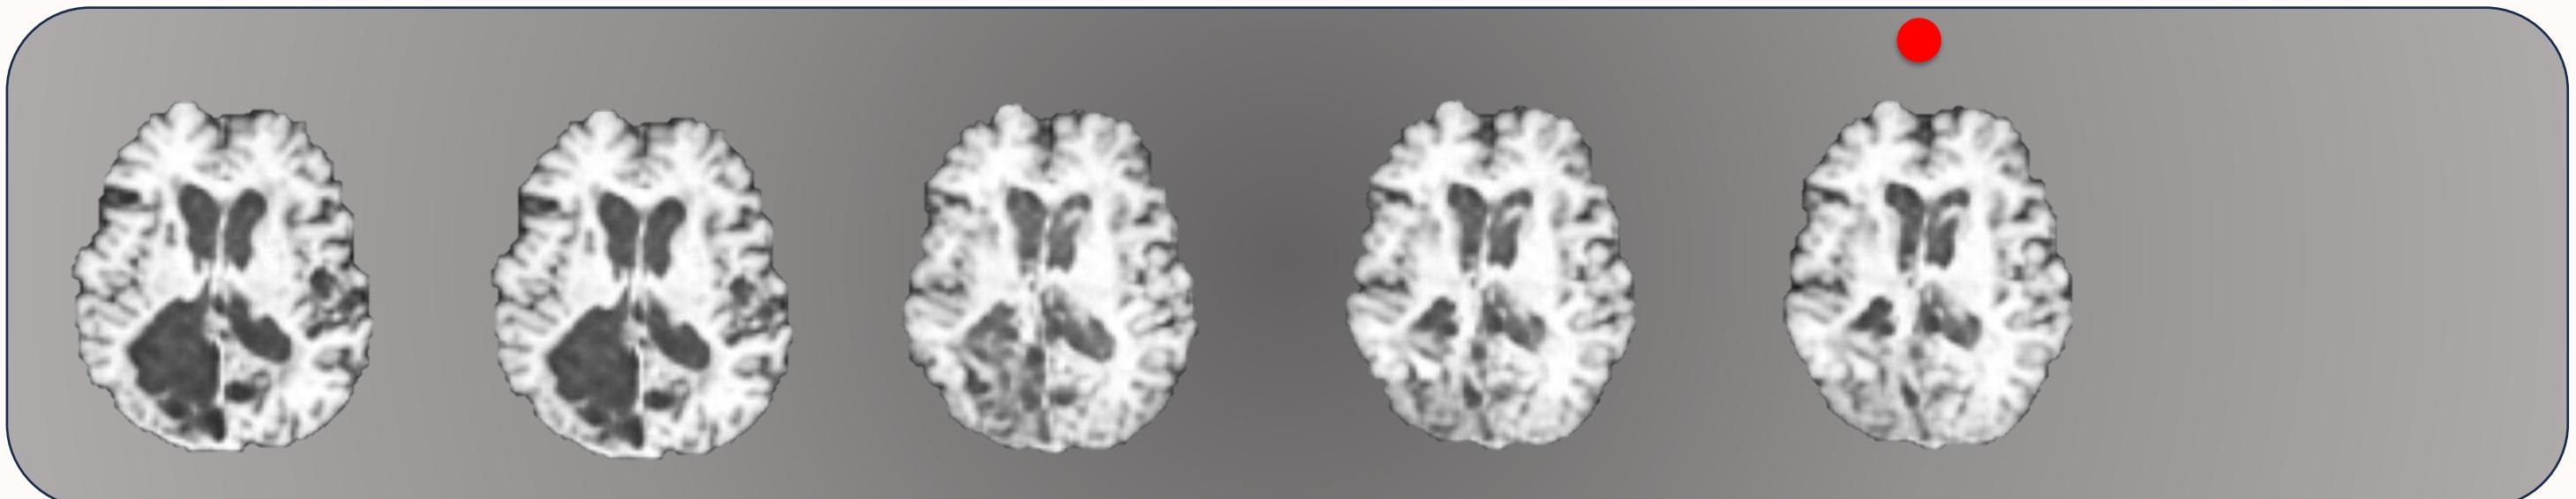
A healthy brain MRI

Image Space



A brain MRI with a tumor

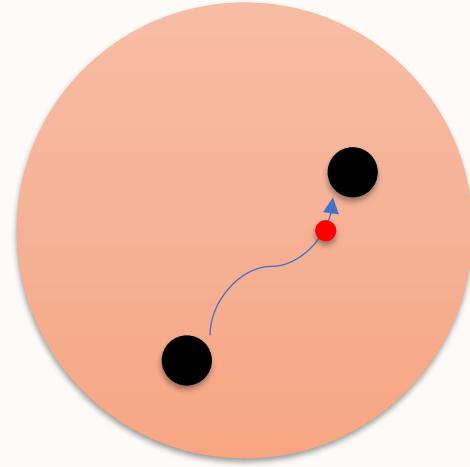
Diffusion





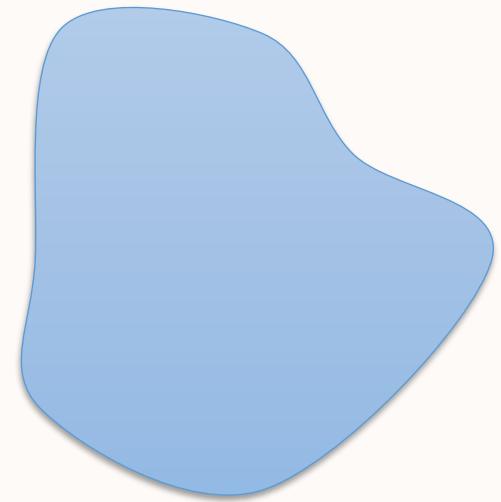
Language-based Image Editing

Language Space



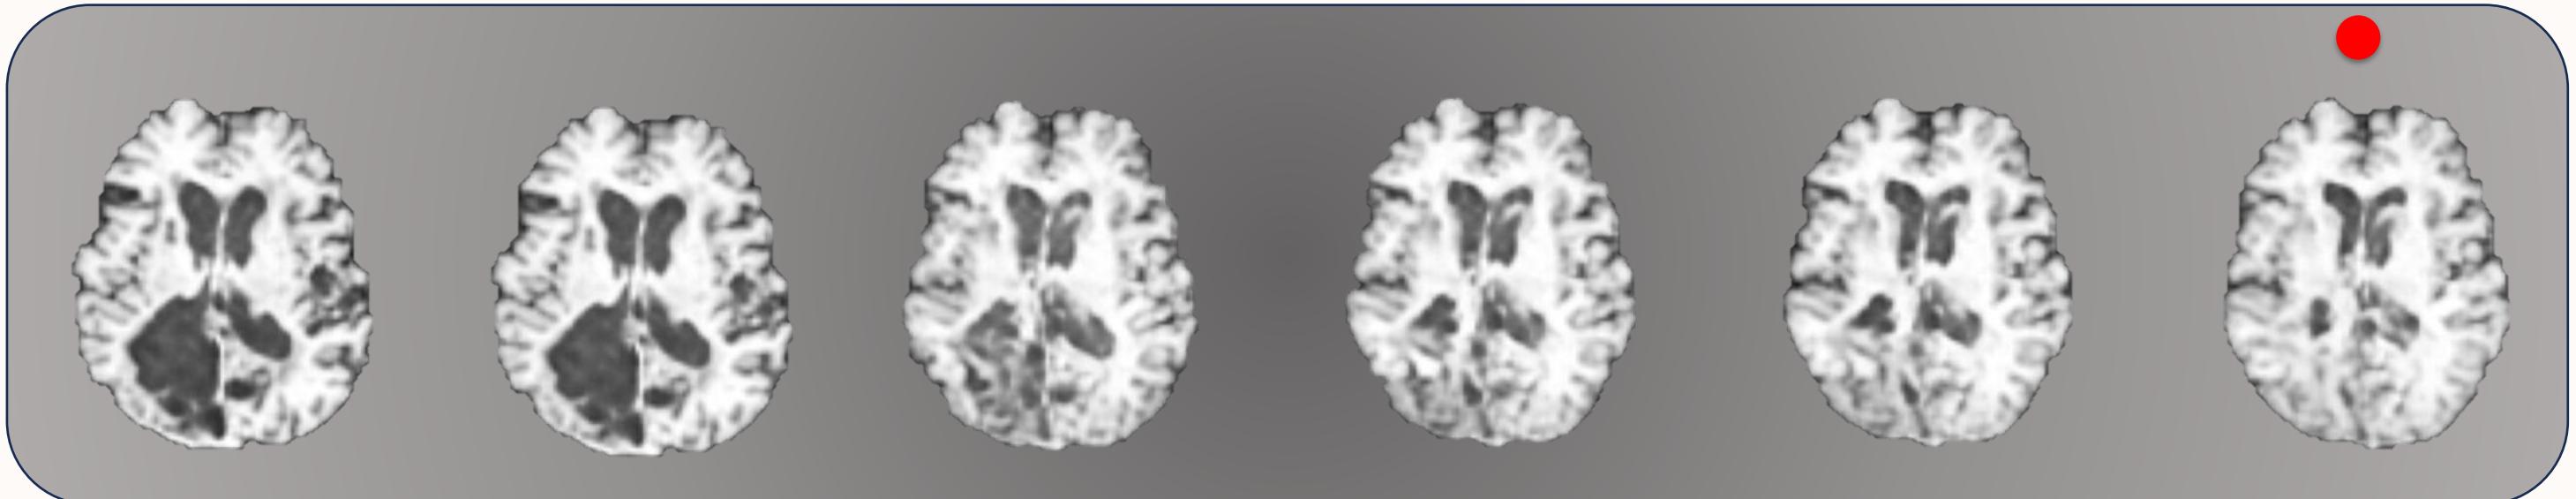
A healthy brain MRI

Image Space

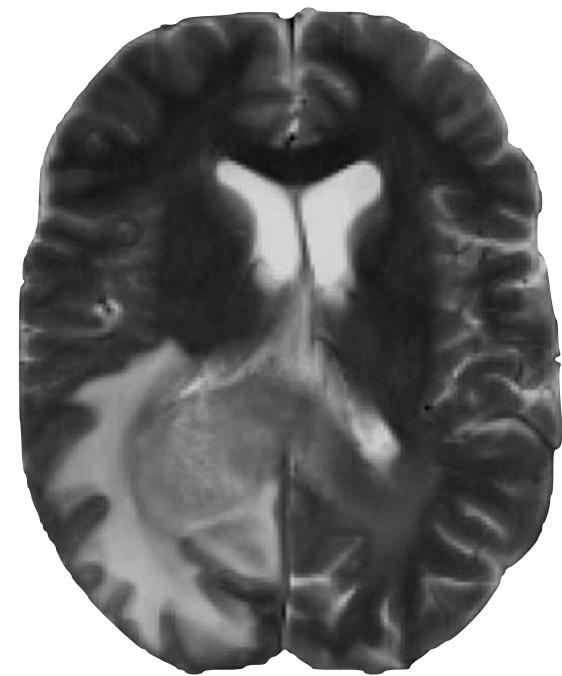


Diffusion

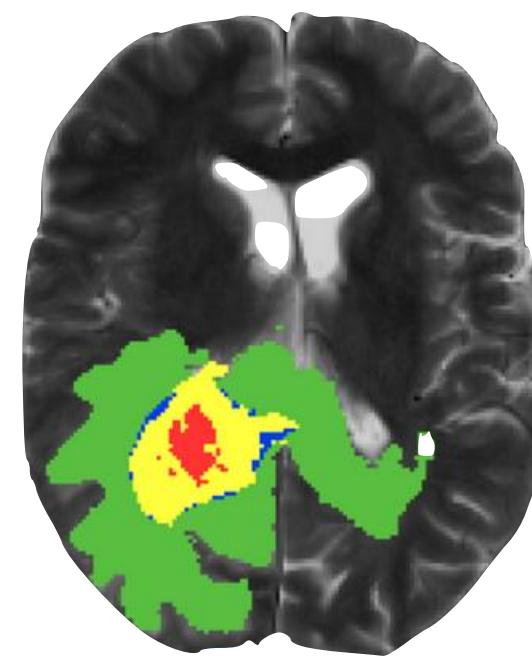
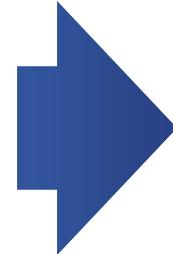
A brain MRI with a tumor



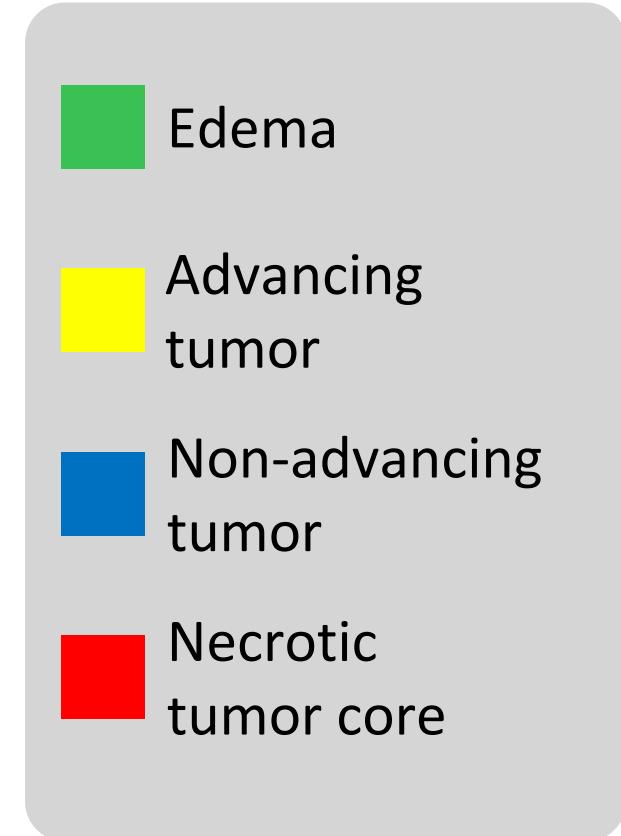
What for? Anomaly Detection (Supervised)



Input Brain MRI

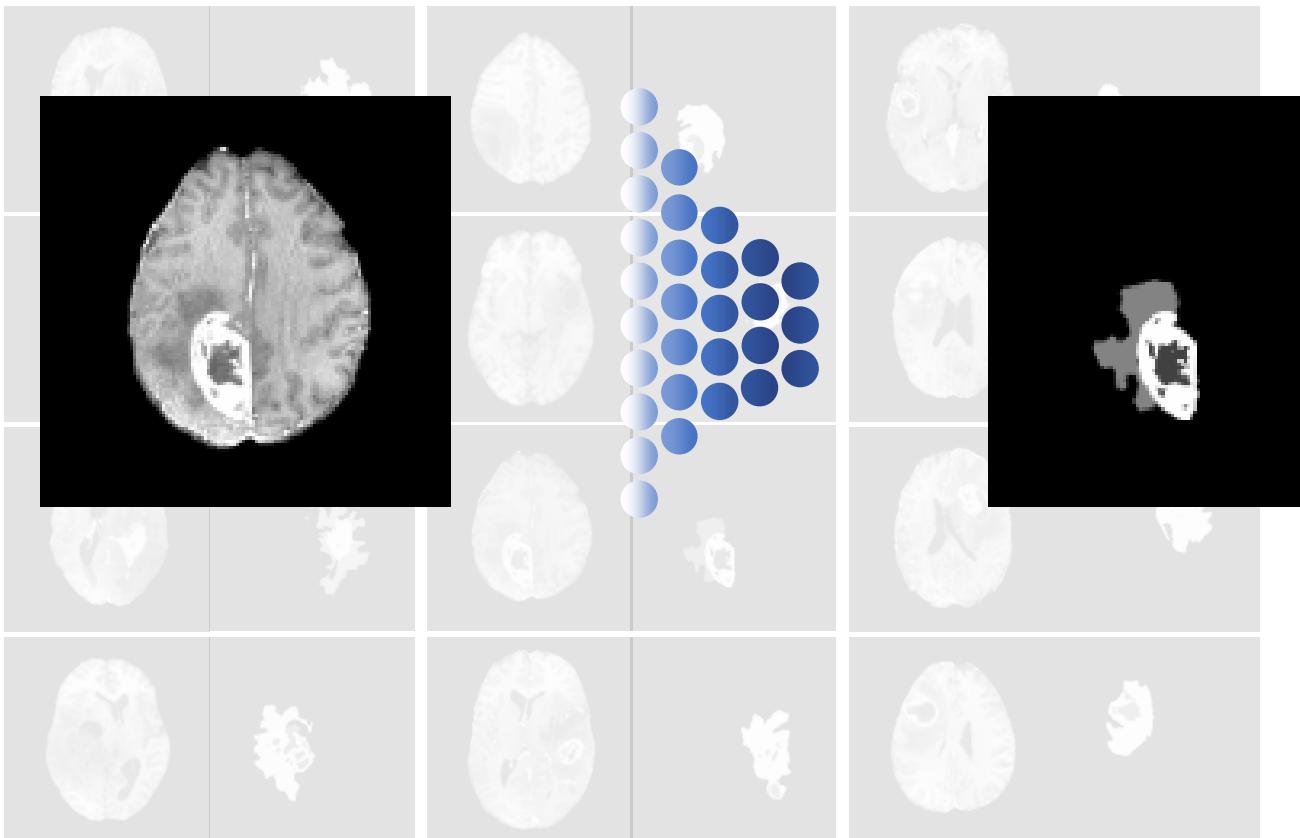


Automatic Segmentation



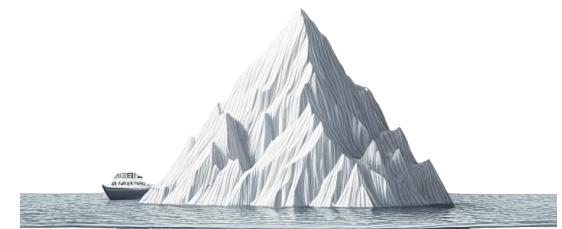
Train

Neural networks are trained to imitate radiologists
and learn image -> anomaly masks mappings



Annotated dataset of images + masks

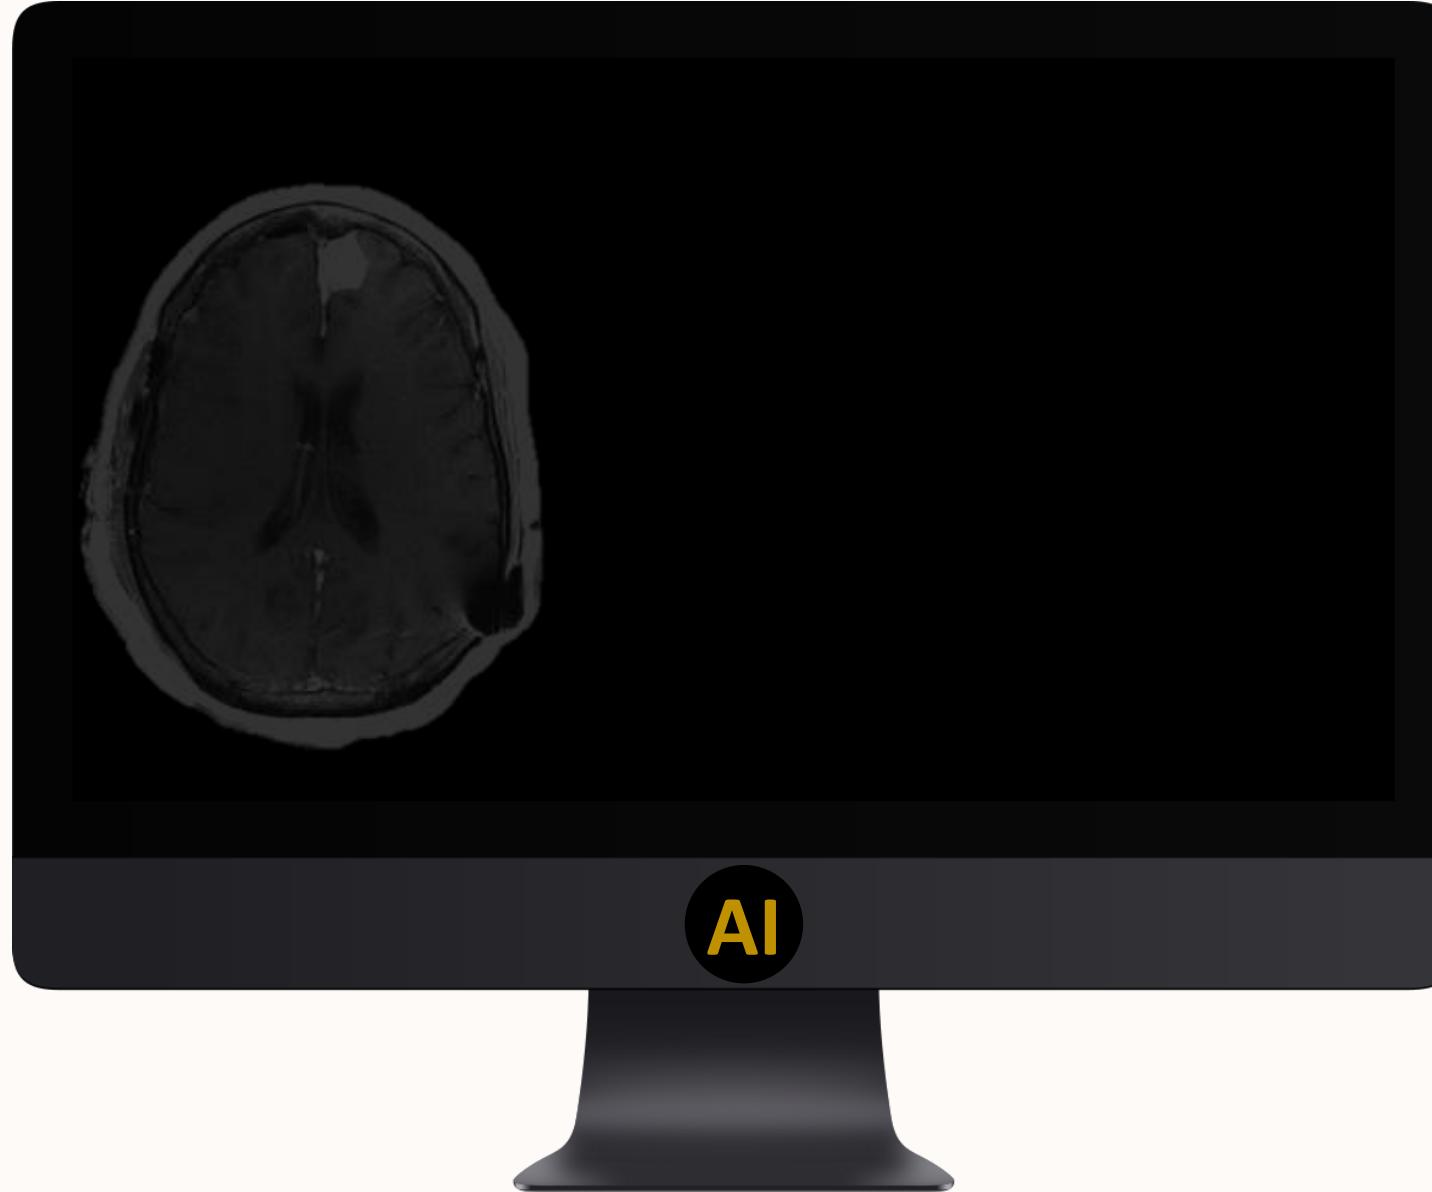
Known Diseases



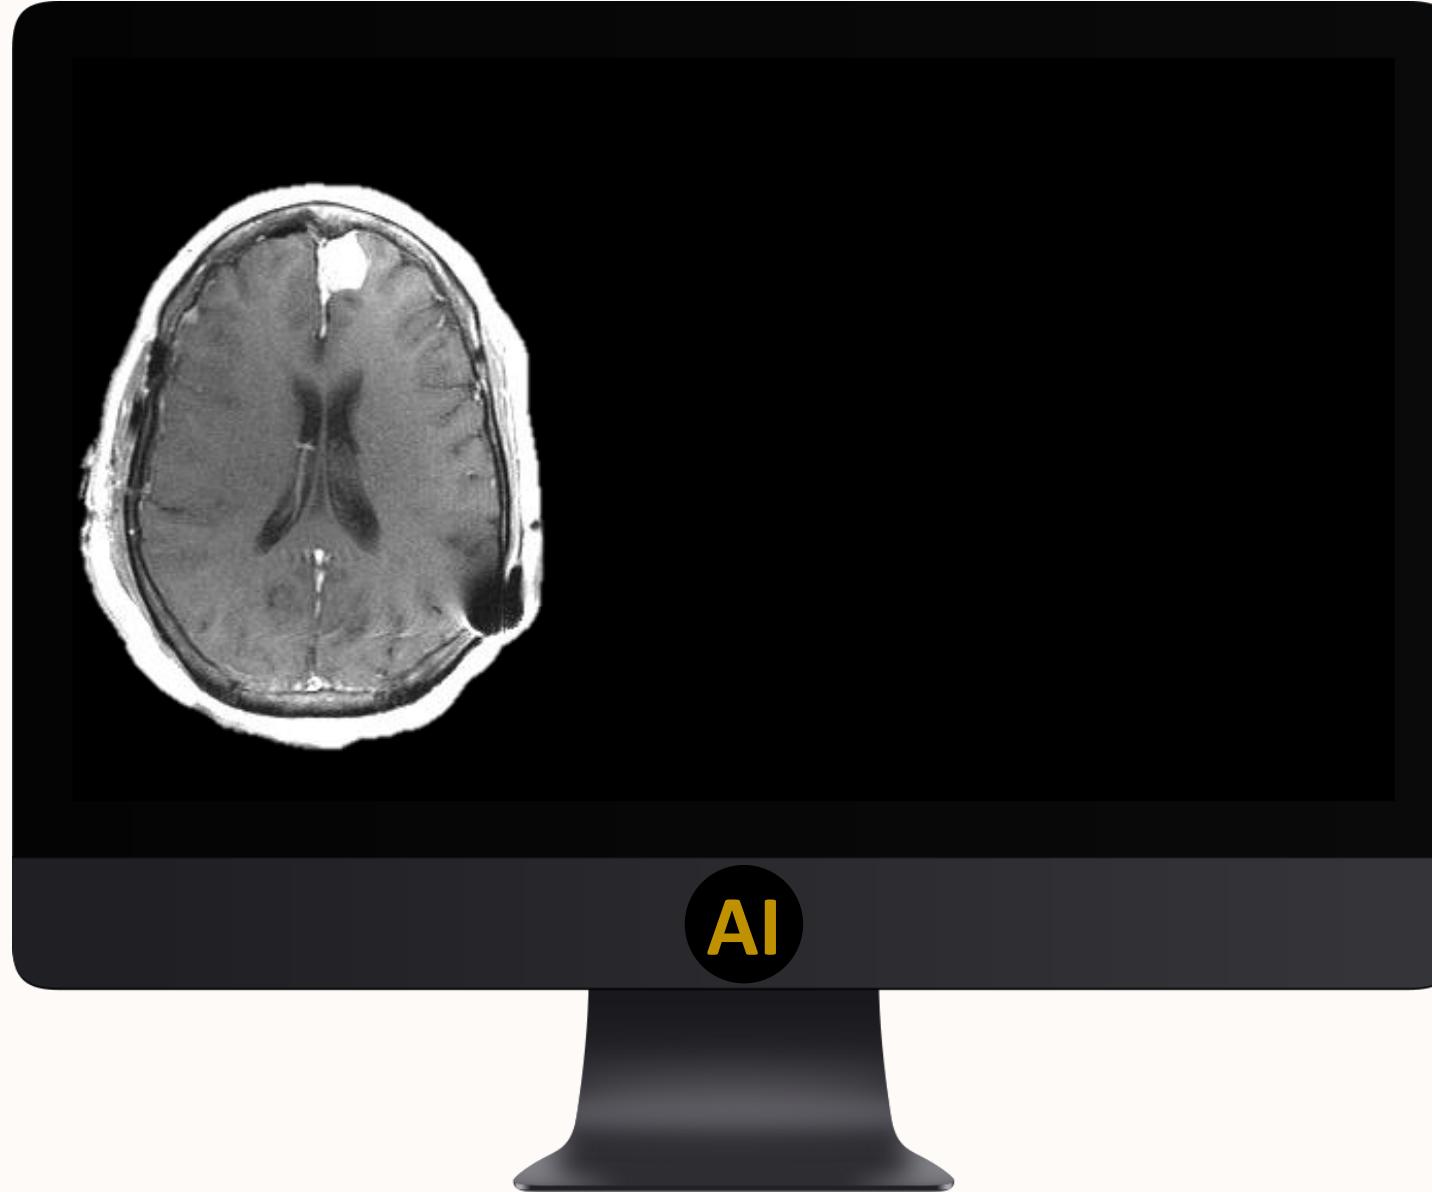
Expert Clinician

AI can detect common diseases

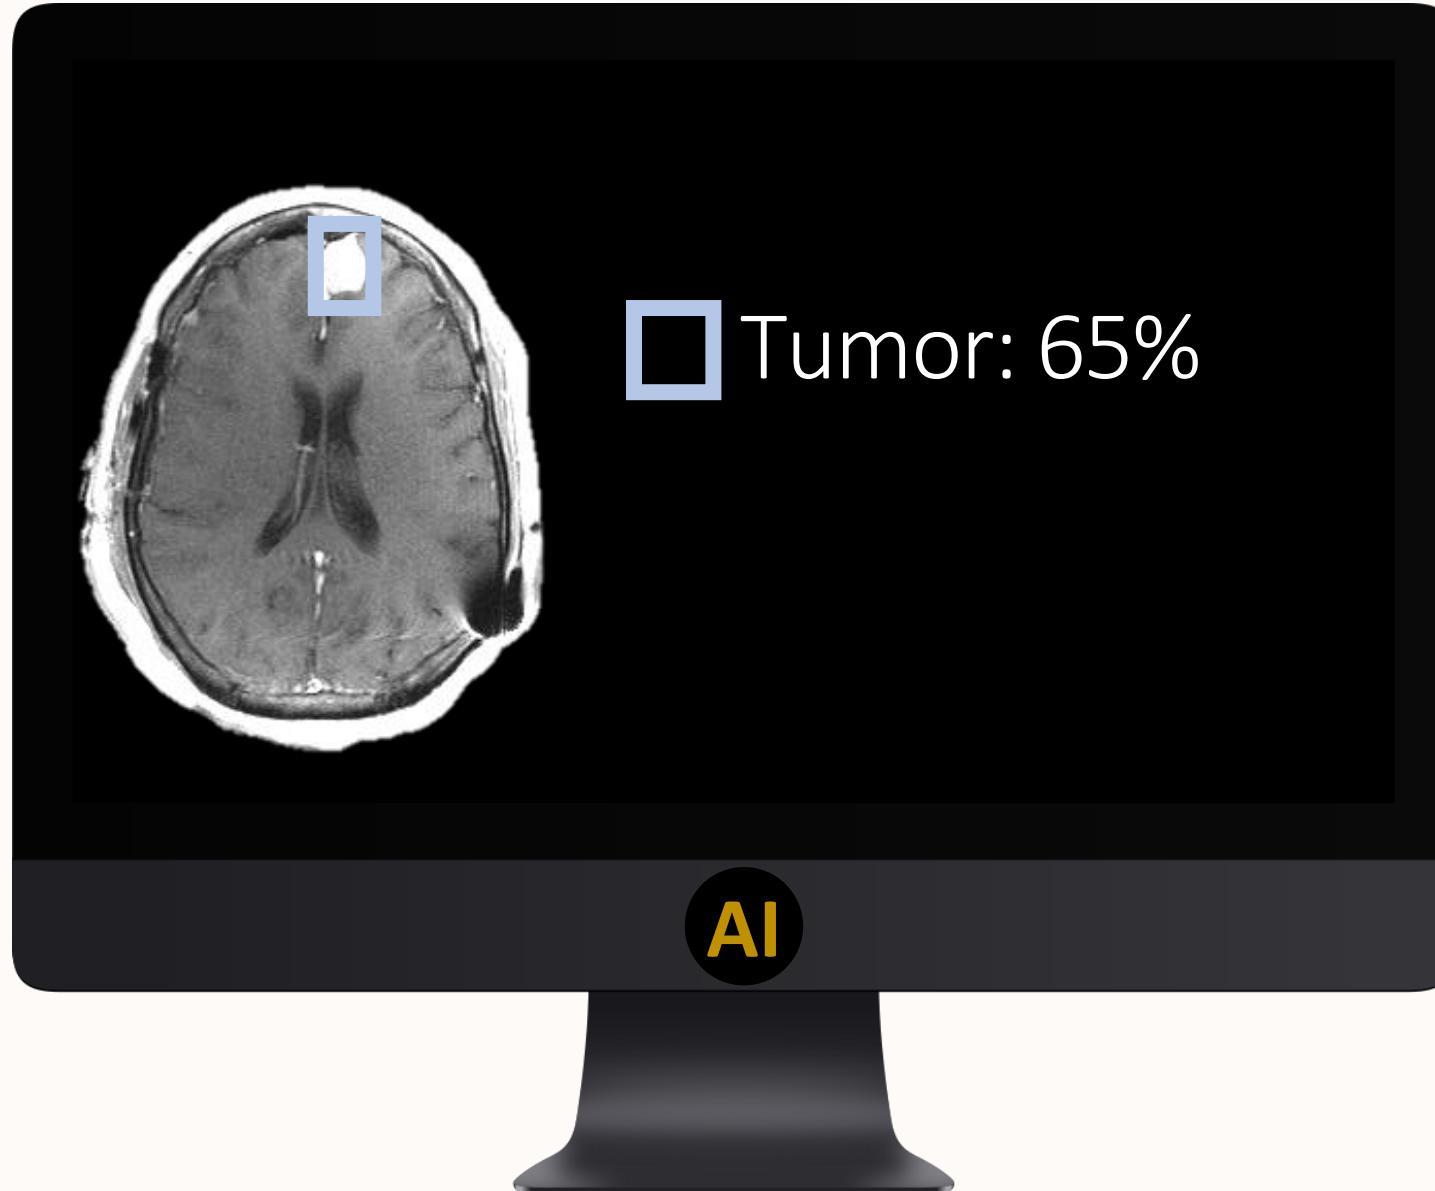
AI can detect common diseases



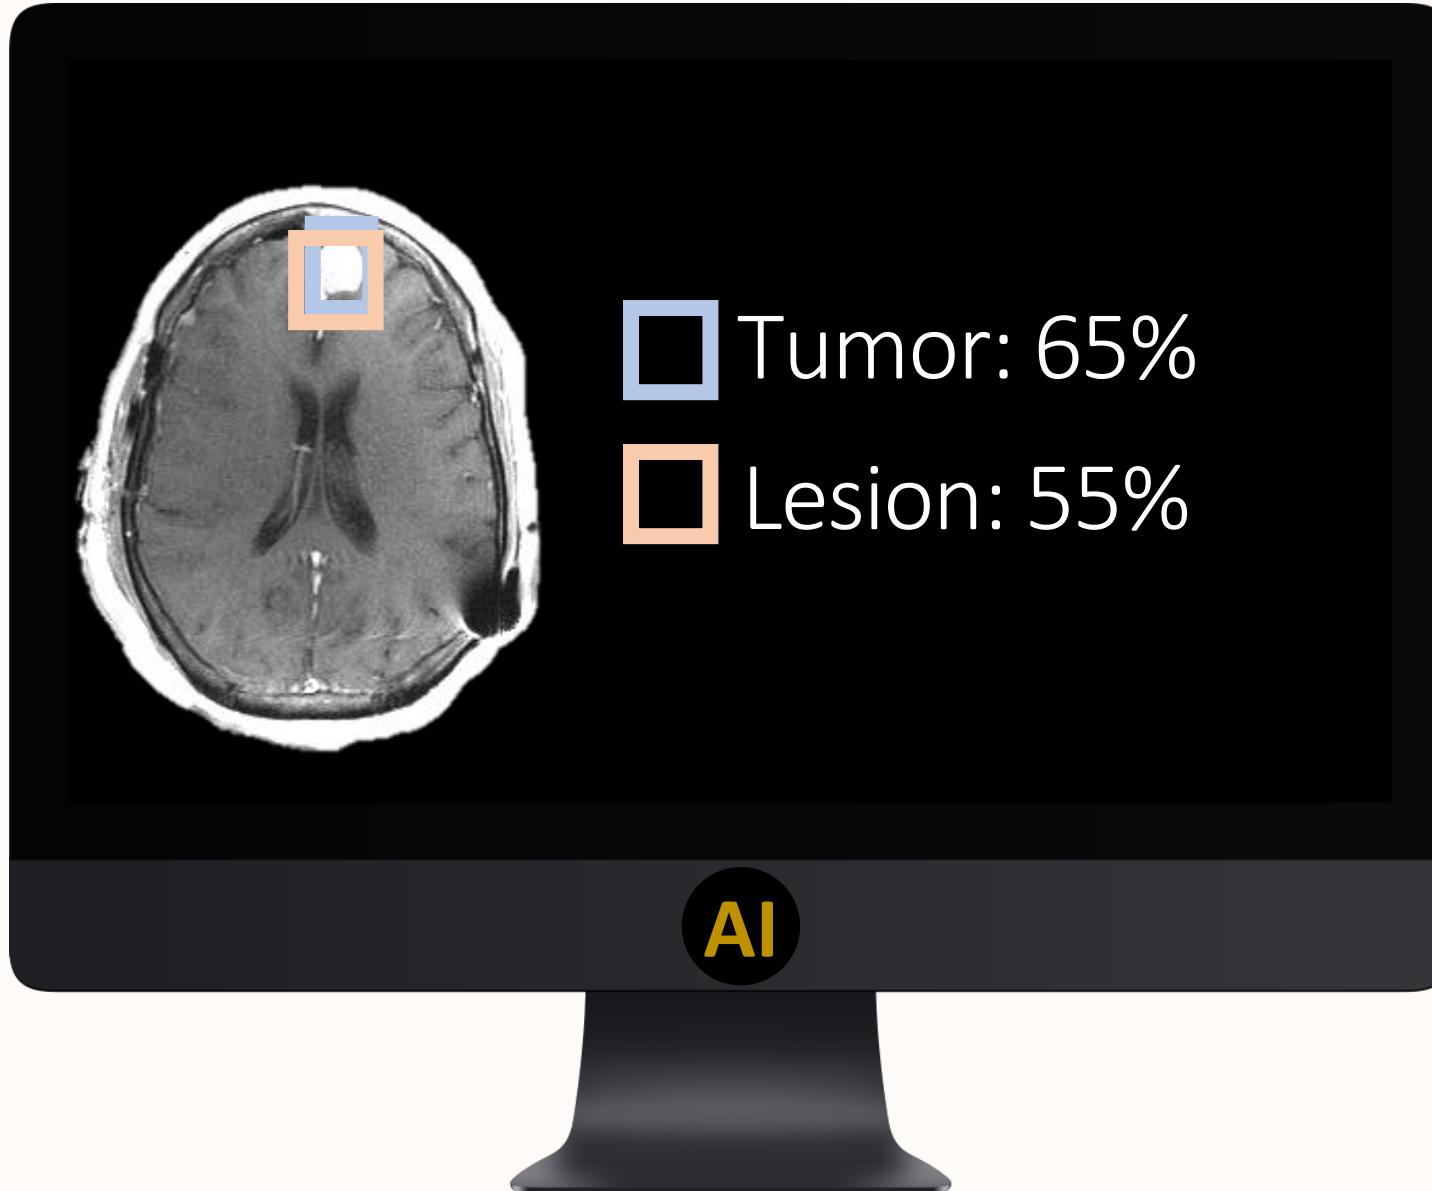
AI can detect common diseases



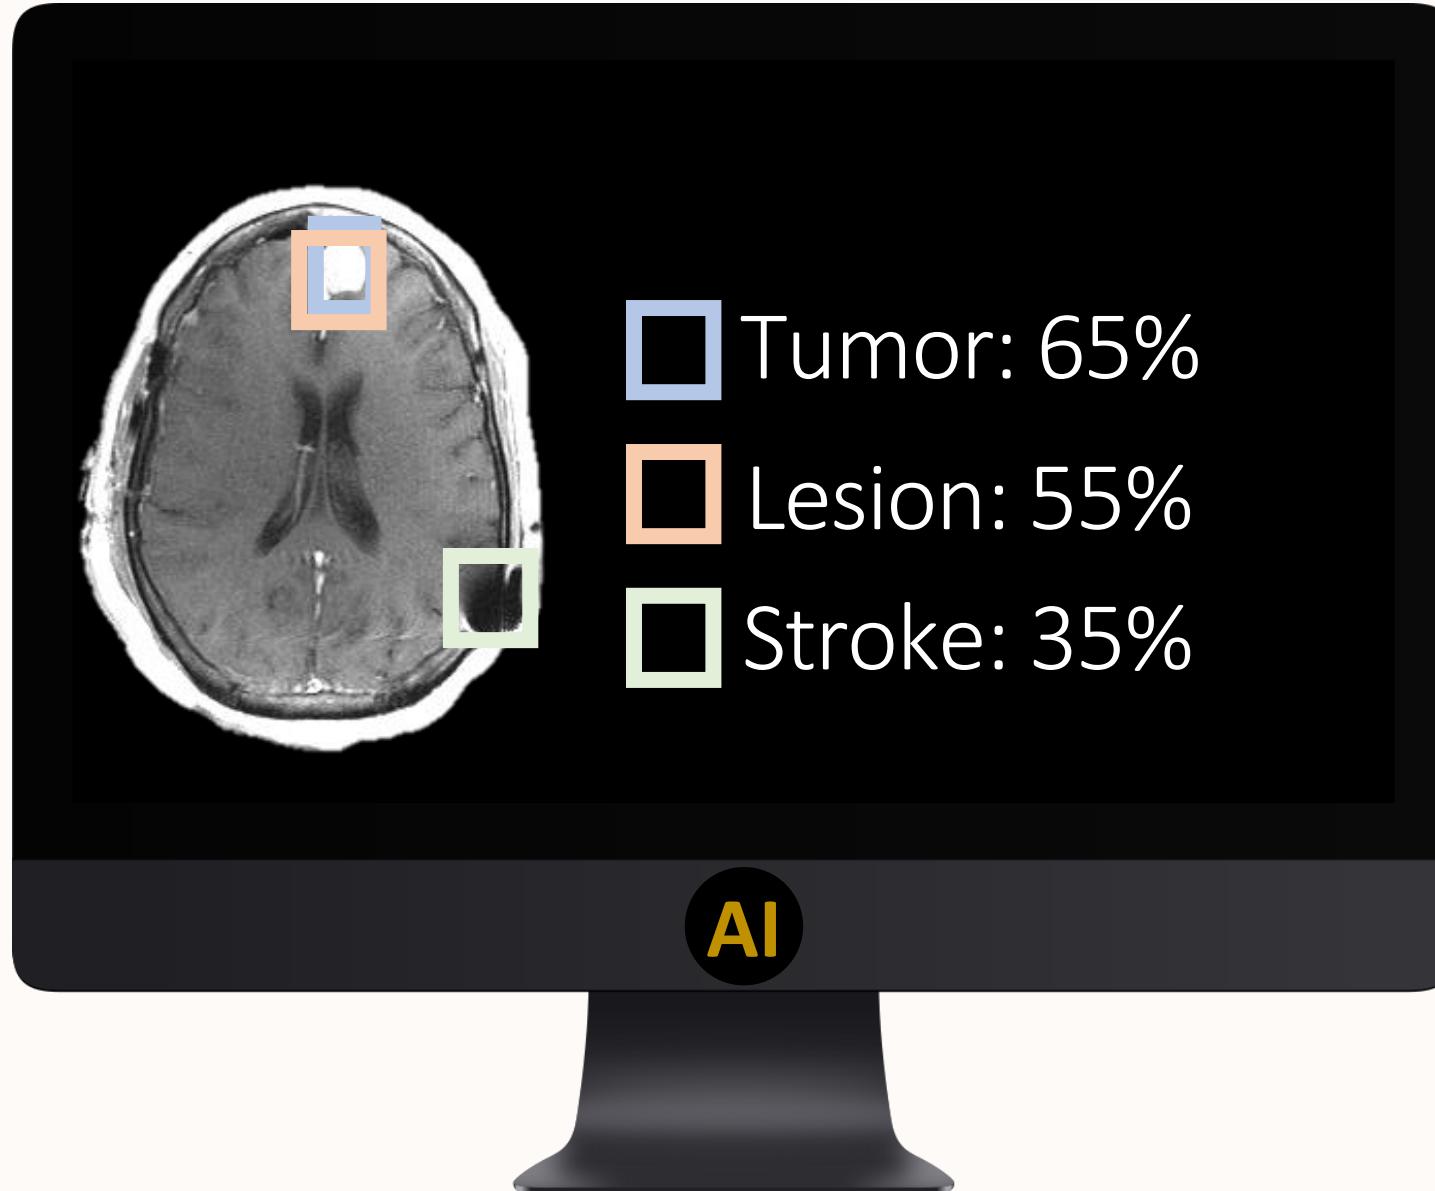
AI can detect common diseases



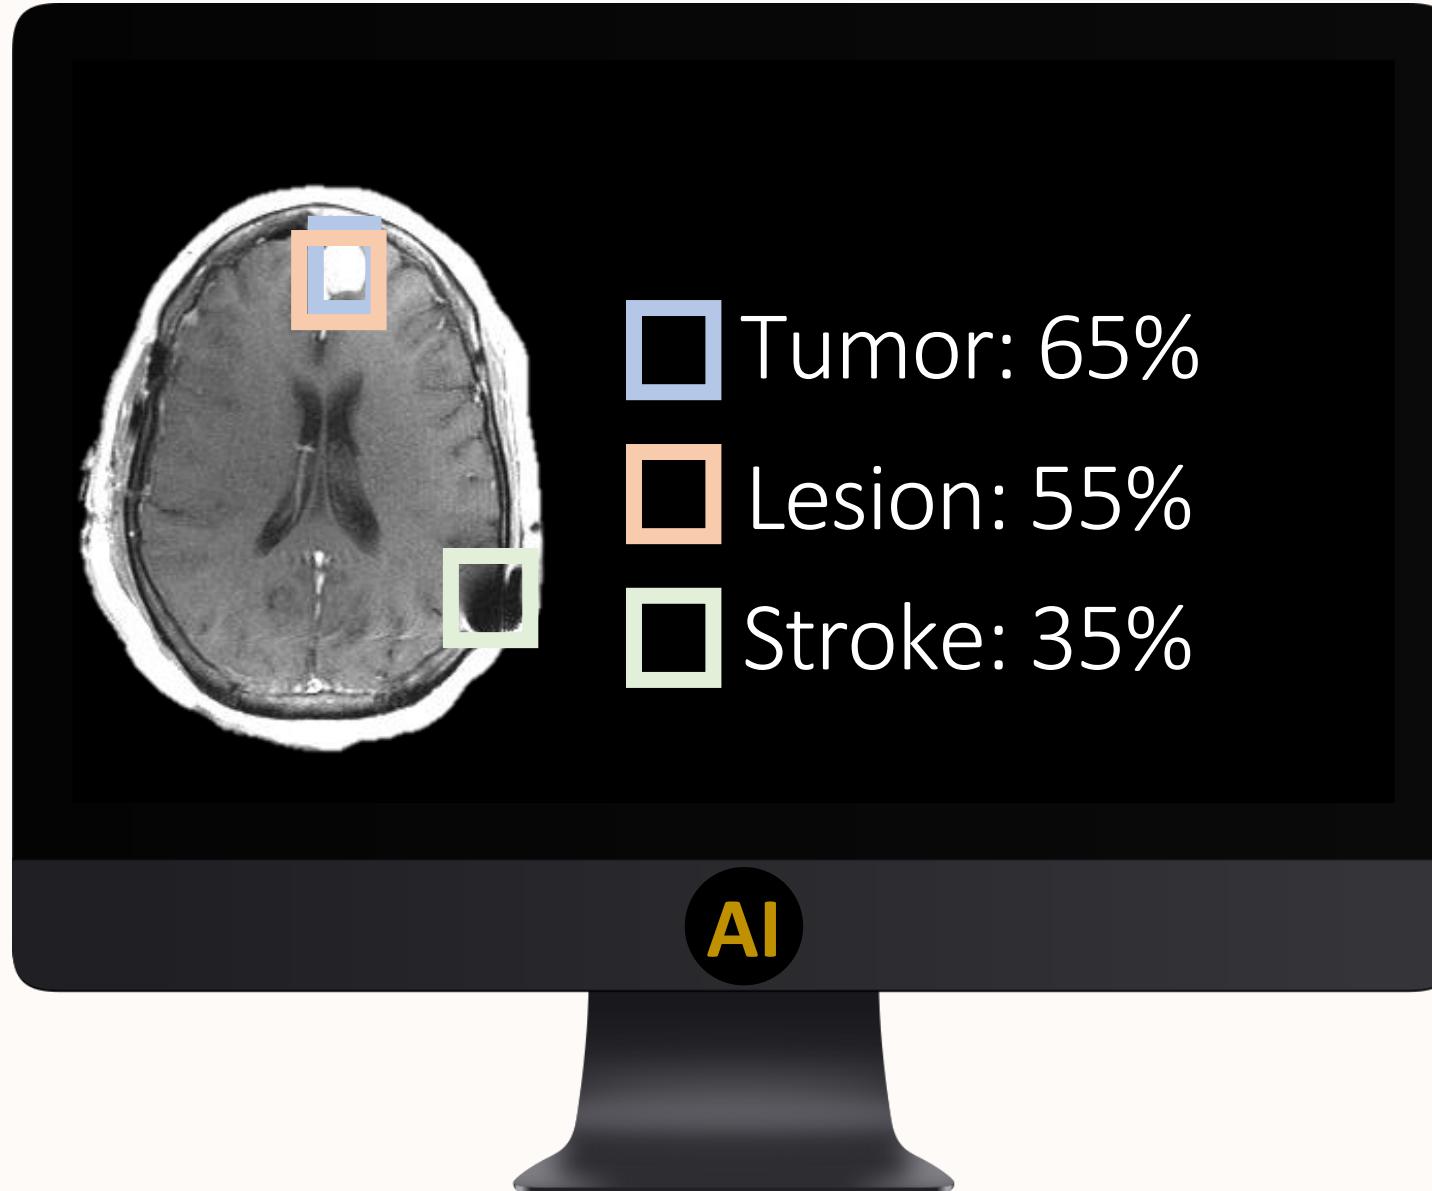
AI can detect common diseases



AI can detect common diseases

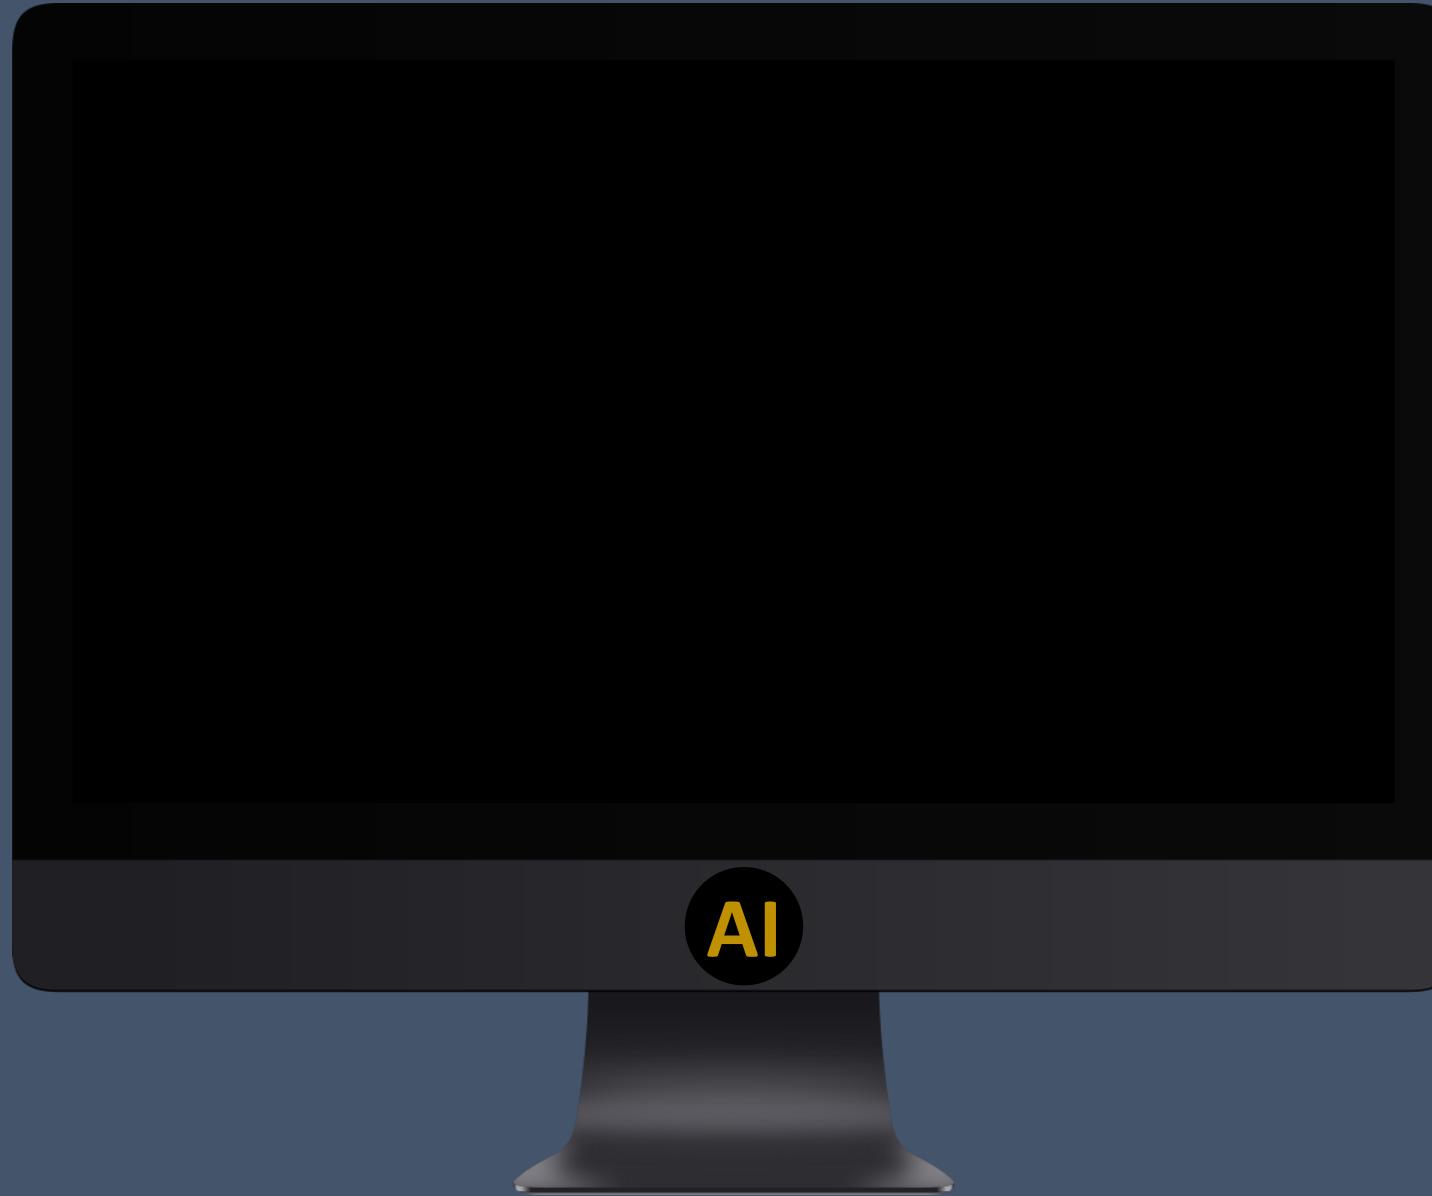


Can AI detect common diseases ?

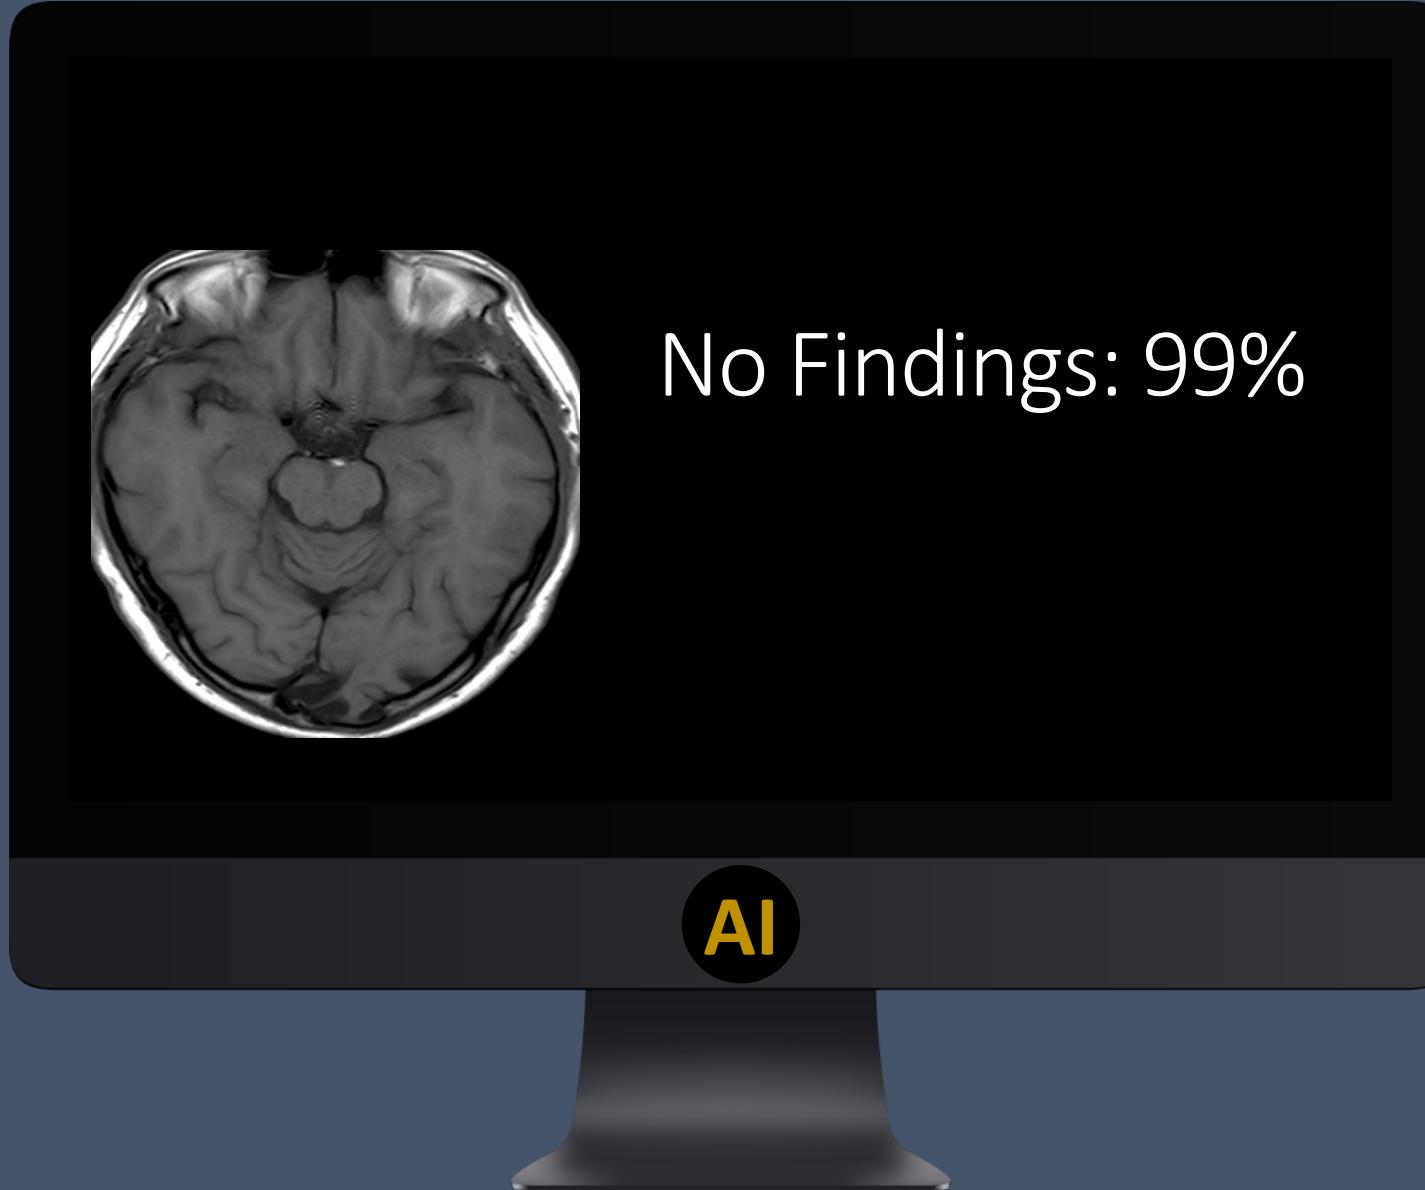


- Tumor: 65%
- Lesion: 55%
- Stroke: 35%
- No explanations
- Unintuitive “confidence”
- Prone to miss-diagnosis

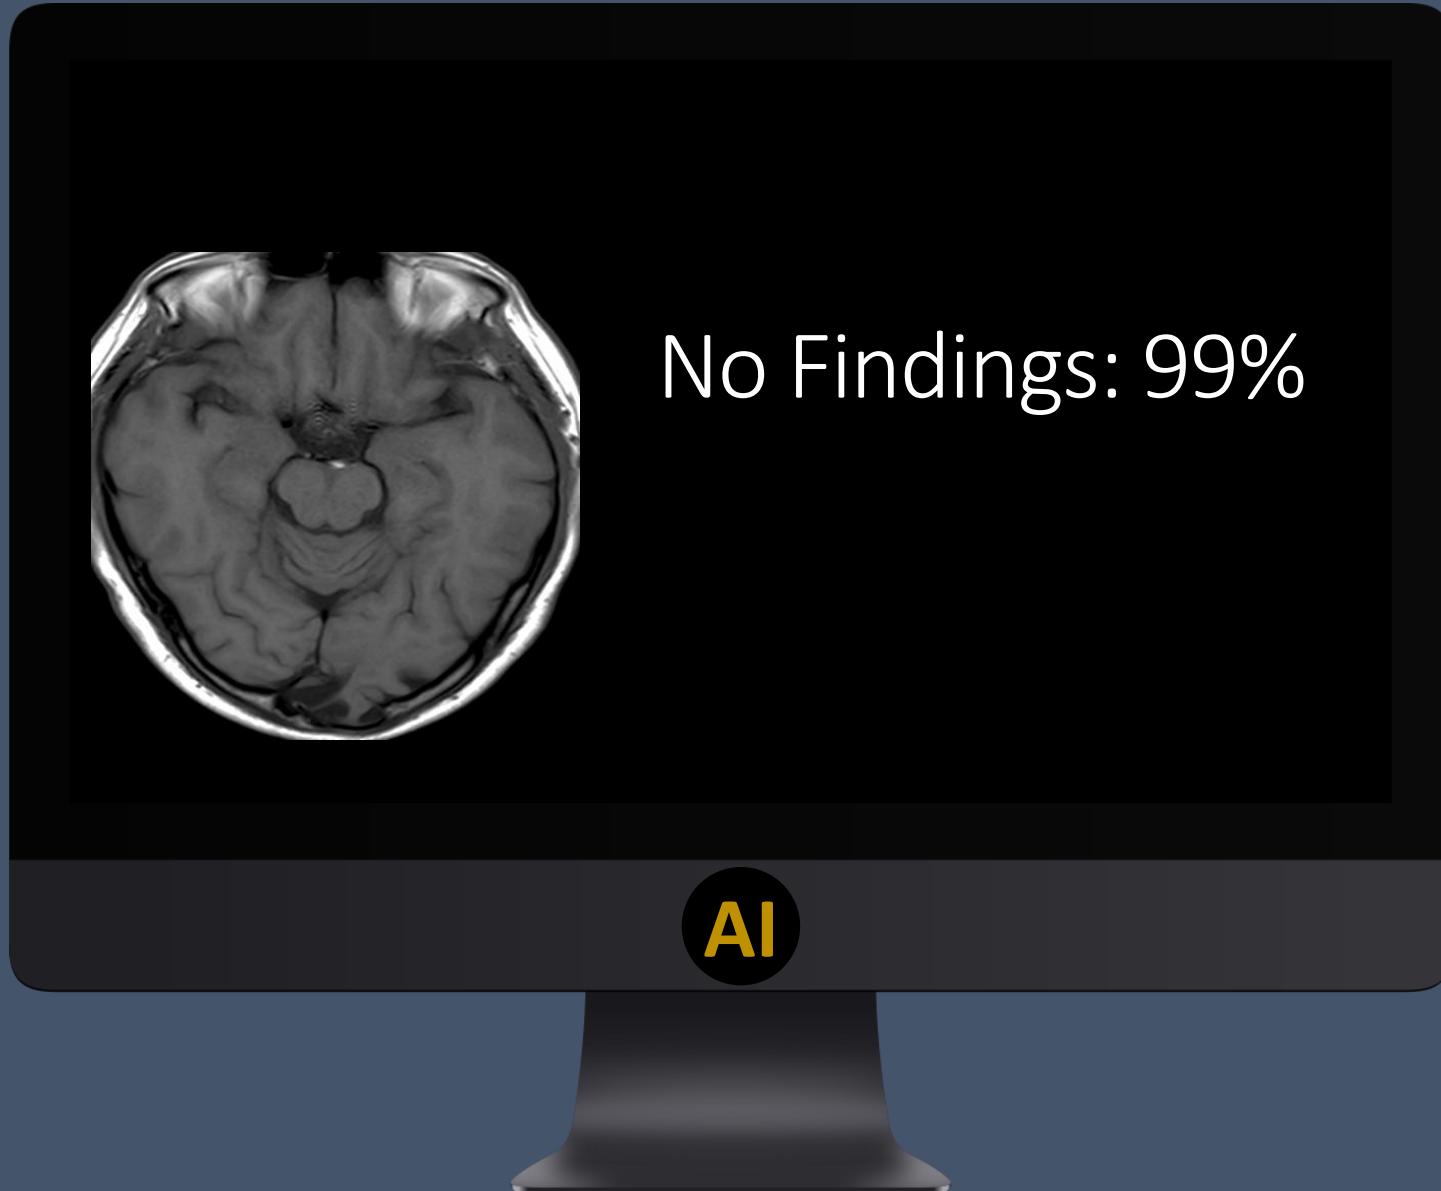
Can AI detect **rare** diseases ?



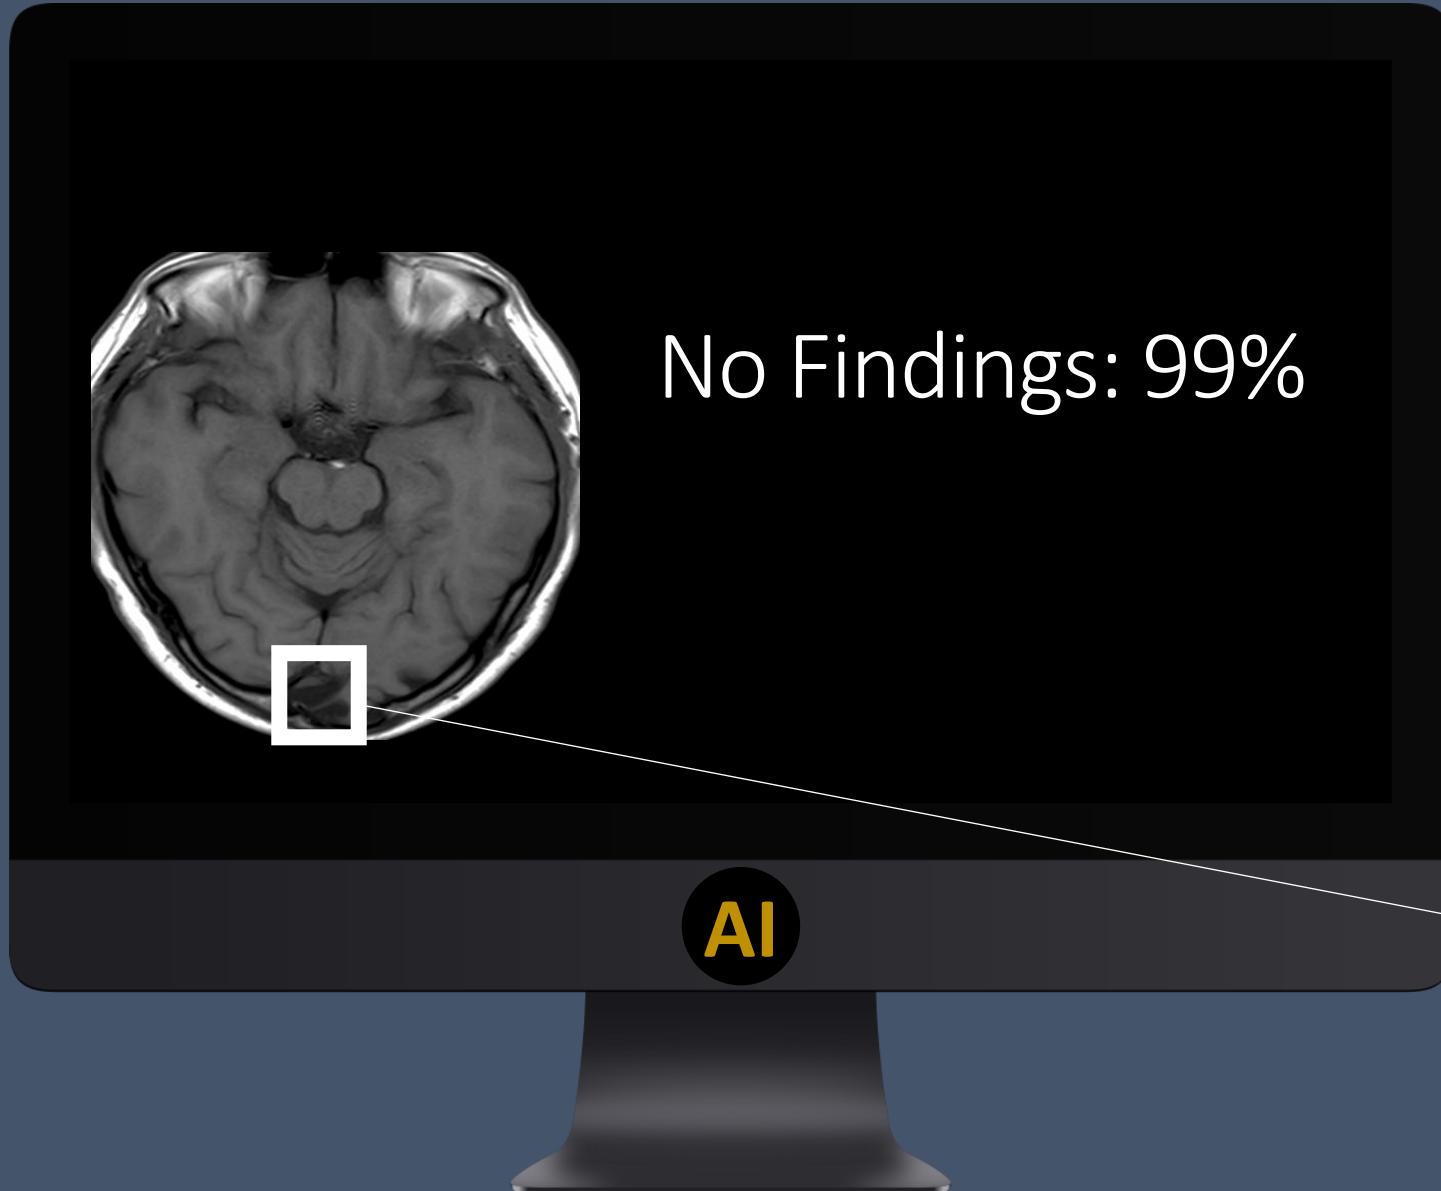
Can AI detect **rare** diseases ?



Can AI detect **rare** diseases ?



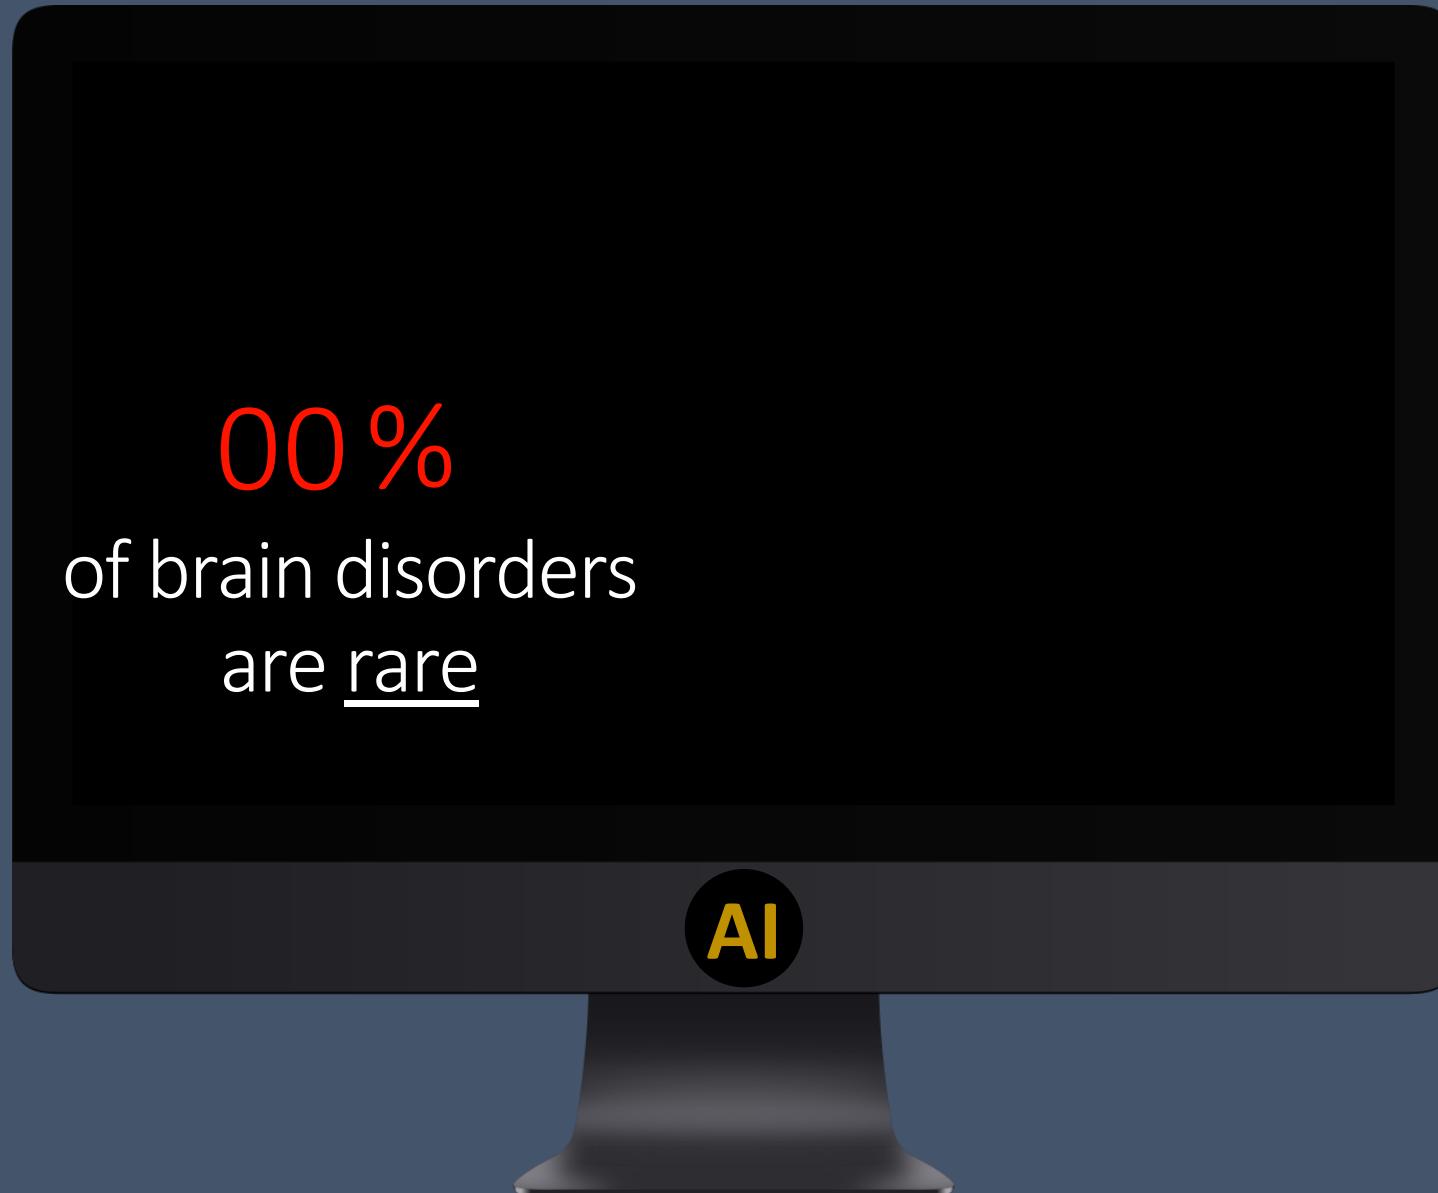
Can AI detect **rare** diseases ?



I think it's *occipital intradiploic encephalocele*



Can AI detect **rare** diseases ?



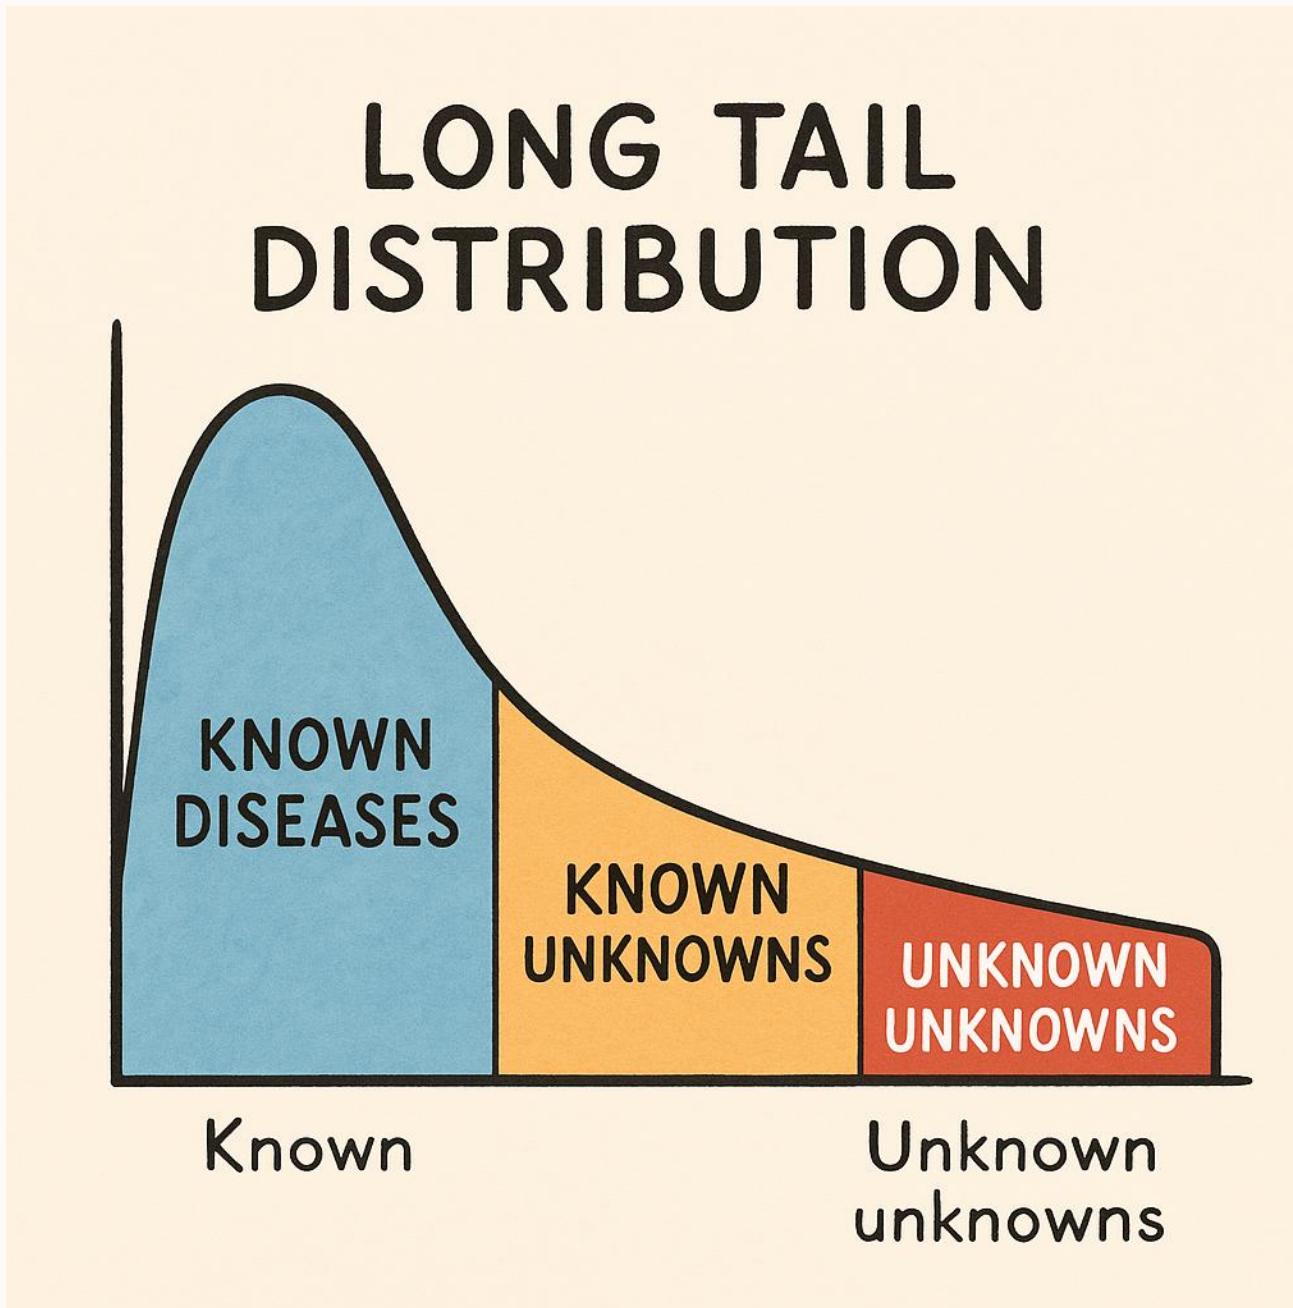
00%
of brain disorders
are rare



Can AI detect **rare** diseases ?



The Long Tail Distribution of Medical Anomalies



<https://bernhard-kainz.com>



How Do Doctors Detect the Unknown?

Doctors recognize subtle **deviations from normal**—not every possible disease.

Learning the Normal to Detect the Unknown

Train on Healthy Only

Model learns to reconstruct normal anatomical patterns

Generate Anomaly Maps

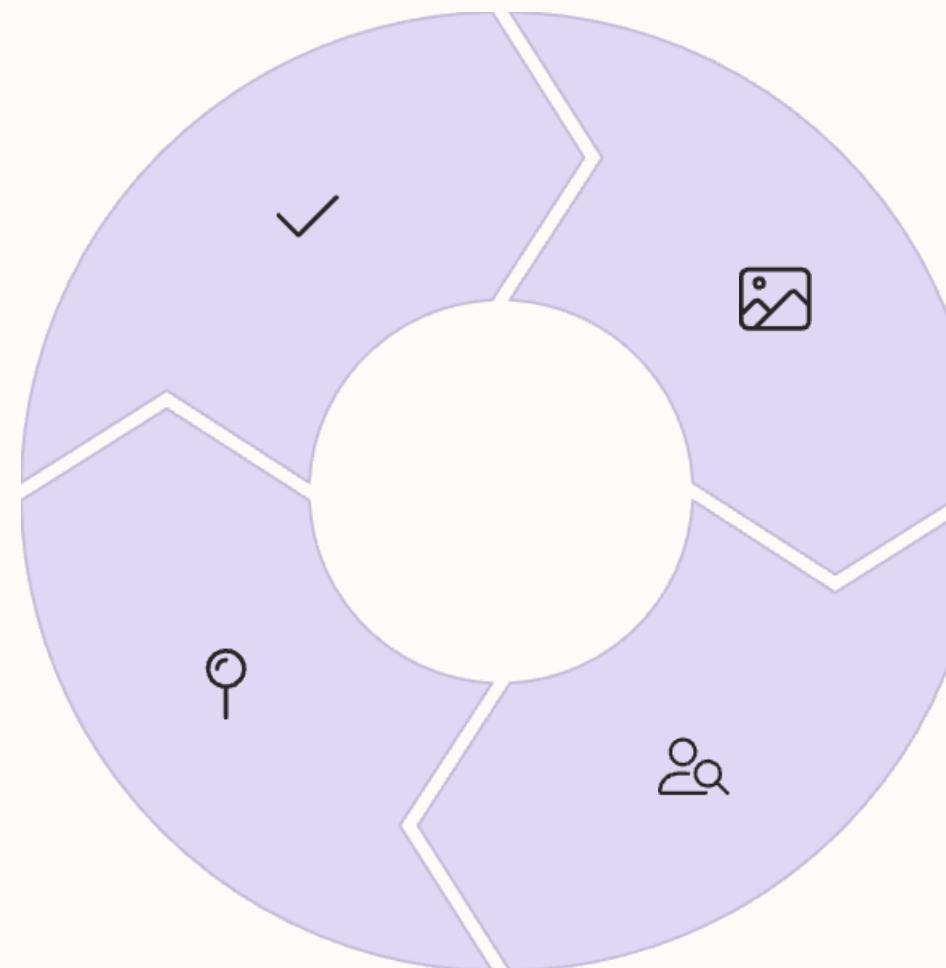
Highlight deviations for clinical interpretation

Test on Any Image

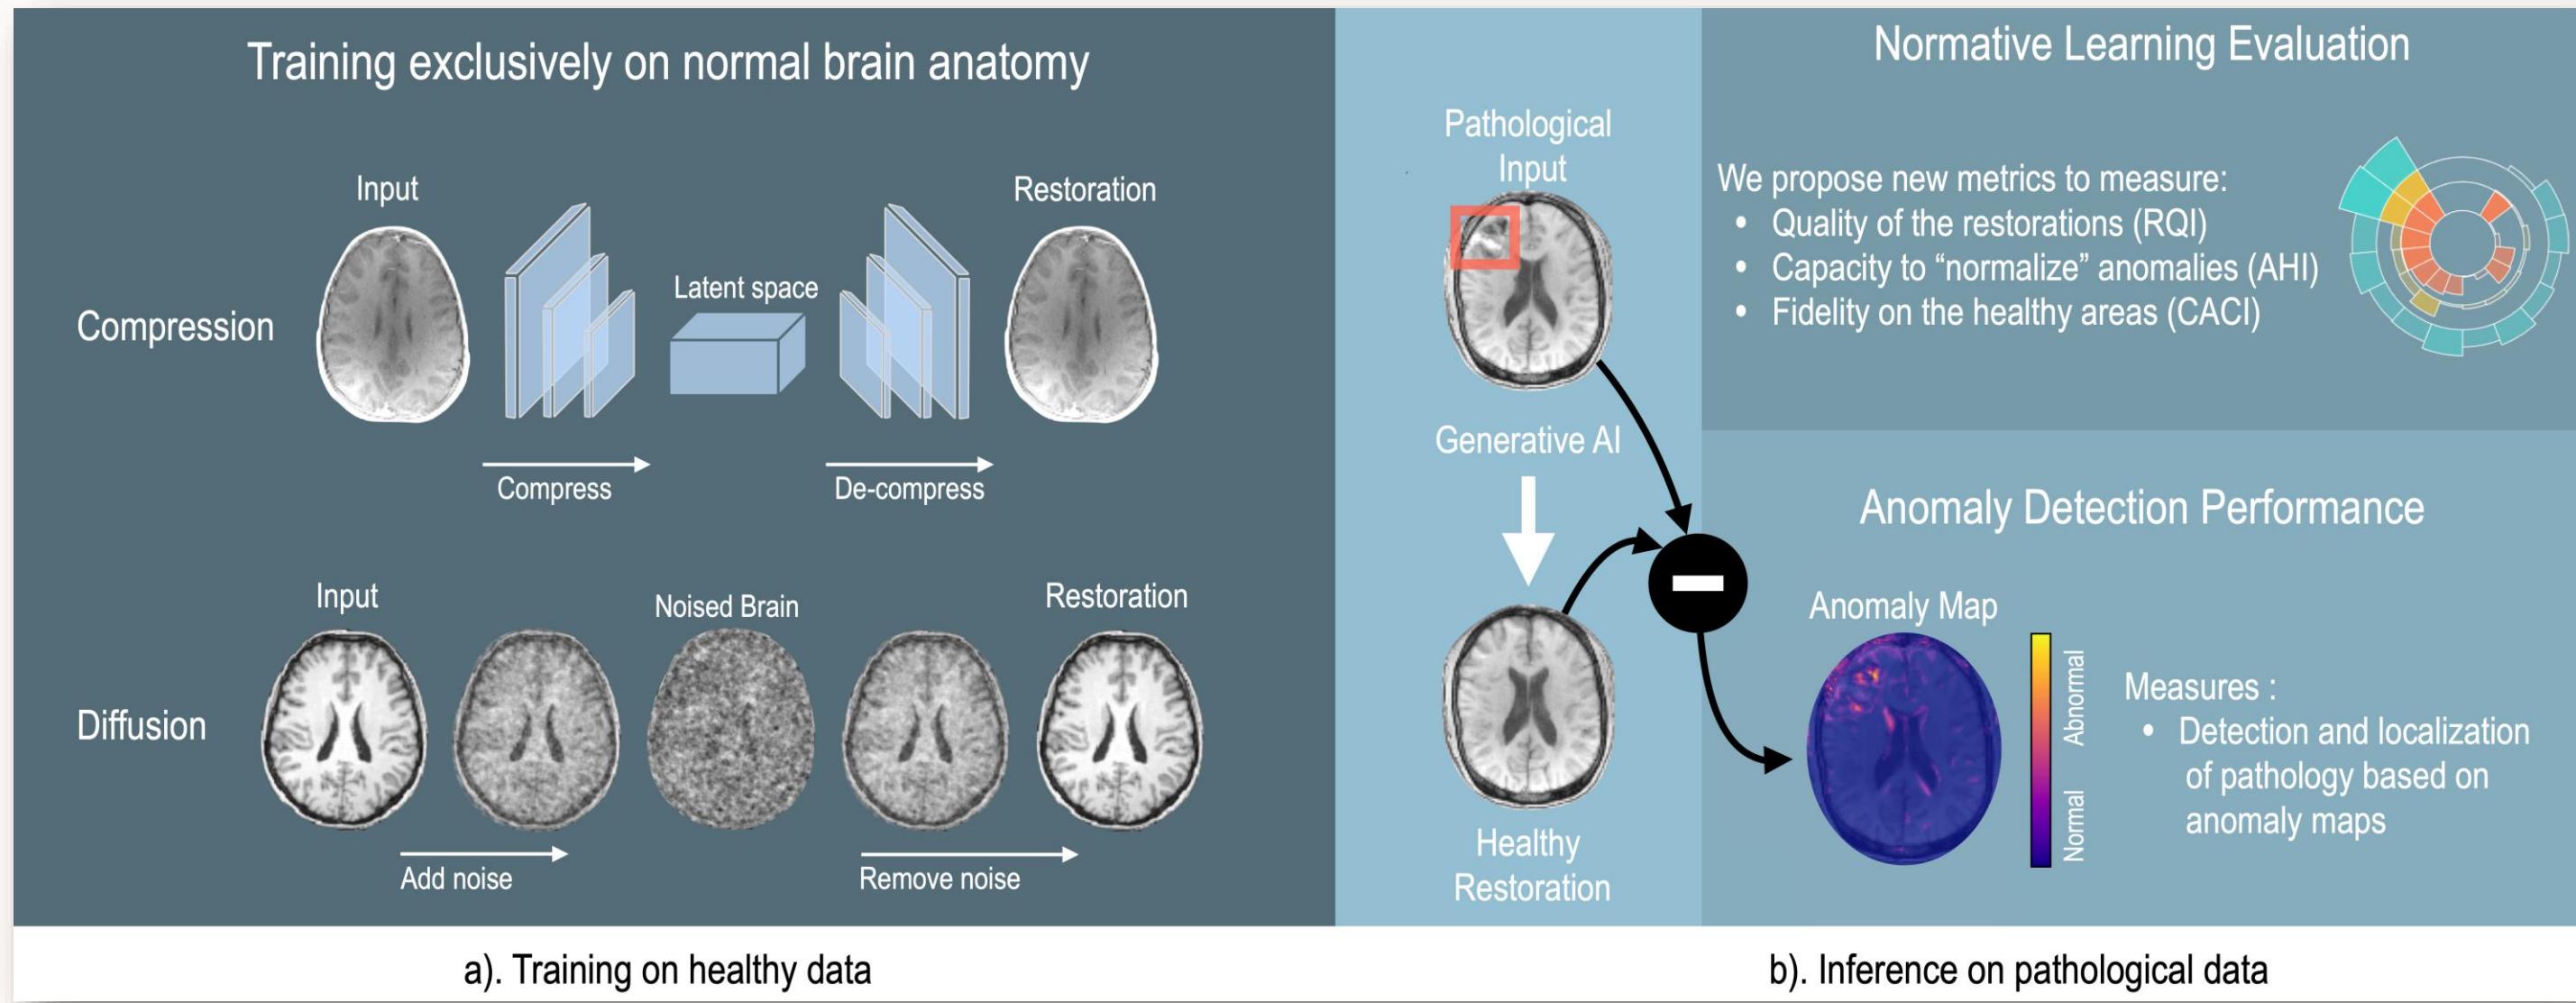
Apply to images that may contain pathologies

Compute Differences

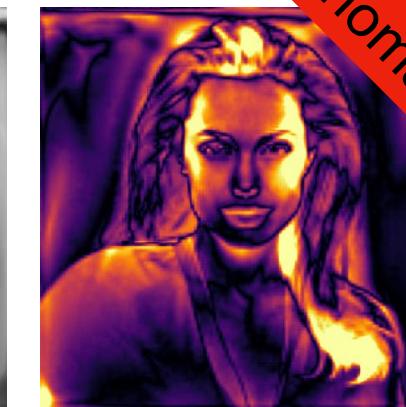
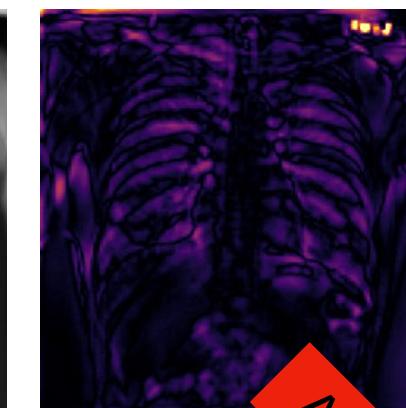
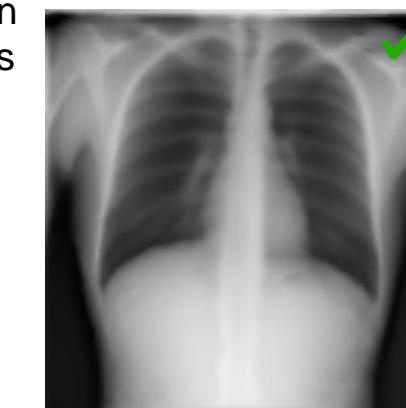
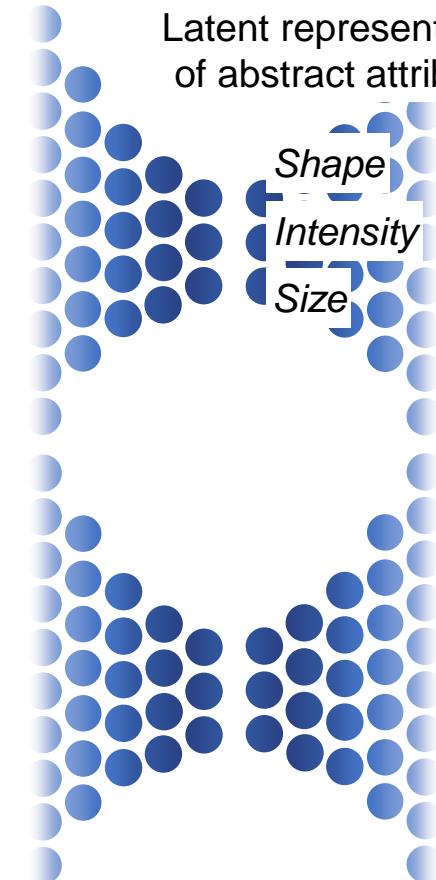
Discrepancies reveal anomalous regions



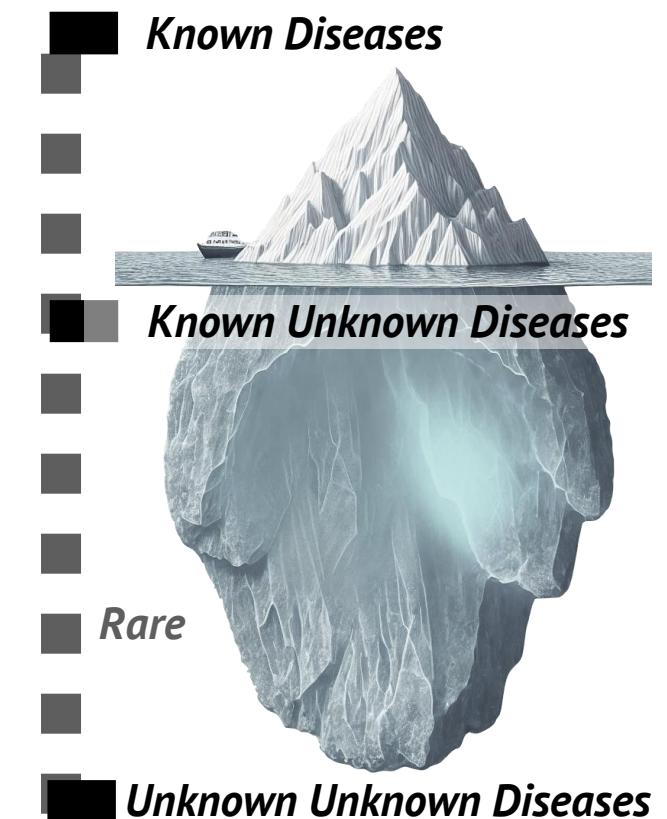
Learning the Normal to Detect the Unknown



Learning the Normal to Detect the Unknown



- ✓ In-distribution, e.g., healthy chest X-rays
- ✗ Out-of-distribution, e.g., pathology



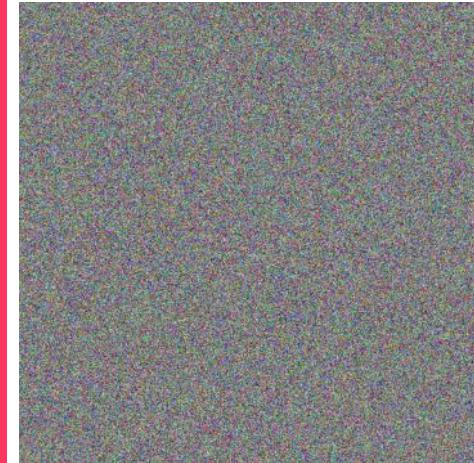
Bercea, C. I., Rueckert, D., & Schnabel, J. A. (2023, October). What do aes learn? challenging common assumptions in unsupervised anomaly detection. In *International Conference on Medical Image Computing and Computer-Assisted Intervention* (pp. 304-314). Cham: Springer Nature Switzerland

Diffusion Models for Anomaly Detection

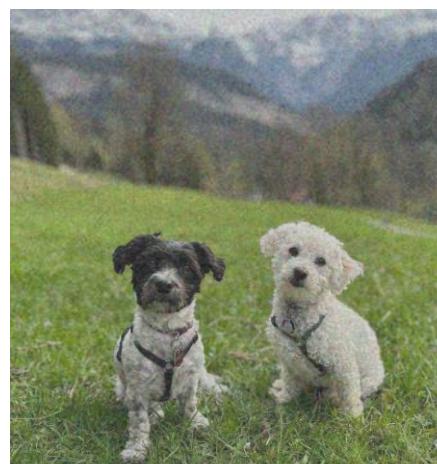
Goal: Learn the data distribution by gradually *adding* and *removing* noise!

In anomaly detection, the **backward** starts from a **selected time t** .

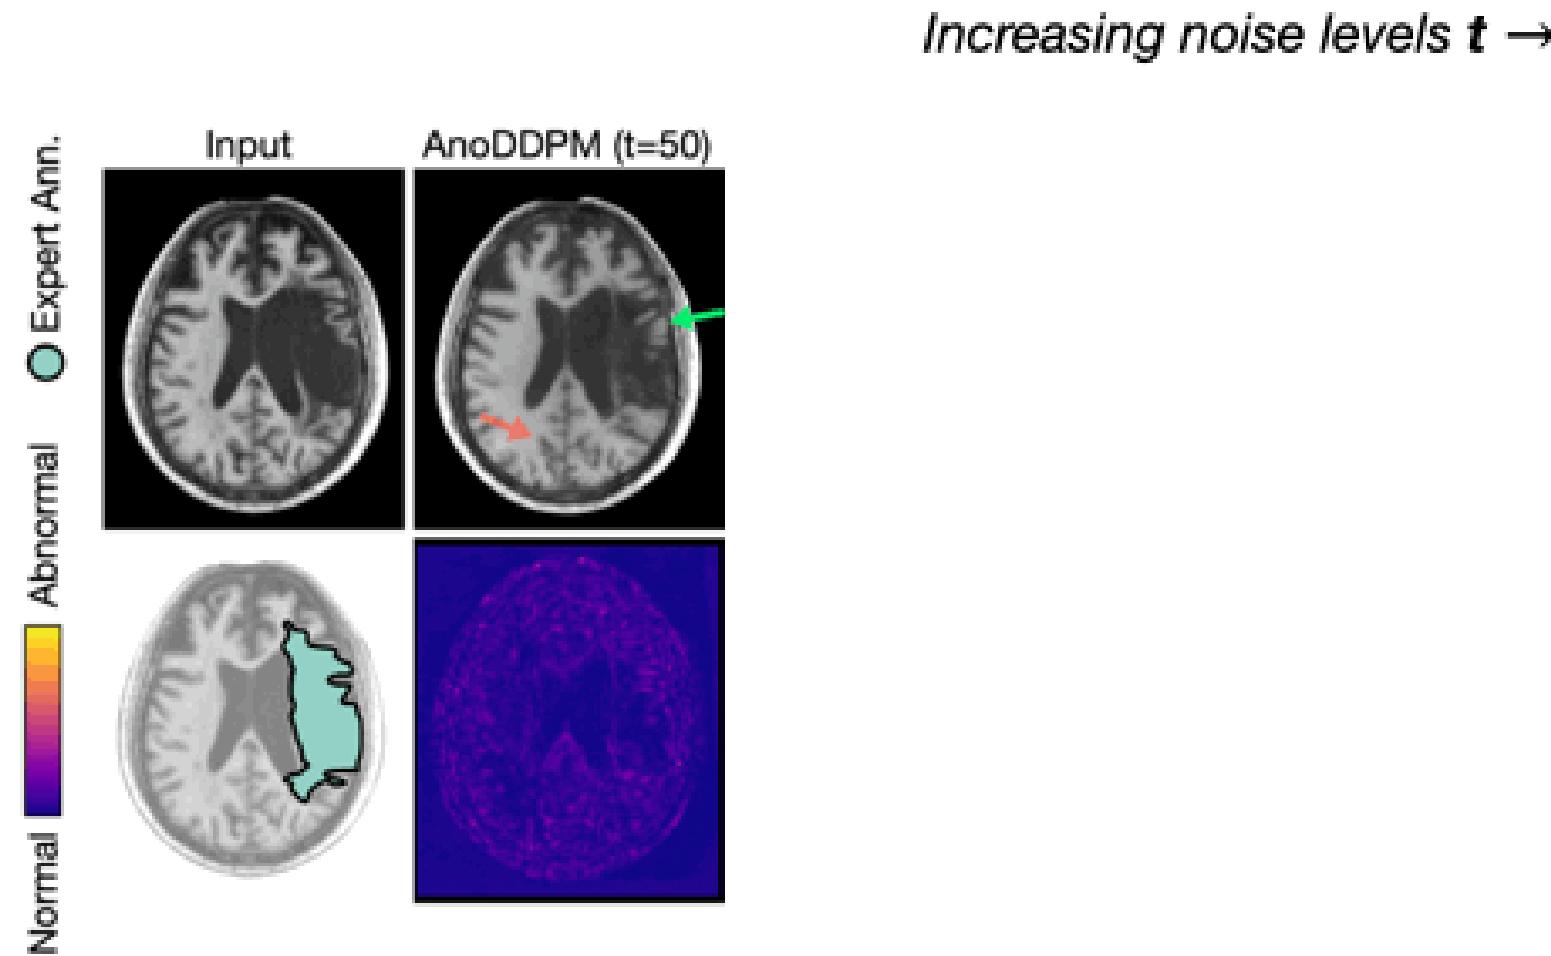
Forward Process



Backward Process



The noise paradox



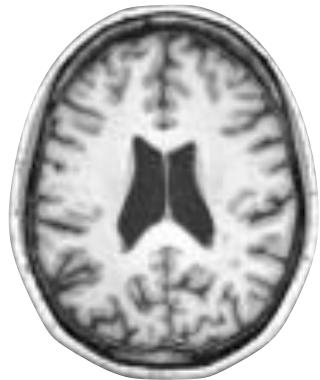
Noise Paradox: Removing anomalies vs. Preserving healthy tissues

Reversing the Abnormal



GANs

Goal: The Generator improves over time until the Discriminator can no longer tell the difference between real and generated data.



Real brain MRI



Artist (Generator)

Creates fake data (e.g., MRI) trying to fool the critic.



Critic (Discriminator)
Evaluates data authenticity and provides feedback.

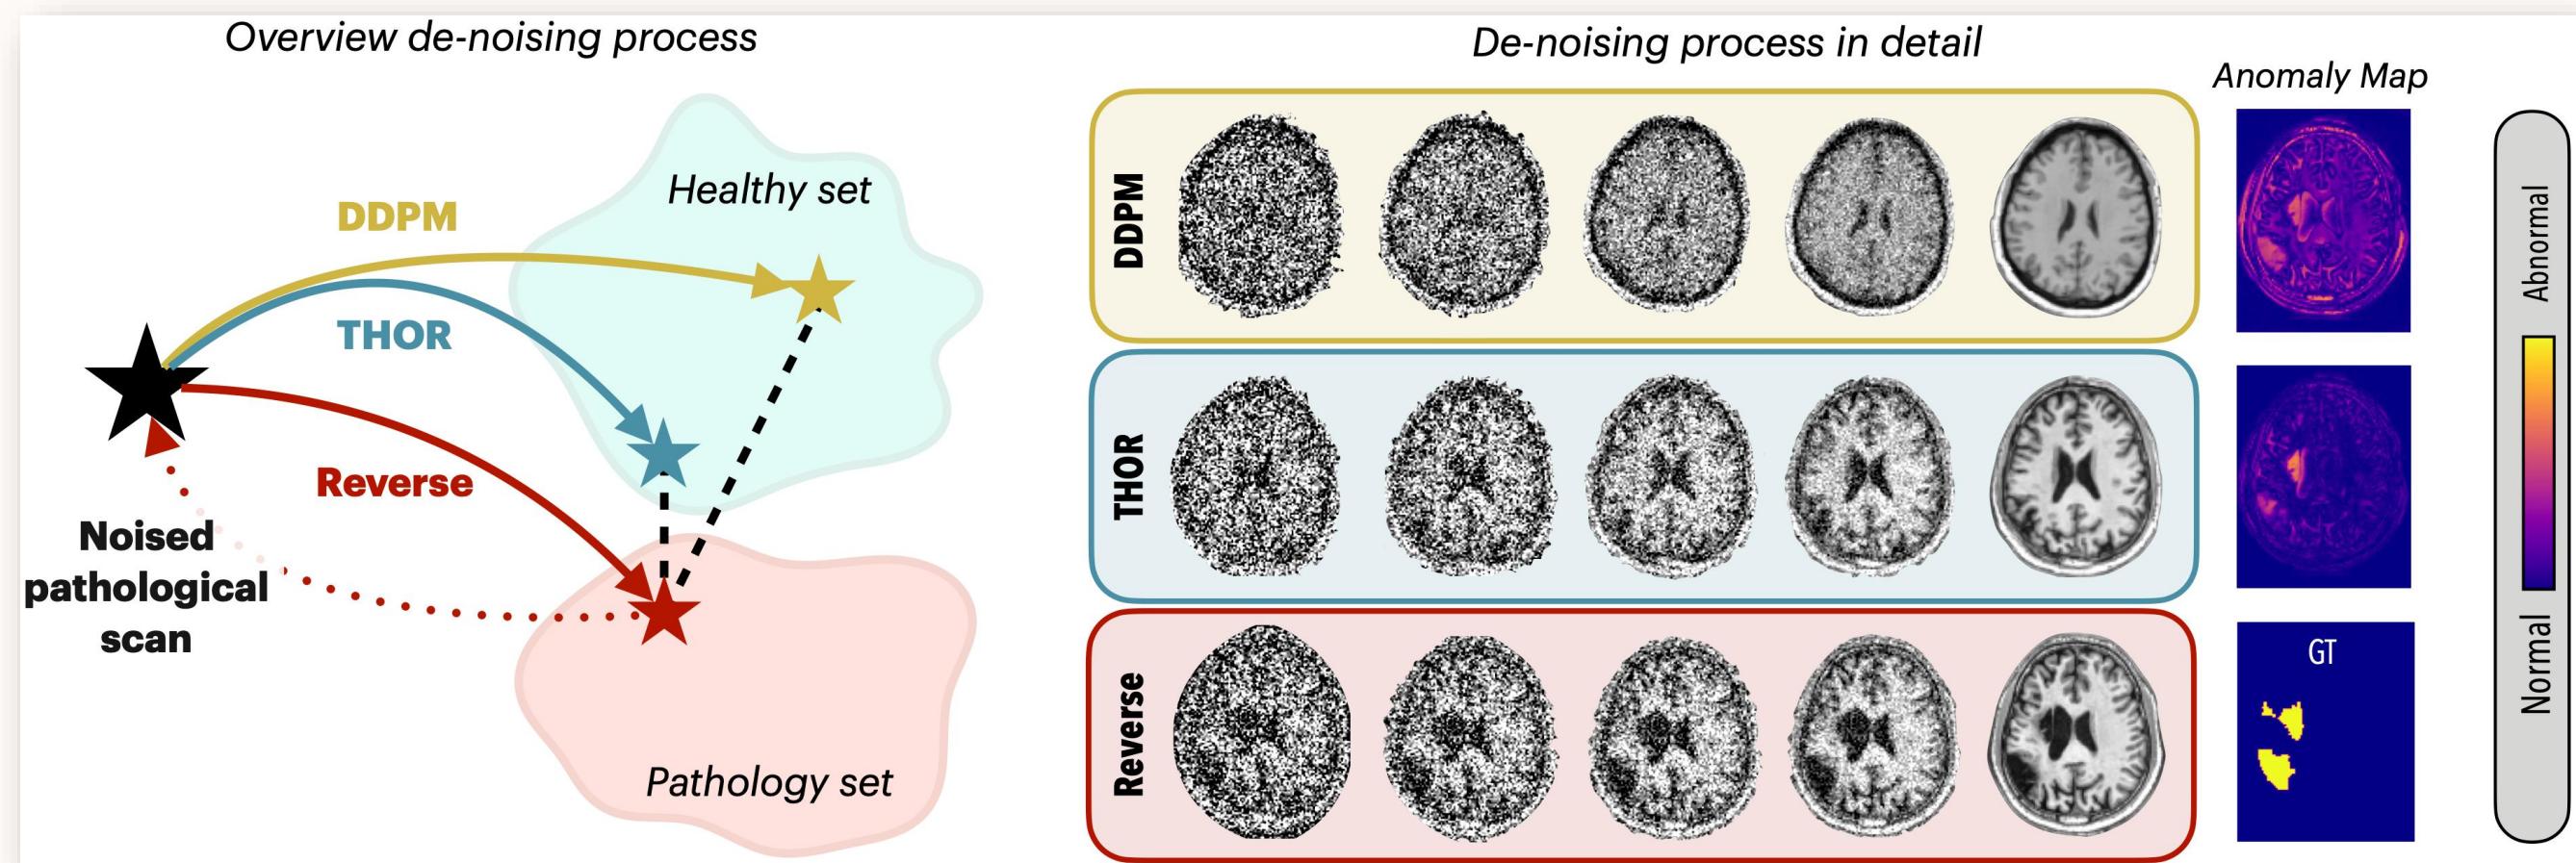
Mask, Stich, and Re-Sample



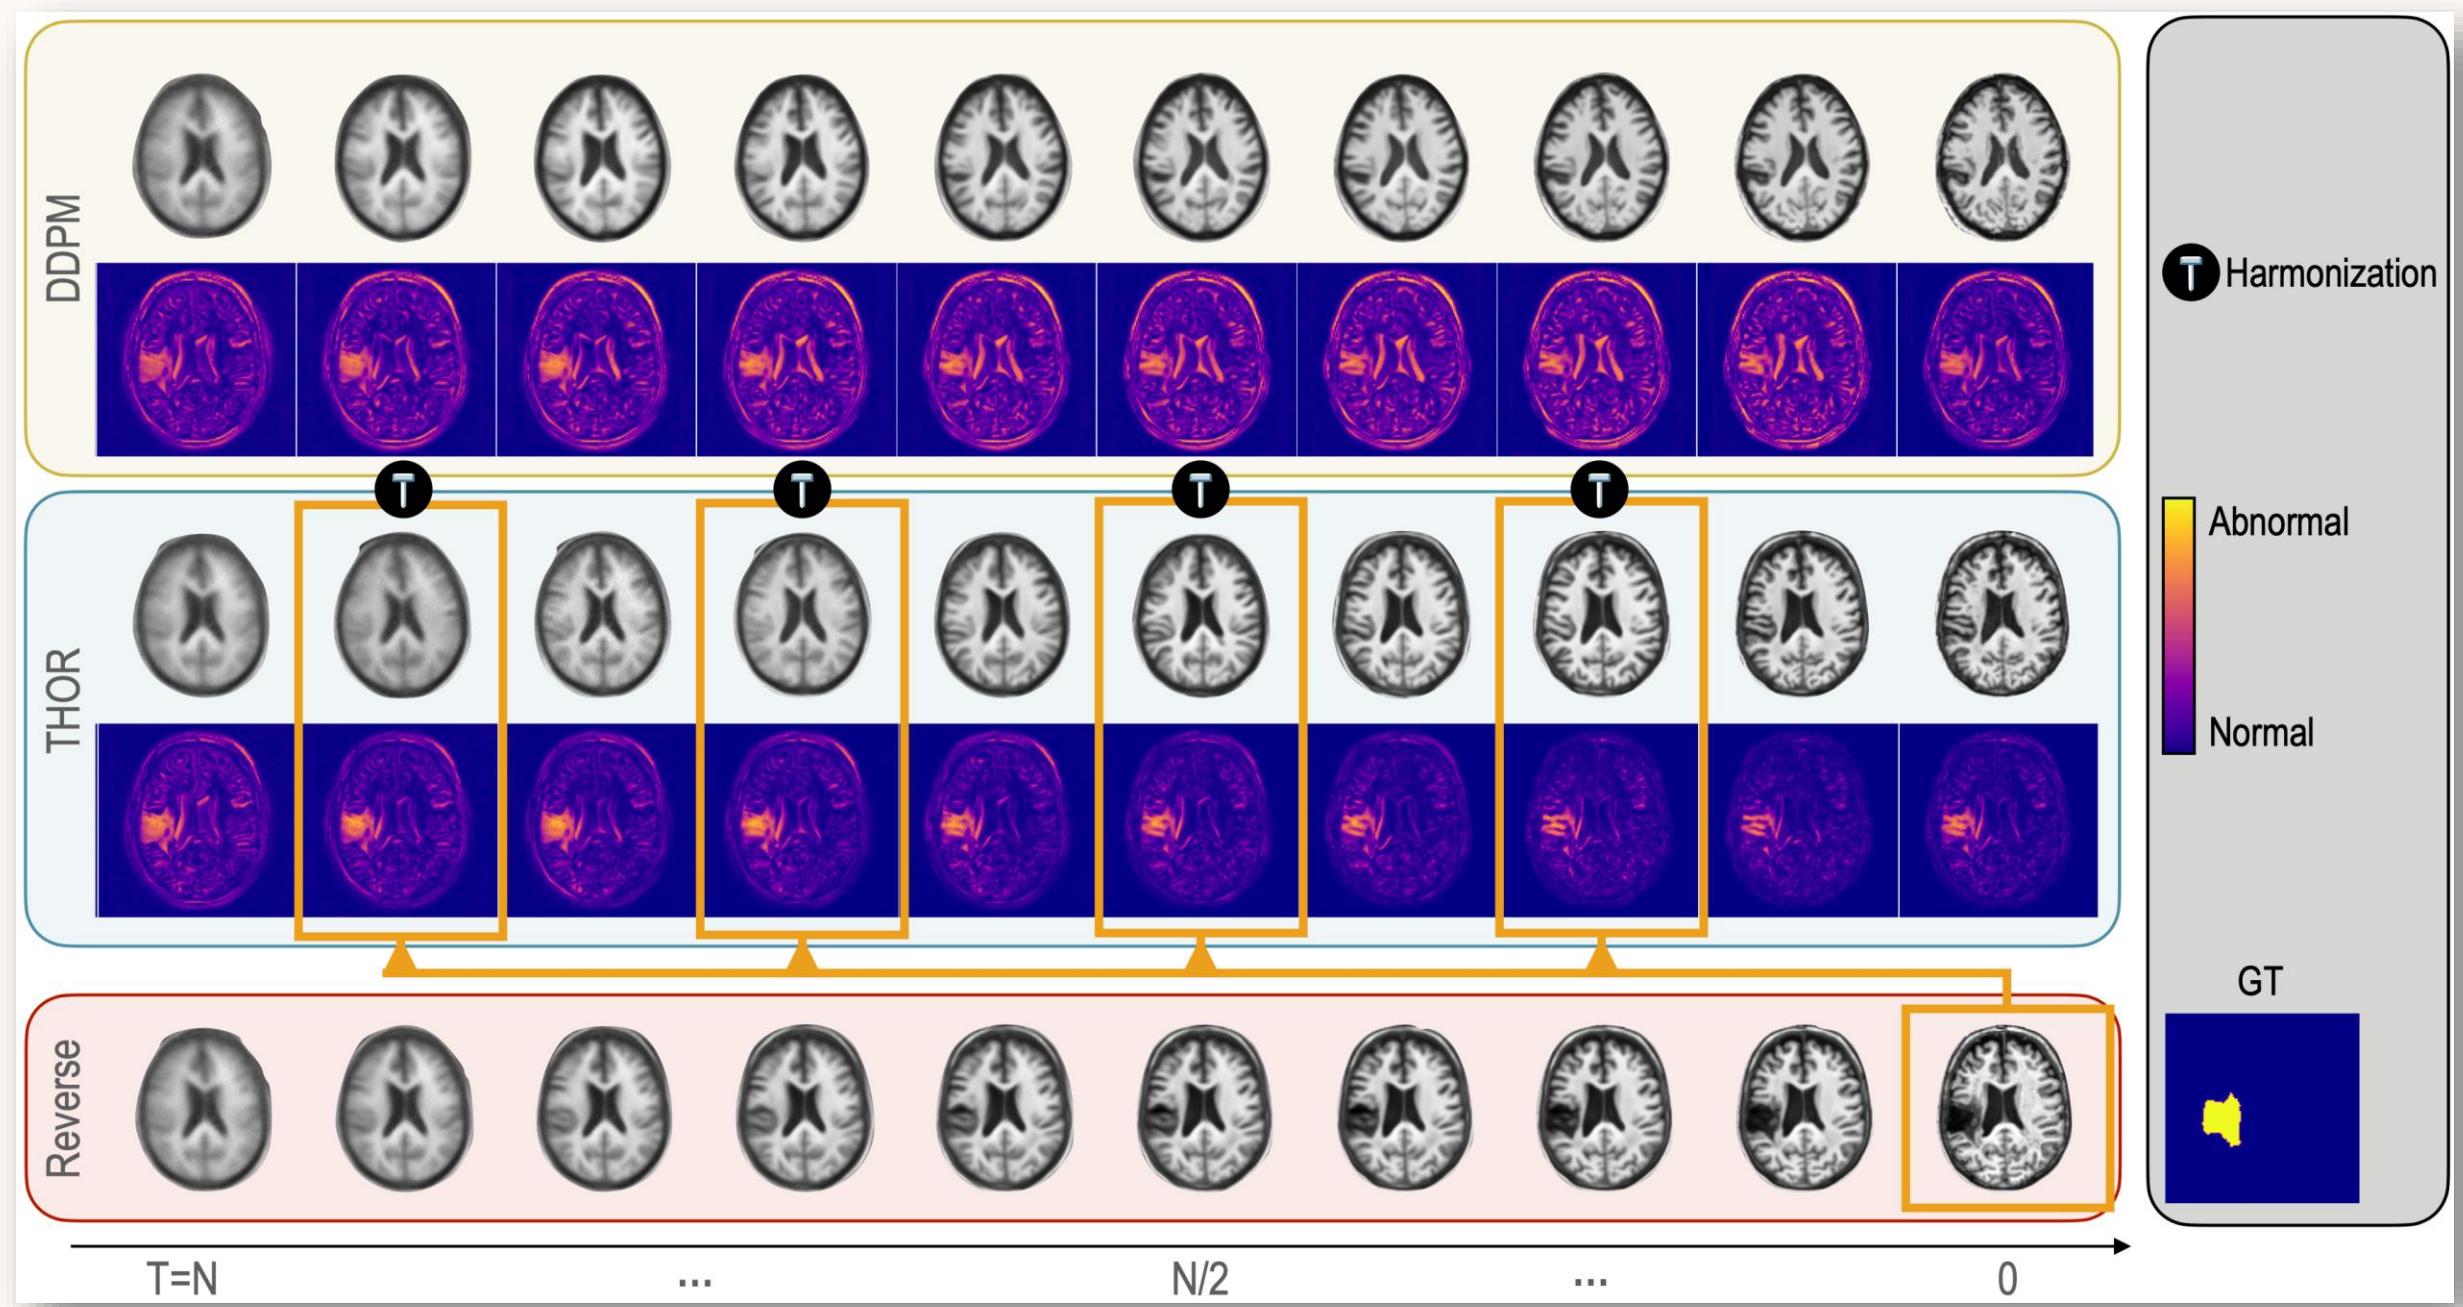
Michael Neumayr



THOR



THOR



Vision-Language Diffusion Models



Text Conditioning

Natural language prompts guide image generation and editing.



Semantic Transformation

Convert descriptions like "healthy" to "showing tumor" visually.



Rare Disease Discovery

Language models help identify unusual patterns in images.





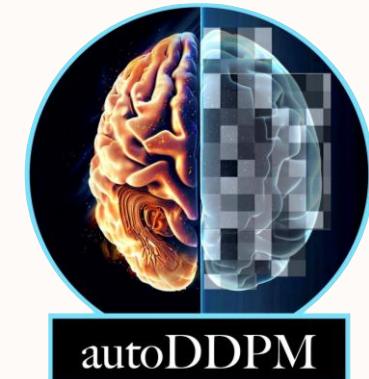
MICCAI 23



MIDL 23



MICCAI 23



ICML W 23

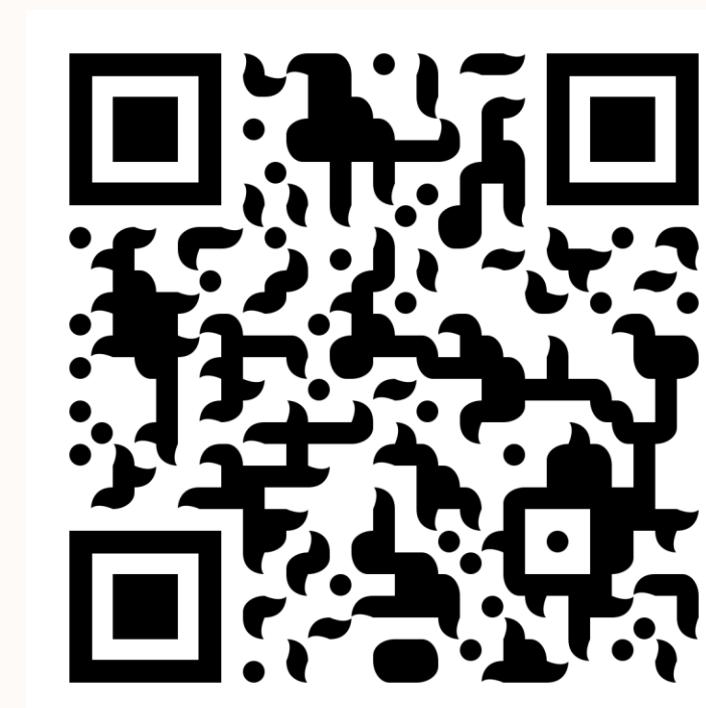


MICCAI 24

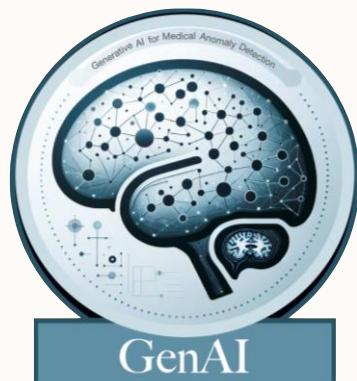


Nature

Machine Intelligence 22



Open source - Code is
available on GitHub:

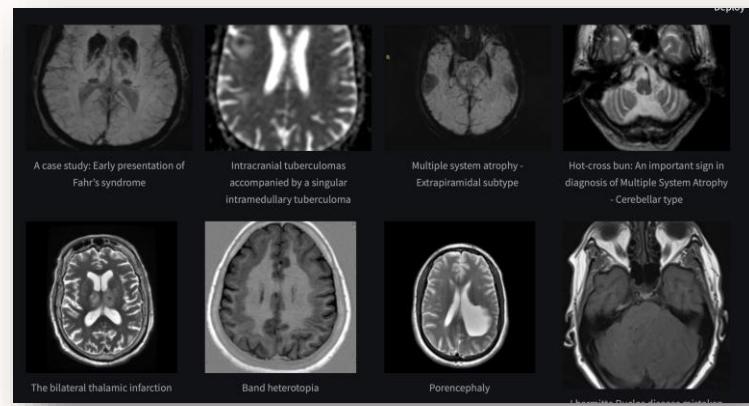


Nature

Communications 24



THE 2ND MICCAI STUDENT BOARD (MSB) WORKSHOP



Medical Out-of-Distribution Analysis Challenge – MICCAI 25

Moo(D) – CoW Track Real rare clinical anomalies

<https://zenodo.org/records/15083914>

From Noise to Knowledge



Controlled Generation

Diffusion models allow complex medical images to be synthesized from noise with greater precision.



Language Guided Editing

Edit medical images based on natural language instructions, enabling precise alterations.



Interactive Question/Answers Explainable AI through Language Diffusion

Answer clinical questions and provide explanations with integrated AI.



Multimodal Integration

Combine different imaging modalities for a comprehensive patient view



Clinician In-The-Loop

Integrate clinician expertise with AI tools for collaborative diagnosis and treatment planning.



Detecting Unknown Unknowns

Identify rare anomalies by learning normal patterns and highlighting deviations.